

Liquidity, Activism Disclosure, and Takeover Premia

A Theory of Exit, Voice, and Corporate Control

Austin Li

This draft: 30 August 2026

Abstract

Disclosure rules change what a market knows while an investor is building a stake. I study how that timing affects bidder entry and the takeover premium. In the model, a blockholder first trades without a public flag. If her stake crosses the legal threshold and the filing window expires, the market sees her stake and intent; otherwise it sees only order flow. This split gives the threshold and the filing window different economic roles. At fixed trading policies, a tighter threshold reduces the sensitivity of the engagement-related premium to liquidity. A shorter window does so only when the change in who remains hidden is not too large. The model therefore gives no unconditional sign for the 2024 US acceleration. Seeded filing data show that median delay fell from seven to five business days. I set out a pre-specified design to test the market and takeover implications.

1 Introduction

A blockholder who learns that a company is worth more than the market price can sell, wait, or try to change the company. If she builds a large enough position while she waits, the law makes her disclose it. In the United States, Section 13(d) of the Williams Act requires a holder who crosses 5 percent of a voting class to file a Schedule 13D. Until February 2024 the initial deadline was ten calendar days; it is now five business days (U.S. Securities and Exchange Commission, 2023). That rule gives me two policy variables. The threshold determines which positions become public. The window determines how much trading can take place before the filing. A lower threshold and a shorter window can therefore change different parts of the market’s information.

The central point is that disclosure divides the histories seen by the bidder and by market makers. After a filing, the market sees the blockholder’s stake and her intention to engage. Before a filing, it sees order flow and must infer whether an informed holder is trading. I call these the flagged and pooled cells. The distinction matters because a bidder responds to both the expected value of the firm and the chance that engagement has raised it. Liquidity changes the inference in the pooled cell, but it need not change what the market learns from a completed filing.

Several literatures contain pieces of this argument. Liquidity can finance monitoring (Maug, 1998), have no effect under free entry (Kahn and Winton, 1998), or weaken the information in an exit threat (Admati and Pfleiderer, 2009). Trading models study pooled and revealed orders without a legal disclosure rule (Kyle and Vila, 1991). With a fixed disclosure horizon, shortening the horizon can be equivalent to reducing noise-trading volatility (Back et al., 2018); other work places a multi-trader game inside a pre-disclosure window (Cetemen et al., 2026). None of these papers combines the two legal margins with liquidity and an endogenous control outcome. That combination is the contribution here.

The model has three main implications. First, a cutoff equilibrium over complete trading plans exists under conditions on the plan menu, prices, and the best-response map. The exact construction and its proof are in the online appendix. Two of the conditions fail at the implemented calibration, so the result does not establish existence there. Twenty-three of twenty-seven numerical sweep nodes converge, and the other four remain unresolved. I make no nonexistence claim.

Second, at fixed trading policies, the expected engagement-related premium has a simple decomposition: only the pooled cell contributes to its liquidity sensitivity. A tighter threshold moves histories into the flagged cell and therefore reduces that sensitivity. A shorter window has a weaker result. It reduces the weight on the pooled cell, but it also changes which histories remain there. The window reduces sensitivity precisely when the weight effect dominates that composition effect. This is the fixed-policy result stated in Theorem 1.

Third, allowing the blockholder’s policy to respond to liquidity adds a feedback term. The fixed-policy sign survives for the threshold margin on a region where the best-response map contracts and the direct effect dominates the feedback bound. The region is a condition of the result, not an object I claim to have found. Eighteen of eighty grid nodes satisfy the pointwise inequalities, with a largest contiguous block of nine. Proposition 2 states this implication.

The calibration gives a useful but limited guide to the 5 February 2024 filing-window change. Holding policies fixed and shortening the window from ten to five days reduces the computed sensitivity to between 18 and 77 percent of its long-window value at every checked node. The composition effect accounts for most of the reduction. At the same calibration, however, the pooled

posterior has 23 to 767 distinct values at each of 180 non-degenerate nodes, rather than the three-point support required by the threshold comparison. The directional window result is therefore numerical evidence, not a theorem about the implemented two-round model. The empirical design keeps attenuation and amplification as live possibilities.

The first descriptive fact is already visible in the filing data. In a seeded sample, the median trigger-to-filing delay fell from 7.0 business days before the February 2024 acceleration to 5.0 afterwards, and the share filed within five business days rose from 35.7 to 75.6 percent. The remaining empirical work is a written specification rather than a set of estimates. It tests the run-up and filing-day jump separately, uses pre-period filing behaviour as a bindingness dose, studies the stake at filing, bounds the accumulation channel’s effect on the bid hazard, and estimates bidder entry by liquidity. The specification reports both signs and treats the preliminary sample as descriptive.

Section 2 describes the rule and the literature. Section 3 lays out the two-round game in words and notation. Section 4 records the equilibrium construction briefly; the technical argument is in the appendix. Section 5 gives the two comparative-statics results and the calibration record. Section 6 presents the descriptives and the pre-specified tests. Section 7 closes. The proofs sit in the online appendix.

2 Institutional setting and related literature

2.1 The disclosure rule and its two margins

Section 13(d) of the Williams Act of 1968 requires any person or group acquiring beneficial ownership of more than 5 percent of a class of registered voting equity to file a Schedule 13D disclosing the stake and its purpose. From 1968 until 2024 the initial deadline was ten calendar days. The SEC proposed shortening it in February 2022, adopted the amendments in October 2023, and made them effective February 5, 2024: the initial 13D deadline became five business days, amendments on material change became two business days, and the parallel changes for Schedule 13G took effect September 30, 2024 (U.S. Securities and Exchange Commission, 2022; U.S. Securities and Exchange Commission, 2023). Item 4 of the form asks the filer to state its purpose; control-intent filers use the 13D, while the shorter 13G is open to three eligible classes (exempt investors, qualified institutional investors, and passive holders below 20 percent who certify no control intent). Which of the two a given block lands in is a choice the rule frames rather than an accident of the holder, and eligibility for the short form is itself a type variable rather than a random one.

Both margins of the rule vary across jurisdictions, and the variation is not subtle. The United Kingdom sets the initial threshold at 3 percent with notification at every additional percentage point, within two trading days. The EU Transparency Directive sets a ladder at 5 percent and above with four trading days, and several member states start at 3 percent. The threshold margin is therefore a cross-section object and the window margin a time-series one. Other jurisdictions have moved their windows too, the EU cutting its Transparency Directive deadline to four trading days in 2013 among them; what recommends the February 2024 acceleration as an anchor is that it is a dated change to one margin, in one jurisdiction, with a clean before and after. That is why it anchors this paper and why it has begun to anchor others (Polk et al., 2024; Trivedi, 2026).

Two institutional facts discipline the model. First, the window is a *legal deadline attached to a partition*, not a random horizon: the clock starts when the stake crosses the threshold, the filing

lands at most T days later, and between the crossing and the filing the market prices the stock without the flag. The model takes the deadline itself as the filing date, so filing early is outside what it claims. Second, the pooled cell is pooled *for the price-setting market*. Zeng (2026) documents that a target’s own managers learn of the accumulation through stock-surveillance vendors before the filing, and that the price run-up begins on the trigger date. The cell is leaky at its edges, which is a measurement problem for the empirics, not a failure of the partition for the market makers who set prices. On the delay distribution itself, Bebchuk et al. (2013) show filers bunching at the old deadline: 45.3 percent of filings by engaging investors in their sample land in the eight-to-ten-*business-day* bucket and roughly a fifth on the tenth business day itself, their delays being counted in business days against a deadline the rule states in calendar days. So the window binds for a large share of filers, which is what makes the February 2024 acceleration a treatment rather than a formality.

2.2 Exit, voice, and liquidity

The classic trade-off originates with Hirschman (1970). Coffee (1991) and Bhidé (1993) argue that liquidity weakens governance by letting blockholders sell instead of fight; Maug (1998) reverses the logic: liquid markets let a monitor recoup the costs of engagement through informed trading, and with an endogenous stake liquidity raises equilibrium monitoring. Kahn and Winton (1998) show the sign of the trading motive depends on whether engagement deepens the blockholder’s information advantage; Faure-Grimaud and Gromb (2004) make price informativeness and insider incentives complements; Admati and Pfleiderer (2009) make exit itself a disciplining threat. On the empirical side, Edmans, Fang, et al. (2013) find more liquid targets shift from voice to exit, and Norli et al. (2015) find liquidity predicts the incidence of costly voice.

Two limitations of this strand shape the present paper. The four classics use four different liquidity parameters (camouflage, noise volume, informativeness, exit noise) with three different signs, so “liquidity and governance” is not one question. And none of the four has a bidder, an offer, or a premium: takeovers appear in Maug (1998) only as a relabelled intervention technology, Kahn and Winton (1998) and Faure-Grimaud and Gromb (2004) set control transfers aside in terms, and Admati and Pfleiderer (2009) simply has no bidder. The object this paper studies, the sensitivity of an expected takeover premium to liquidity and what the disclosure rule does to that sensitivity, is not stated in any of them.

2.3 Trading, disclosure, and the market for control

The trading model builds on discrete order flow in the Kyle tradition (Kyle, 1985; Glosten and Milgrom, 1985), in the tractable form of Edmans, Goldstein, et al. (2015), and on the feedback between prices and real decisions (Edmans, Goldstein, et al., 2012; Bond et al., 2012; Goldstein, 2023). Grossman and Hart (1980) make the takeover premium the welfare object for dispersed minority shareholders, and Burkart, Gromb, et al. (1998) show why higher premia protect them.

Against the closest theory papers the position is specific. Kyle and Vila (1991) have an endogenous pooled/revealed split in mixing equilibria but disclaim disclosure requirements on their first page, and their traders disclose nothing at a legal deadline; the defensible contrast is that their split is endogenous, un-keyed, and available only where the trader happens to mix, while the rule-

keyed partition here is imposed by a stake trigger and a deadline in every parameter configuration. Collin-Dufresne and Fos (2015) show measured price impact falls when informed blockholders trade against liquidity, and Collin-Dufresne and Fos (2016) fix a trading horizon, but the presence of the informed trader is common knowledge in their model; they name the inference problem and set it aside. Back et al. (2018) embed an engaged blockholder in a continuous-time Kyle model with a fixed horizon at which the stake is disclosed; for them a shorter horizon is a re-parameterisation of noise-trading volatility (“reducing the trading horizon T is isomorphic to reducing noise trading volatility”), there is no acquirer and no premium, and endogenizing the horizon is their own declared out-of-scope extension. Cetemen et al. (2026) study multiple blockholders timing trades and coordinating through market signals inside a pre-disclosure window they leave unmodelled (their own reading of their model), with liquidity the comparative static and no premium object. Corum (2025) has the closest reduced form of the rule, two scalars measuring how much of the stake is tradable before disclosure bites, but collapses threshold and window into them and has no second player. Corum and Levit (2019) put a takeover probability and a disclosure threshold in one model, but the threshold does three jobs at once (liquidity break, 13D trigger, pill trigger), no comparative static in it is stated, and the flagged state never occurs on path. Ordóñez-Calafí and Bernhardt (2022) study the optimal threshold level (crossing reveals the position, so in equilibrium the engaging blockholder’s stake is the smaller of the threshold and its unconstrained optimum, the threshold acting as an upper bound on the position and hence on trading profits), and they assert, without a proposition, the interaction between liquidity and the threshold on managerial discipline; the object here is a control outcome and the interaction is proved. Burkart and Lee (2022) have the premium and explicitly hand over the missing piece: they do not endogenize the toehold in anonymous pre-disclosure markets and leave the interaction of pre- and post-disclosure decisions to future work.

2.4 Structural and empirical work on activism

A mature structural literature estimates activism along complementary margins: Gantchev (2013) the sequential costs of campaigns; Johnson and Swem (2021) a dynamic reputation model in which buying and filing are one move; Albuquerque et al. (2022) the binary 13D-versus-13G choice jointly with announcement returns, with liquidity a control variable and no control outcome; Celentano and Levine (2025) an equilibrium model of activism, takeovers, and managerial discipline in which the 5 percent threshold enters once, as a calibration constant, and the window never appears. None of these models the trading-and-disclosure environment itself. Brav et al. (2008) and Greenwood and Schor (2009) document that the targets of engagement campaigns are unusually likely to be acquired, and that fact makes bidder entry the natural control outcome; Becht et al. (2017) show the returns internationally; Zeng (2026) measures what leaks inside the window.

On the February 2024 acceleration itself, two papers precede this one and neither occupies the cell. Polk et al. (2024) use pre-rule data only, a cross-section of self-selected delays offered as a projection with no control group, no post-rule data, and no standard errors anywhere in the paper. Their treatment variable counts calendar days against a rule they themselves state in business days. Trivedi (2026) runs a difference-in-differences on 333 filings with 13G controls and finds a first stage on compliance timing (the share filed within five business days rises by 34.8 percentage points, standard error 13.0) and nulls on spread and Amihud changes that carry no minimum

detectable effect; there is no control outcome in the paper, its sampling frame is never stated, and it is not peer reviewed. The cell this paper’s empirics occupies (a rule margin, a liquidity split, and a control outcome, with the timing split as the identity claim) remains open, and the specification in Section 6 is written so that every estimate it will report was specified before any of them existed.

3 The model

3.1 Object and partition

The disclosure rule divides the market’s information into two cases. A stake threshold τ and a filing window T determine whether the filing has arrived by the control decision. In the *flagged* cell, the market sees the stake and the intention to engage. In the *pooled* cell, it sees only order flow and must infer whether an informed blockholder is trading. The cells are exhaustive. The outcome of interest is the expected engagement-related premium $\Delta^{\text{act}}(\kappa, \tau, T)$. I also track the run-up before the filing and the jump when the filing arrives. A lower threshold is a tighter threshold margin; a shorter window is a tighter window margin.

The model uses the discrete order-flow structure of Kyle (1985) and Edmans, Goldstein, et al. (2015), competitive prices as in Glosten and Milgrom (1985), and the takeover premium as the dispersed shareholder’s control-outcome measure in the spirit of Grossman and Hart (1980). The contribution is the legal partition between those familiar pieces.

3.2 Timing: the two rounds

Trading runs over business days $d = 0, \dots, H$. Nature draws the firm’s standalone value and the blockholder’s private signal. The blockholder then chooses a complete contingent plan from a finite ordered menu. The plan fixes the stake path and whether it is meant to engage the firm.

In the first round, market makers see the plan’s order flow mixed with ternary noise. The blockholder does not revise the path in response to realised prices or order flow. The filing arrives exactly T business days after the first crossing of τ , provided that date is before the horizon. That filing ends the pooled round. In the second round the blockholder trades the residual position, the market updates its price, and the bidder decides whether to enter. If no filing arrives, the bidder acts on the pooled history. The technical appendix spells out the off-path beliefs and the regularity conditions needed for this timing.

3.3 Primitives

The firm’s standalone value is $v \sim N(\mu_v, \sigma_v^2)$. The blockholder observes $s = v + \varepsilon$, where $\varepsilon \sim N(0, \sigma_\varepsilon^2)$ is independent of v . Her posterior mean is $\mathbb{E}[v \mid s] = \mu_v + \beta(s - \mu_v)$, with $\beta = \sigma_v^2 / (\sigma_v^2 + \sigma_\varepsilon^2) \in (0, 1)$. A potential bidder has an independent synergy shock $\xi \sim N(0, \sigma_\xi^2)$, mean synergy $\bar{S}$, and entry cost $K > 0$. Engagement changes the takeover premium from m_0 to m_1 , where $\Delta_m = m_1 - m_0 > 0$, and creates non-takeover value $\Delta_V \geq 0$.

Noise trading takes the values $-\bar{z}$, 0 , and $+\bar{z}$. The zero-noise probability is $1 - \kappa$ and each nonzero value has probability $\kappa/2$, so $\kappa \in [0, 1]$ measures noise-trading intensity. The disclosure policy is (τ, T) , with $T \in \{1, \dots, H\}$; b_0 and $\bar{b}$ are the initial and maximum stakes.

The blockholder chooses a plan from a finite menu $\mathcal{J}$, ordered from least to most aggressive. Plan j has an engagement flag a_j , equal to one for Voice and zero for Exit or Hold. Its pooled stake at date d is $B_j(s, d)$, with $B_j(s, -1) = b_0$, and its terminal target is $b_j^*(s) = B_j(s, H)$. The first crossing of the threshold is $c_j(s; \tau) = \inf\{d : B_j(s, d) \geq \tau\}$, with value $+\infty$ when no crossing occurs. The filing date is $f_j = c_j + T$, and disclosure occurs when $a_j = 1$ and $f_j \leq H$. At a filing, the reported stake is $B_j^F = B_j(s, f_j)$ and the remaining order is $Q_j^F = b_j^*(s) - B_j^F$. A finite coarsening Γ maps the stake increment into the pooled order mark $q_{jd}(s) = \Gamma(B_j(s, d) - B_j(s, d - 1))$, so observed pooled flow is $X_d = q_{jd} + z_d$.

Market makers observe the pooled history $\mathcal{H}_d^P = (X_0, \dots, X_d; \text{flag landed by } d)$. A filing reports $F = (B^F, a = 1)$, and the flagged tuple adds the residual order, $S_F = (B^F, Q^F, a = 1)$. At the control node, public information is $\mathcal{H}_H^P$ in the pooled cell and S_F in the flagged cell; the bidder's ξ remains private. I write these cells as $\mathcal{C}_P$ and $\mathcal{C}_F$. They are exclusive and exhaustive. The engagement posterior is $\pi(\mathcal{I}) = \mathbb{P}(a = 1 \mid \mathcal{I})$, and equals one after a filing. The bidder enters with probability

$$p(\mathcal{I}) = 1 - \Phi\left(\frac{P + K + m_0 + \pi\Delta_m - \bar{S}}{\sigma_\xi}\right) \in (0, 1). \quad (1)$$

With B the entry indicator, the terminal shareholder payoff is

$$Y = (1 - B)(v + a\Delta_V) + B(P + m_0 + a\Delta_m), \quad (2)$$

and prices satisfy the inner fixed point $P(\mathcal{I}) = \mathbb{E}[Y \mid \mathcal{I}]$. Two conventions are part of the model rather than notation. First, $P_{-1}^P := \mathbb{E}[Y]$ is the pre-trading pooled price, which is needed whenever $c = 0$, as $T = H$ forces on every flagged history. Second, $P_{\text{ND}}(\mathcal{H}_{f-}^P) := P_{f-}^P$: the not-yet-disclosed price is the last pre-filing pooled price at the *same* realised order flow, whose history already carries “flag not landed by $f - 1$ ”, and not a never-disclosed counterfactual. Throughout, c^- and f^- denote the business days immediately before the crossing date and the filing date. The price-path objects are the run-up path $R_d = P_d^P - P_{c-}^P$, the cumulative run-up $R = P_{f-}^P - P_{c-}^P$, and the filing-day jump $J = P^F - P_{\text{ND}}$; R_d , R and J are unsigned, and J is not claimed to be κ -invariant.

The genuine fixed point sits at the control nodes. At an earlier pooled date $d < H$ the price is a tower expectation of already-solved control-node values, with no self-reference; only the control-node map is a fixed point to be solved.

The blockholder evaluates a plan by the expected terminal value of the position it builds, less the trading outlay and the cost of engagement:

$$U_j(s) = \mathbb{E}\left[b_j^*(s)Y - \mathcal{C}_j^{\text{trade}} - a_j C_j(s) \mid s, j\right], \quad (3)$$

where $\mathcal{C}_j^{\text{trade}}$ is the execution outlay, with increments valued at the pooled prices P_d^P up to the plan's last pooled date and $Q_j^F(s)P^F$ when $D_j = 1$, and $C_j(s) \geq 0$ is the engagement cost. Only two properties of (3) are ever used: *plan-locality*, that U_j depends on j only through the executed stake path, the prices paid on it, the terminal stake, the engagement flag and the cost; and *integrability*, $\mathbb{E}[\max_j |U_j|] < \infty$ under the finiteness conditions of Section 3.3.

The cost in (3) is not tied to a particular date. It can be booked when the plan is completed or treated as sunk once the filing arrives. These readings give the same comparison among flagged round-two deviations, so the existence argument does not depend on the convention.

The primitive draws are mutually independent, with strictly positive variances. The menu, order-mark support, noise support, and calendar are finite; prices and payoffs are locally bounded on the maintained parameter set; and $\mathbb{E}[\max_{j \in \mathcal{J}} |U_j|] < \infty$ throughout the cutoff set.

Four restrictions on the primitive table travel together, and several results below consume them as a block. Clause (TR-i), *distributional forms and signs*: the distributions above hold as stated, with $\sigma_v^2, \sigma_\varepsilon^2, \sigma_\xi^2 > 0$, $K > 0$, $\bar{z} > 0$, $\Delta_m > 0$, $\Delta_V \geq 0$, and

$$m_0 \geq 0, \tag{4}$$

so that $\bar{m}(\mathcal{I}) = m_0 + \pi(\mathcal{I})\Delta_m \geq 0$. Restriction (4) is what makes the inner pricing root exist, be unique and be continuous; when it is dropped, executed counterexamples produce both no root and three roots. The initial stake satisfies $b_0 < \tau$, so a crossing that predates the model is outside the core analysis. Clause (TR-ii), *stake-path regularity*: Voice stakes are weakly increasing in the calendar date and in the signal, Hold stakes are constant, Exit stakes are weakly decreasing in the date, and every map $s \mapsto B_j(s, d)$ is Borel. Borel regularity is automatic for Voice and for Hold and is a genuine addition for Exit, where the primitives supply monotonicity in the date alone; without it the pooled prices are not defined, because pooled pricing integrates over every type. Stake levels themselves remain continuum-valued: the finiteness above covers the menu, the image of Γ , the noise support and the calendar, not the stake level. Clause (TR-iii), *legal-clock objects*: the crossing, filing, stake-at-filing, residual-order and terminal-target definitions hold as stated, together with $D = 1 \Rightarrow a = 1$, $Q^F \geq 0$ for Voice plans, and, at fixed policies, $T' < T$ implying $B^F(T') \leq B^F(T)$ and $Q^F(T') \geq Q^F(T)$. Clause (TR-iv), *pricing and entry*: the terminal payoff is (2), prices obey $P(\mathcal{I}) = \mathbb{E}[Y \mid \mathcal{I}]$, entry obeys (1), the control-node information set is the fill given above, and the two price conventions $P_{-1}^P = \mathbb{E}[Y]$ and $P_{\text{ND}} = P_f^P$ hold.

3.4 Equilibrium notion

I use a cutoff perfect Bayesian equilibrium, and it carries six requirements: (i) a weakly ordered cutoff vector $k = (k_1, \dots, k_{J-1})$ selects the plan as a function of the signal; (ii) the pooled and flagged actions are sequentially optimal at every history; (iii) beliefs obey Bayes' rule at every history reached with positive probability; (iv) prices solve their competitive fixed points; (v) the bidder follows (1); and (vi) off-path beliefs are limits along one fixed full-support perturbation family over plans. Weak inequalities allow an action region, including Hold, to collapse. The outer cutoff map is solved on a compact ordered polytope, and uniqueness is not claimed.

The timing and primitive restrictions accompany every result. The cutoff construction requires adjacent plan payoff differences to cross at most once in the signal, with the preferred plan weakly increasing in the signal. It also requires a common compact bracket containing an indifference signal for every adjacent plan pair at every point of Θ , and a continuous, self-mapping outer best-response map $\mathcal{T}(k; \vartheta)$. The inner pricing root is unique and continuous in its belief summaries under (4); continuity of the composed price family in the cutoff vector is the part retained as a condition.

On a flagged history, $(B^F, Q^F, a = 1)$ identifies the informed plan component, so the remaining pooled order is independent noise conditional on that tuple. For the fixed-policy comparisons, the composed terminal target $s \mapsto b_{j(s)}^*(s)$ must be strictly increasing on the flagged signal region for every $k \in \Theta$. The existence construction uses the stronger statement that $(j, s) \mapsto (B_j^F, Q_j^F, a_j)$ is

injective over all flagged pairs, including pairs not selected by the cutoff policy.

The pooled-cell comparison uses a strong support restriction. Its posterior law must have the three-point form

$$\mathbb{E}[h] = A_0(\kappa) h(0) + A_{1/2}(\kappa) h(\bar{\pi}/2) + A_1(\kappa) h(\bar{\pi}), \quad (5)$$

with $A'_0 = A'_1 = A'_\kappa$, $A'_{1/2} = -2A'_\kappa$, and a non-positive chord whose absolute value is weakly increasing in $\bar{\pi}$. Two clauses carry the restriction's remaining bite: (τ -i), *kernel through posterior*, that h depends on the information set through the posterior alone; and (τ -ii), *κ -free support*, that the three support points, the upper point $\bar{\pi}$ included, do not move with κ . I write the restriction $A(\tau)$. Whether the two-round pooled cell satisfies the support condition is open; the implemented calibration does not.

A threshold comparison $\tau' < \tau$ at fixed policies and a common κ needs a bridge between the two pooled classes, written $A(\text{br})$, in five clauses. (br-i), *representation at both policies*: (5) holds under τ and under τ' , with endpoints $\bar{\pi}(\tau)$ and $\bar{\pi}(\tau')$ and coefficients $A'_\kappa(\tau)$ and $A'_\kappa(\tau')$. (br-ii), *κ -localisation*: all κ -dependence of M_P sits in the weights, so that $\partial_\kappa M_P = \Delta_m A'_\kappa C_h(\bar{\pi})$ exactly, with no composition remainder; that last step is derived rather than assumed, and the clause names the same object as (τ -i) against the reading $h = \pi p(\hat{v}, \pi)$ rather than adding an independent restriction. (br-iii), *coefficient stability*: $|A'_\kappa(\tau')| \leq |A'_\kappa(\tau)|$, whose weakest sufficient form is equality, reclassification changing which histories are pooled rather than the κ -responsiveness of the pooled weights. (br-iv), *endpoint linkage*: $\bar{\pi}$ is the chord endpoint of (5) and is the same weakly increasing function of the pooled prior engagement share $\bar{\pi}_{\text{pr}} = \mathbb{P}(a = 1 \mid D = 0)$ at τ and at τ' , the identity branch $\bar{\pi} = \bar{\pi}_{\text{pr}}$ being excluded as degenerate. (br-v), *comparability of the chord functional*: $C_h(\cdot)$ and the kernel behind it are the same functions of the posterior at both thresholds; without it the threshold comparison sets $|C_h|$ against a different functional and means nothing. These clauses are used only for the threshold sensitivity result. For the general-equilibrium result I require a twice continuously differentiable outer map on a candidate region $\mathcal{R}$, a contraction bound $L_{\mathcal{R}} < 1$, and a constant sign for the equilibrium liquidity derivative.

3.5 The cutoff polytope and the premium objects

The cutoff vector is $k = (k_1, \dots, k_{J-1})$ with weak ordering. The compact, convex set of such vectors is Θ , and ϑ collects the remaining parameters. The outer map $\mathcal{T}(k; \vartheta)$ gives the cutoff vector implied by a conjectured vector k . On a candidate region $\mathcal{R}$, write $L_{\mathcal{R}} = \sup_{\mathcal{R}} \|D_k \mathcal{T}\|$ and use $r_\tau = -\tau$ and $r_T = -T$ so that a larger r means a tighter rule. The fixed-policy attenuation margin is

$$g_r^{PE} = -\text{sgn}(d\Delta^{\text{act}}/d\kappa) \partial_{\kappa r} \Delta^{\text{act}}, \quad (6)$$

and the inversion-free cutoff bounds are $\bar{k}_x = |\partial_x \mathcal{T}|/(1 - L_{\mathcal{R}})$ and $\bar{k}_{\kappa r}$. The general-equilibrium remainder is

$$\mathcal{B}_r^{GE} = |\Delta_{\kappa k}| \bar{k}_r + (|\Delta_{kr}| + |\Delta_{kk}| \bar{k}_r) \bar{k}_\kappa + |\Delta_k| \bar{k}_{\kappa r}. \quad (7)$$

Here the Δ terms are the gradient and cross-partial of Δ^{act} in (k, κ, r) along the equilibrium branch. The dominance-and-contraction region $\mathcal{R}_r$ has slack $\eta_r = g_r^{PE} - \mathcal{B}_r^{GE}$ and may be empty.

The engagement component of the premium is built from $h(\mathcal{I}) = \pi(\mathcal{I})p(\mathcal{I})$. The kernel satisfies

$h \geq 0$ and $h(0) = 0$, and the expected engagement-related premium is

$$\Delta^{\text{act}} = \Delta_m \mathbb{E}[h(\mathcal{I}_H)] \geq 0. \quad (8)$$

The cell averages are $M_F = \Delta_m \mathbb{E}[h \mid D = 1]$ and $M_P = \Delta_m \mathbb{E}[h \mid D = 0]$ when the corresponding cell has positive mass. The flagged weight is $\Omega = \mathbb{P}(D = 1) = \mathbb{P}(a = 1)\omega_a$, where $\omega_a = \mathbb{P}(D = 1 \mid a = 1)$. In the three-point representation, $\bar{\pi}$ is the upper support point, not the mean pooled engagement share. The latter is κ -invariant under the support restriction and is strictly below $\bar{\pi}$ away from the degenerate case; it equals $\bar{\pi}/2$ only when $A_0 = A_1$. Define $\mathcal{S} = |\partial_\kappa \Delta^{\text{act}}|$ and $\mathcal{S}_P = |\partial_\kappa M_P|$. The chord is

$$C_h(\bar{\pi}) = h(0) - 2h(\bar{\pi}/2) + h(\bar{\pi}), \quad (9)$$

maintained non-positive with $|C_h|$ weakly increasing in $\bar{\pi}$, and A'_κ is the common derivative of the weights in (5), bounded on $[0, 1]$. The weight-effect ratios are W_τ and W_T , for instance $W_T = (1 - \Omega(\tau, 5))/(1 - \Omega(\tau, 10))$, which is at most 1 when Ω rises. The composition-effect ratios are C_τ and C_T , for instance $C_T = \mathcal{S}_P(\tau, 5)/\mathcal{S}_P(\tau, 10)$, which are unsigned. The margin subscript on a composition ratio is always written, because C is otherwise overloaded by the chord C_h , the engagement cost $C_j(s)$ and the cells $\mathcal{C}_F, \mathcal{C}_P$.

Reading $\bar{\pi}$ as the mean of the pooled posterior law rather than as its upper support point forces a point mass at $\bar{\pi}$ with $A'_\kappa = 0$ and zero interior motion for every kernel. That is degenerate, and it is excluded throughout.

4 Equilibrium existence

Existence is standing ground here, not the claim. The claim the paper is making is the attenuation result of Section 5; what this section supplies is the conditional on which the whole construction rests.

Under the hypotheses (P-1)–(P-13) below, a cutoff perfect Bayesian equilibrium over complete contingent plans exists at *every* $\kappa \in [0, 1]$: there is $k^* \in \Theta$ with $k^* = \mathcal{T}(k^*; \vartheta)$, prices at their inner fixed points, and beliefs Bayes-consistent on path. Uniqueness is not claimed, in any form.

Five things travel with that conclusion. Off-path beliefs are limits of *one* full-support perturbation family over plans, fixed once and used to define the price system at every $k \in \Theta$ and not only at k^* , because the deviation payoffs that define $\mathcal{T}$ read off-path pooled histories. They are supplied at every pooled history reachable with positive probability under some plan profile. The boundary values of κ are handled by *extension* rather than by restriction: at $\kappa \in \{0, 1\}$ the noise support degenerates to $\{0\}$ and to $\{\pm \bar{z}\}$, and a pooled history needing a mark outside the surviving support is null under every profile, so it lies off nature's path rather than off the players', carries no requirement under clause (vi) of the equilibrium definition, and is read by no step. No cut to $\kappa \in [0, 1)$ is taken. Flagged-tuple beliefs are the point mass at the unique generating pair, supplied at every tuple in the image of $(j, s) \mapsto (B_j^F, Q_j^F, a_j)$, on path and off, since that image includes tuples the cutoff vector does not select; the point mass is a *version* of the conditional law at every image tuple, which is what a conditional law is when the signal is continuous, and any a.e.-equal version serves clauses (iii) and (vi) equally. No tuple outside the image arises, because the round-two action-set hypothesis (P-9) leaves no off-menu order to produce one. The bidder follows (1).

And there is a sequentially optimal flagged component at *every* flagged pair (j, s) , selected or not: the flagged price is invariant across the round-two deviation set, so the flagged order cancels out of the continuation and (P-10) makes what remains constant.

The hypotheses. The list is long, and every clause of it is load-bearing, so it is enumerated rather than gathered into a sentence.

- (P-1) Independent primitives.
- (P-2) The finite model and integrable objective.
- (P-3) Ordered plans and single crossing.
- (P-4) Legal-clock discipline.
- (P-5) The compact continuous outer self-map, read as set out below.
- (P-6) *Joint tuple injectivity*: the map $(j, s) \mapsto (B_j^F, Q_j^F, a_j)$ is injective on the whole flagged-pair set $\{(j, s) : D_j = 1\}$, *including pairs no cutoff vector selects*. This is strictly stronger than the on-path form, and it is the form the argument consumes where it pins *off-path* flagged beliefs.¹
- (P-7) The disclosure partition and price-path decomposition of Section 5.1, by statement, *with its own hypotheses travelling*.
- (P-8) The no-feedback timing, read together with the flag-termination clause.
- (P-9) The *definitional* round-two action-set hypothesis: the flagged-round action set *is* the plan-generated set $\{Q_{j'}^F(s)\}$ taken over menu elements agreeing with j on everything already played. This is not a closure condition; the closure form is jointly unsatisfiable with finiteness, by cardinality.
- (P-10) *Continuation-cost equivalence on that same set*: menu elements sharing j 's pooled path up to $f_j(s)$ and carrying $a_{j'} = a_j$ have the same engagement cost, $C_{j'}(s) = C_j(s)$. It holds trivially on any single-Voice menu, where the set is a singleton, and is live only on menus with two or more Voice plans sharing a pooled path. What it buys, under the plan-completion reading of the cost timing, is this: on that set the flagged price does not move and the flagged order cancels, so the engagement cost is the only thing that can differ between staying and deviating, and at a flagged pair the cutoff vector does not select there is no date-0 optimality to fall back on, so without the clause the deviator takes the class member with the smallest cost and clause (ii) of the equilibrium definition fails at that node. Under the sunk reading the continuation is constant on the deviation set with no clause at all and the hypothesis is not consumed; it is listed because the statement does not commit to a reading, and it is what makes the conclusion hold under both.
- (P-11) The restriction $m_0 \geq 0$ of (4).

¹The distinction is not cosmetic. An earlier recorded statement of this result carried the on-path form while the proof consumed the joint form. The joint form is stated here because the proof uses it to pin off-path flagged beliefs; the on-path form alone would not suffice. Which form is meant is named every time either appears.

- (P-12) The blockholder’s objective (3), whose $-a_j C_j(s)$ term carries the engagement cost, together with the cost timing convention: the cost may be booked on completing the plan or treated as sunk once the filing has landed, the two give the same round-two comparison on the round-two deviation set, and the result does not depend on the choice.
- (P-13) The primitive table restrictions the argument consumes, in particular: the terminal payoff (2) with the price convention $P(\mathcal{I}) = \mathbb{E}[Y \mid \mathcal{I}]$ and the entry rule (1), both from clause (TR-iv); clause (TR-ii)’s Borel regularity for *every* plan including Exit, needed here *directly* and not by way of the partition result, whose conclusion is measurability of D and of the cell map; clause (TR-iii)’s implication $D = 1 \Rightarrow a = 1$, the definitions of c , f , B^F , Q^F and b^* , and $\partial_s B_j \geq 0$ for Voice; and clause (TR-i)’s distributional forms with $\Delta_m > 0$.

Two readings. Both belong to the statement rather than to the proof. First, the pricing assumption of Section 3.4 *is not assumed here*: its existence and uniqueness content is derived from $m_0 \geq 0$, its continuity content *in the belief summaries* $(\hat{v}, \pi)$ from the same scalar reduction, and its measurable-selection content from joint tuple injectivity together with clause (TR-ii). What is *not* derived is its cutoff clause, the continuity of the composed pooled price family in the cutoff vector k , which runs through the conditioning pair $(\hat{v}, \pi)$ rather than through the pricing map. That continuity enters only through the second reading, and Section 5.6 records it measured to fail at the implemented calibration. Second, the outer-map assumption *is read* as asserting that $\mathcal{T}$, under a named tie-break-and-corner selection, without which a correspondence cannot be called continuous, is a well-defined, single-valued, continuous self-map of the nonempty polytope Θ ; and, in its *bracket* half, that there is a fixed compact signal interval, **the common bracket** $[s_{\text{lo}}, s_{\text{hi}}]$ at the given ϑ , common in $k \in \Theta$ and in the adjacent plan pair, of which Θ is the ordered part $(s_{\text{lo}} \leq k_1 \leq \dots \leq k_{J-1} \leq s_{\text{hi}})$, such that for every $k \in \Theta$ and every adjacent plan pair the best-response indifference signal **exists** and lies in the bracket. That is what makes the finite cutoff vector represent the best-response plan map on the whole Gaussian signal line and not merely inside the bracket. The clause excludes every menu-and-price-system at which some adjacent indifference signal fails to exist, never-selected middle plans as well as dominated extreme plans; the stronger form is taken so that the cutoff vector carries interior data at every k , and the exclusion is a named domain narrowing.

Finally, at any such equilibrium at which the interior-crossing condition holds, both cells carry strictly positive probability and both are on path; restating that condition as a single signal threshold needs two further hypotheses, Voice stake monotonicity across plans and engagement flags equal to one *exactly* on an upper set of the ordered menu.

How the proof goes. The argument builds the equilibrium object rather than characterising it, and the online appendix carries it in full. A conjectured cutoff vector induces a measurable plan-selection map, under which every date-0 object is a deterministic Borel function of the pair (j, s) . The inner price system at that conjecture separates a finite pooled layer from a continuum-indexed flagged layer: the scalar reduction behind (4) gives existence, uniqueness and continuity of the inner root, and the flagged layer is priced by inverting the flagged-pair map, which joint injectivity makes invertible. Beliefs are then supplied at every history that carries a requirement, Bayes where the reaching weight is positive and the fixed perturbation family’s limit elsewhere,

and sequential optimality of the flagged component is discharged at every flagged pair, selected or not. The outer best-response map is built on the common bracket; that it is weakly ordered and carries Θ into itself follows from single crossing, while the bracket and the continuity of the map are taken rather than derived. Brouwer’s fixed-point theorem then applies, and the six requirements of the equilibrium definition are assembled at the fixed point, with the boundary values of κ handled by extension. The proof closes with six disclaimers: uniqueness in no form; the interior-crossing addendum conditional at the equilibrium and nowhere else; boundary extension, which is where a withdrawn claim of a belief at every pooled history at $\kappa = 1$ sat; one fixed convention; flagged-order optimality without uniqueness; and two hypotheses that are sufficient rather than necessary and are not verified menu by menu.

5 Liquidity and control outcomes

This section states what the model delivers about the control outcome. The order is the order the model builds it in: the partition and its accounting identities, then what liquidity does inside each cell, then the attenuation theorem, then what survives the move from fixed policies to equilibrium. Section 5.6 then reads the whole stack against the implemented calibration and says plainly what it makes of the antecedents. Each statement carries its conditionality with it; the full numerical records sit in the online appendix.

5.1 The partition and its accounting identities

Three facts about the partition come first, and the rest of the section rests on them. They hold under the independence, finiteness, legal-clock and pricing conditions of Section 3, the cutoff selection map of clause (i) of the equilibrium definition, and the two price conventions $P_{-1}^P = \mathbb{E}[Y]$ and $P_{\text{ND}} = P_{f-}^P$; part (c) also consumes clause (TR-ii)’s Borel regularity of $s \mapsto B_j(s, d)$ for *every* plan including Exit, because pooled pricing integrates over every type, and a content for the control-node information set. (a) The disclosure indicator $D = \mathbf{1}\{a = 1, c(\tau) + T \leq H\}$ is measurable, and it maps every control-node history into exactly one cell. (b) For every Voice plan, $f_j \leq H$ if and only if $B_j(s, H - T) \geq \tau$. (c) Each flagged history yields B^F , R_d , R and J , and these satisfy the exact identity

$$P^F - P_c^P = R + J. \quad (10)$$

Part (b) is the *clock equivalence*, and the threshold and window results both consume it: it converts a statement about the filing date into a statement about the stake path T days before the horizon, which is what makes window comparisons tractable. Identity (10) is the model counterpart of the timing split the empirics of Section 6 measure: the total price move from just before the crossing to the flagged price decomposes exactly into the cumulative run-up plus the filing-day jump.

The premium then decomposes across the two cells. Given the partition, the definitions of Section 3.3, the version of $\mathbb{E}[Y \mid \mathcal{I}]$ pinned by the pricing assumption, finiteness of Δ_m , and one

common probability space, whenever $0 < \Omega < 1$,

$$\Delta^{\text{act}} = \Omega M_F + (1 - \Omega) M_P. \quad (11)$$

At $\Omega = 1$ this degenerates to $\Delta^{\text{act}} = M_F$ and at $\Omega = 0$ to $\Delta^{\text{act}} = M_P$; in each degenerate case the average over the null cell is *undefined rather than imputed*. That last clause is a non-identification statement, and it is proved rather than asserted.

5.2 Liquidity in the flagged cell

Liquidity does nothing inside the flagged cell. At fixed cutoff *and* execution policies, under the same standing conditions together with the on-path injective identification condition consumed almost surely on the flagged set,² the partition result above, the no-feedback timing, $\Omega > 0$, and an explicit bidder-entry rule, either (1) or any rule with the two properties named in the proof and carried as bookkeeping, the flagged tuple $S_F = (B^F, Q^F, a=1)$ makes the pre-filing pooled history conditionally independent of (v, s, ξ) on the flagged set. The flagged posterior, the flagged price, the entry probability and M_F are therefore all invariant to κ .

The weak identification wording of the identification assumption is not sufficient here: it permits two pairs (j, s) with different pooled paths, which is this result's first failure case. The form consumed is the on-path injective one, named as such. This is the theoretical statement behind the empirics' identification claim: the flagged half of the timing split should not move with liquidity, and the pooled half should carry the entire liquidity derivative.

5.3 Liquidity in the pooled cell

Inside the pooled cell liquidity does move the premium, but only under the support restriction $A(\tau)$, and the restriction is the substance of the result rather than a caveat on it. Assume $A(\tau)$, both clauses $(\tau\text{-i})$ and $(\tau\text{-ii})$; $h(0) = 0$; κ -free pooled mass and κ -free pooled engagement moment at fixed policies; the partition result by statement; and regularity stated minimally, h continuous on $[0, \bar{\pi}]$ and twice differentiable on $(0, \bar{\pi})$, with Darboux's theorem doing the rest and no continuity of h'' required. Then (a) $\partial_\kappa \mathbb{E}_\kappa[h] = A'_\kappa C_h(\bar{\pi})$, exactly; and (b) $C_h(\bar{\pi}) = \frac{1}{4} \bar{\pi}^2 h''(\zeta)$ for some $\zeta \in (0, \bar{\pi})$, an identity rather than an approximation. Adding a second-order Peano expansion of h at $0+$ and one and the same kernel along a shrinking family gives (c) $C_h = \frac{1}{4} h''(0) \bar{\pi}^2 + o(\bar{\pi}^2)$, so the interior motion vanishes at rate $\bar{\pi}^2$ as $\bar{\pi} \downarrow 0$; the seam at which the threshold result consumes part (c) needs $|A'_\kappa|$ bounded *uniformly in $\bar{\pi}$* along the limit. The statement is an “if” and never an “iff”: $A'_\kappa = 0$ also kills the interior motion.

Whether the two-round pooled cell of Section 3.2 satisfies the support condition of $A(\tau)$ is **open**: the online appendix shows that the support condition is the assumption's entire remaining bite, and exhibits a one-round market that satisfies it and a no-disclosure structure that does not.

Tightening the threshold then moves the pooled cell in three steps. Take $b_0 < \tau' < \tau$ at a common window T and a common κ , with $\Omega(\tau', T) < 1$, at fixed policies. Leg 1, which is unconditional: $\mathcal{C}_F(\tau, T) \subseteq \mathcal{C}_F(\tau', T)$, with every newly flagged history generated by a Voice plan, so $\Omega(\tau', T) \geq \Omega(\tau, T)$. Leg 2, also unconditional: the pooled engagement *share* falls, $\bar{\pi}_{\text{pr}}(\tau') \leq \bar{\pi}_{\text{pr}}(\tau)$,

²“Almost surely” is the permissive reading, and it is the only coherent one when B^F is continuum-valued, since then no individual flagged tuple has positive probability.

with an exact identity for the difference between the two shares. Leg 3, under A(br): $\mathcal{S}_P(\tau', T) \leq \mathcal{S}_P(\tau, T)$, with equality whenever $C_h(\bar{\pi}(\tau)) = 0$. Legs 1 and 2 need only the clock equivalence, the no-feedback timing, fixed policies, $b_0 < \tau' < \tau$ imposed at *both* thresholds, independence, legal-clock discipline, clause (TR-iii)'s implication $D = 1 \Rightarrow a = 1$, and $\Omega(\tau') < 1$; nestedness is a *conclusion* of leg 1, not a hypothesis of it. Leg 3 also needs the pooled interior-motion result by statement, the maintained *magnitude* monotonicity of $|C_h|$ in $\bar{\pi}$ from A(τ), the sign half $C_h \leq 0$ being unused at this leg, and clauses (br-i)–(br-v) of A(br).

Leg 3 inherits the conditionality of A(br), and through clause (br-i) it inherits the representation (5) at *both* compared thresholds; the numerical record in the online appendix therefore bears on leg 3 as directly as it bears on the interior-motion result. At the implemented calibration leg 3's antecedent is not satisfied, and at that calibration leg 3 says nothing about the implemented pooled cell. Legs 1 and 2 are unaffected: they consume no chord restriction. Clause (br-iii) is the clause with the least justification behind it and is the one to attack first.

5.4 Disclosure attenuation

Theorem 1 (Disclosure attenuation at fixed policies). *At fixed plan and cutoff policies, with $0 < \Omega < 1$ and $\mathcal{S}_P > 0$, and under the hypotheses (T-1)–(T-15) listed below:*

- (A) Factorisation. $\mathcal{S} = (1 - \Omega) \mathcal{S}_P$, exactly; and the same factorisation holds for the total-variation aggregate of Δ^{act} over any κ -grid, with no differentiability required.
- (B) Threshold margin. *Threshold tightening attenuates:*

$$\frac{\mathcal{S}(\tau')}{\mathcal{S}(\tau)} = W_\tau C_\tau \leq 1, \quad (12)$$

because both ratios lie in $[0, 1]$. No dominance condition is needed.

- (C) Window margin. *Window tightening attenuates if and only if $W_T C_T \leq 1$, where $W_T \leq 1$ is proved — from the clock equivalence of Section 5.1 and the monotone Voice stake path — and C_T is unsigned. The equivalent form*

$$\frac{\partial_{r_T} \mathcal{S}_P}{\mathcal{S}_P} \leq \frac{\partial_{r_T} \Omega}{1 - \Omega} \quad (13)$$

holds on average along the tightening path, integrated over $[-T, -T']$, and exactly in the infinitesimal limit; read pointwise, (13) is false.

The hypotheses are: (T-1) fixed policies; (T-2) the interior-crossing condition at each compared policy; (T-3) $\mathcal{S}_P > 0$; (T-4) the two-cell decomposition (11); (T-5) the flagged-cell invariance of Section 5.2, with its own hypotheses travelling; (T-6) the partition result of Section 5.1; (T-7) $PE\text{-}\Omega$, that is $\partial_\kappa \Omega = 0$ at fixed policies — derivable rather than assumed, and it is exactly what fails in general equilibrium, which is the term Proposition 2 bounds; (T-8) κ -differentiability of M_P , which no standing hypothesis of Section 3 supplies and which is therefore carried in the proof; (T-9) the support restriction A(τ) at both compared policies, with $\bar{\pi}$ read as the upper support point; (T-10) the pooled interior-motion result of Section 5.3; (T-11) clauses (br-i)–(br-v) of A(br) at the threshold pair; (T-12) the three-leg threshold result of Section 5.3; (T-13) the no-feedback timing;

(T-14) a smooth window interpolation, for the local form (13) only; and (T-15) threshold-side smoothness, which has been confirmed non-load-bearing. No unconditional window sign is claimed.

Part (A) is the weight effect in its exact form: the liquidity sensitivity of the expected engagement-related premium is the pooled cell’s sensitivity, scaled by the pooled cell’s weight. Part (B) is the threshold-margin result and needs no dominance condition, because the weight ratio and the composition ratio both lie in $[0, 1]$. Part (C) is the window-margin result, and it is an “if and only if” rather than a sign: the composition ratio C_T is unsigned, so the model does not deliver window-margin attenuation as a theorem. The two margins of the disclosure rule therefore enter the policy discussion with different epistemic status, and the February 2024 acceleration, a window change, lands on the margin where the theory signs nothing.

The disclosure-regime comparison, and why it is not a window test. A separate numerical comparison is sometimes read as evidence about the window, and it is not one. In the one-round predecessor model the expected engagement-related premium can be computed twice: once with the market seeing the disclosure flag alongside order flow, once with order flow alone. The cutoffs k_1 and k_0 stay at the baseline equilibrium throughout and only the disclosure cutoff k_D moves; moving it across four values generates the flagged weights $\Omega = 0.037, 0.129, 0.286$ and 0.50 . The ratio of liquidity sensitivities (flagged divided by pooled, total variation over $\kappa \in [0.15, 0.85]$) is $1.064, 1.184$ and 1.136 at the first three of those weights, falling to 0.378 at $\Omega = 0.50$, with the sign crossing located by bisection at $\Omega^* = 0.343$. Above one, the flag makes premia *more* liquidity-sensitive. Figure 1 plots the two curves at the baseline. This is a *disclosure-regime* comparison (the flag observed versus hidden, at a fixed filing window), and its ratios are regime-comparison composition outcomes: they are not $W_T C_T$, they measure no window pair, and the predecessor model has no window to vary. The comparison earns its place for one reason: it shows that a composition factor can exceed one, which is what motivates the genuine window-margin “if and only if” of Part (C) and warns against reading attenuation into any flag comparison.




Figure 1: The disclosure-regime comparison in the one-round predecessor model, at fixed cutoffs ($\Omega = 0.037$). The solid line with circle markers prices the blockholder’s engagement-driven premium with the market seeing order flow and the disclosure flag (threshold disclosure, baseline); the dashed line with square markers prices it from order flow alone (no disclosure, counterfactual). The two curves cross near $\kappa = 0.73$. At this baseline the flagged cell is *more* liquidity-sensitive than the pooled one (total variation ratio 1.064), so this comparison refutes attenuation in κ for the regime margin at this weight. It is a composition fact, not a window test; the window record is Table 1.

5.5 From fixed policies to equilibrium

Theorem 1 holds at fixed policies, and hypothesis (T-7) is exactly what fails once the cutoff vector is allowed to move with κ . The next result asks what survives that move. It carries the region as a *named hypothesis* rather than exhibiting one.

Proposition 2 (The dominance-and-contraction implication in general equilibrium). *Assume:*

- (C-1) the along-the-path contraction $L_{\mathcal{R}} < 1$ of Section 3.4, in one fixed norm convention — an induced operator norm with its dual pairings; a mismatched pairing silently voids the implication;
- (C-2) $\mathcal{R}_r$ relatively open in both coordinates, with $\kappa \notin \{0, 1\}$;
- (C-3) an interior single-branch equilibrium;
- (C-4) twice continuous differentiability of Δ^{act} in (k, κ, r) ;
- (C-5) a non-vanishing equilibrium liquidity derivative — the equilibrium sensitivity $\mathcal{S}^{GE}$, which is distinguished from the fixed-policy sensitivity $\mathcal{S}$ of Section 3.5;
- (C-6) strict dominance $g_r^{PE} > \mathcal{B}_r^{GE}$ on the region, with g_r^{PE} as in (6) and $\mathcal{B}_r^{GE}$ as in (7);

(C-7) the threshold margin r_τ only: the window coordinate is an integer, and nothing local is claimed there.

Then the fixed-policy attenuation sign survives in equilibrium on the region:

$$\partial_r S^{GE} \leq -\eta_r < 0. \quad (14)$$

The sign-coherence hypothesis — that the fixed-policy attenuation sign agrees with the equilibrium sign — is confirmed unused in (14): its work is one of naming, fixing when that conclusion may be called a survival of the fixed-policy sign. The sign-constancy clause of the general-equilibrium condition is likewise unused.

5.6 The model at the implemented calibration

The results above are conditionals, and this subsection says plainly what the implemented calibration makes of their antecedents. The online appendix carries the full measurements.

Existence. The two main hypotheses about the equilibrium object are measured to fail at the calibration the implementation runs. The single-crossing clause fails at two independently found loci: at one, an adjacent-plan payoff difference crosses zero three times on an open set of the second cutoff; at the other, the preferred plan decreases in the signal across a cell edge. The outer map’s continuity fails for the polytope the proof constructs, with measured jumps of 6.33×10^{-3} to 2.83×10^{-2} across steps in the cutoff of at most 2×10^{-9} at the baseline node, and jumps reaching 0.16 at a second node. In plain words: the proof needs the outer best-response map to be continuous and single-valued, and at the implemented calibration that map has measured jumps and stretches where no weakly increasing selection exists at all. Nonexistence is neither claimed nor shown: twenty-three of twenty-seven equilibrium-sweep nodes converge, the remaining four stay unresolved after thirty seeds each, and a discontinuous self-map may still have fixed points.

The chord restriction. The restriction behind the pooled-cell results fails at the same calibration. The pooled cell’s engagement-posterior law, enumerated exactly over all 4^{H+1} order-flow paths at $H = 10$, at two hundred nodes, carries between 23 and 767 distinct posterior values where the restriction’s representation asks for three, at every one of the 180 non-degenerate nodes; the remaining twenty nodes are degenerate and decide nothing. The interior motion of the pooled premium, leg 3 of the threshold-nesting result and Part (B) of the attenuation theorem remain proved as conditionals; at this calibration they say nothing about the implemented pooled cell. The open question concerns the restriction’s domain: a different menu, a different horizon, or a different calibration could still satisfy the support condition.

The window record. At frozen policies, with cutoffs pinned at the baseline equilibrium, the flagged weight is $\Omega = 13.84$ percent, the disclosed share of engagements is $\omega_a = 61.15$ percent, and the cell averages are $M_F = 0.553$ and $M_P = 0.223$ premium percentage points. Cutting the window from $T = 10$ to $T = 5$ then cuts the premium’s liquidity sensitivity to between 18 and 77 percent of its long-window value at every checked node, and the composition leg carries almost all of it: C_T runs from 0.21 to 0.79 while the weight leg shaves 3 to 14 percent (Table 1). The same record confirms the flagged cell’s invariance exactly, with every flagged object invariant in κ to machine zero over the grid while the pooled M_P moves 0.0722 premium percentage points, and gives the threshold margin the same direction, with zero of eight compared pairs above one. Two caveats

travel with the record. At $H = 10$ the long window is the corner $T = H$; the $H = 12$ column of the table, where $T = 10$ is strictly interior, points the same way but travels the chord closed form whose support condition fails at this calibration, so it is directional corroboration rather than a second independent magnitude. And the record is node evidence at one calibration, not a sign theorem: Part (C) claims none. At this calibration the model’s directional prediction for the February 2024 acceleration is attenuation, and the empirics of Section 6 treat that as one of two live branches, not as a derived sign.

Table 1: The window-margin record at fixed policies: cutting the window from $T = 10$ to $T = 5$ at the implemented calibration. Weight leg W_T , composition leg C_T , product $W_TC_T \leq 1$ (attenuation), with no τ node above one. Verified on the numerical grid; sensitivity measured as the total variation of Δ^{act} over the κ grid. The $H = 12$ column re-runs the comparison with $T = 10$ strictly interior to the horizon, where the 0.9-quantile node returns 1.0000 because the τ ladder stops biting and the two flagged weights coincide, not because the product crosses one.

τ quantile	W_T	C_T	W_TC_T	W_TC_T at $H = 12$
0.1	0.8559	0.2124	0.1818	0.1099
0.3	0.8559	0.2124	0.1818	0.1099
0.5	0.8622	0.2384	0.2055	0.2406
0.7	0.9176	0.4686	0.4299	0.5772
0.9	0.9730	0.7939	0.7724	1.0000

The equilibrium region. For the move from fixed policies to equilibrium, eighteen of eighty grid nodes satisfy the two pointwise inequalities, the contraction bound below one and positive dominance slack, with the largest contiguous block, nine nodes, at $T = 5$. A node is a pointwise numerical fact, not evidence of a region: whether any parameter vector satisfies the full antecedent of Proposition 2 is open.

6 Empirics: motivating facts and the written specification

The empirical half of the paper has two layers with different statuses, and the draft keeps them apart. The motivating facts in Section 6.1 have been run: they are descriptives computed from the seeded filing sample, with the sample stated exactly and no standard errors claimed. The specification in Section 6.2 is prospective: it states the variables, windows, samples, standard errors, placebos, and decision rules before any estimate exists. The timing split is the specification’s default headline; final selection remains pending in Section 6.3.

6.1 Motivating facts: the delay distribution before and after the acceleration

The anchor is the February 2024 acceleration: the initial Schedule 13D deadline moved from ten calendar days to five business days for trigger dates from 2024-02-05 (U.S. Securities and Exchange Commission, 2023). The first question is whether filers moved at all. They did. Drawing a seeded sample of roughly one hundred fifty initial 13D filings from each of two windows (2023Q2–Q3, the last full pre-rule quarters, and 2024Q3–Q4, well after the effective date) and measuring the

trigger-to-filing delay in business days with the federal-holiday calendar, the distribution compresses sharply (Table 2 and Figure 2).

Table 2: The trigger-to-filing delay in business days, before and after the February 2024 acceleration. Seeded samples of initial 13D filings drawn from EDGAR, with the trigger date read from the filing itself and the delay counted in business days on the federal-holiday calendar. Of the 300 sampled filings (150 in each window), 198 carry a trigger date that parses: 102 pre and 96 post, which are the parse rates of 0.68 and 0.64 against the 150 drawn per window. An admissible band of 0 to 60 business days then screens out ten of those, 4 pre (delays of 129, 143, 526 and 567 business days) and 6 post (delays of 62, 73, 133, 262, 407 and 1,365 business days), and the statistics below are computed on the 188 that remain (98 pre, 90 post). Because the parse rates are below one, every count is a floor. Descriptive statistics; no standard errors are claimed.

Window	<i>n</i>	Mean	Median	p90	Share ≤ 5 bd	Share ≤ 10 bd
Pre (2023Q2–Q3, 10-calendar-day rule)	98	9.63	7.0	23.0	35.7%	80.6%
Post (2024Q3–Q4, 5-business-day rule)	90	6.40	5.0	11.1	75.6%	88.9%




Figure 2: The filing-delay distribution, pre (2023Q2–Q3) versus post (2024Q3–Q4) acceleration, in business days. The mass under the old ten-day regime piles between five and ten business days, with a long right tail of late filers; under the five-business-day rule the mass concentrates inside the new deadline. Descriptives from the seeded samples of Table 2.

Three readings of the table, each bounded. First, the compression is real and large: the median delay falls from seven business days to five, and the share of filings inside the new five-day deadline more than doubles, from 35.7 to 75.6 percent. Second, the pre-rule distribution was not bimodal. It was already spread across the old window in the way Bebchuk et al. (2013) document on filings by engaging investors: the mass rises into the deadline bucket, where 45.3 percent of their filings sit

at eight to ten business days, with 16.5 percent inside three business days and a tail of 7.1 percent beyond fifteen. The acceleration bit a population that was heterogeneous in exactly the dimension the rule targets; the design below exploits that heterogeneity as a dose. Third, the numbers agree in kind with the Commission’s own arithmetic: the release expected earlier filing for about 59 percent of timely 13D reports and put about 29 percent of 2022’s initial filings already inside the amended deadline, against our 35.7 percent already inside before the rule. Different samples and different windows, but the same direction (U.S. Securities and Exchange Commission, 2023). The honest caveats travel with the table: the samples are seeded draws, not the universe; the parse rates are floors; ten filings whose measured delay falls outside a nought-to-sixty business-day band are screened out and counted in the table note rather than cut silently; and these are descriptives, with no standard errors, because none are claimed. The delay is counted in business days throughout; counting calendar days against a business-day rule, as Polk et al. (2024) do, misclassifies the 48 percent of their sample that piles at their day-ten marker.

6.2 The written specification for the full empirics

This subsection summarises a design fixed before estimation; nothing in it has been run. Every estimate the December draft will report is specified here or in the underlying document, and the rule the design imposes on itself is that no specification may change after an estimate is seen, and that a null is reported with its minimum detectable effect rather than as “no effect”.

Sample and split. The universe is all initial Schedule 13D filings from 2022-01-01 to 2025-12-31, enumerated from EDGAR quarterly indexes, matched to CRSP. The treatment split is on the *trigger* date, not the filing date: the deadline attaches to the crossing, so a block crossed on 2024-01-30 and filed on 2024-02-08 belonged to the old regime. Filings that straddle are excluded and counted. The main pre window is 2022-01-01 through 2023-10-09. The adoption date, 2023-10-10, begins a separately reported adoption-to-effect stub through 2024-02-04; the stub is not pooled into “pre”. The structured-data mandate on 2024-12-18 is a second measurement break, so the main sample ends on 2024-12-17 and 2025 is reported as an extension with its own parse-rate table. The EDGAR dissemination cut-off moved from 5:30 pm to 10:00 pm ET on the same date the window changed, which shifts the flagged day by one for part of the post sample; the effective filing date FD^* is constructed to absorb this, and re-running with $FD^* \equiv FD$ is a required robustness row. A full re-fetch and re-parse with the current parser is a stated precondition for any estimate, and every count in the design is a floor until it has run.

(a) The timing split. For every filing, two abnormal returns from a market model estimated over $[TD - 252, TD - 22]$: the *run-up*, from the trigger to the day before the flag lands, and the *jump*, on the filing day (a three-day window around the effective filing date). This is the identity (10) and the flagged cell’s κ -invariance turned into a regression pair. The partition test (H1) is sign-free: the run-up’s liquidity slope should be non-zero, the jump’s zero, and the two different. It is estimated as one stacked regression with a flagged-half interaction, so the difference carries a standard error; inference is two-way clustered on subject firm and on calendar month, with a wild cluster bootstrap on the month dimension and the more conservative of the two reported. The February 2024 acceleration test (H2) asks whether the run-up’s liquidity slope changed after

February 2024, on a fixed five-trading-day window from the trigger so the outcome’s length does not change mechanically with the rule; both signs are written out in advance: attenuation (the shorter window moves mass to the κ -invariant flagged cell and flattens the composite slope) and amplification (the same informed trading concentrated into fewer days steepens it). One check can kill the identity: if the jump’s slope moves the same way as the run-up’s, there is no partition to write about. Liquidity is a pre-treatment Amihud-type illiquidity measure, inverted and standardised within quarter, with the orientation printed in every table note. The run-up *path* from trigger to filing, cut by liquidity tercile and period, is reported as a figure in the house style, descriptive only; the benchmark shape is a 0.9 percent abnormal return on the trigger date and 1.6 percent on the day after (Polk et al., 2024), with the run-up beginning on the trigger date itself (Zeng, 2026).

(b) The bindingness dose. The rule binds filers unevenly: those already filing inside five business days cannot be moved by it. The primary dose is a filer attribute built from pre-period behaviour only and is computed for filers with at least two pre-period initial 13Ds: the share whose delay exceeded five business days, in $[0, 1]$. For a single-filing filer, a leave-one-out pre-period stratum mean, by filer type, size tercile and illiquidity tercile, is imputed only in a robustness row and never in the primary estimate. The dose specification interacts the dose with Post and clusters on filer as well. Its first stage is mechanical (a fast filer cannot get faster) and is labelled a validity check, not a finding; the only working-paper estimate of the same first stage is Trivedi (2026)’s compliance share, +0.348 with standard error 0.130, whose own primary outcome, mean lag, is null, and whose sampling frame is never stated. Because slow filers are plausibly smaller and less liquid, the correlation between dose and liquidity is reported, and the dose specification is re-run within liquidity tercile.

(c) Stake at filing. The percent of the class printed on the 13D is the empirical counterpart of the model’s stake at filing B^F , an object the two-round model produces and the predecessor model could not. The specification asks whether the stake at filing fell after the acceleration and whether the fall is ordered by liquidity. One hazard is stated before the estimate because it points the other way: in the cross-section of delays, Bebchuk et al. (2013) find the *negative* gradient: filers who take longer end up with smaller stakes, 0.6 percentage points per ten days. The cross-section of delays is therefore never read as a window effect; only the 5 February 2024 filing-window change is. The historical days–stake gradient implies a mechanical stake change of roughly 0.12 percentage points from the two-day median-delay cut, against the stake leg’s minimum detectable effect of 0.85 percentage points on the post-re-parse projection, so a detectable estimate would have to be behavioural rather than arithmetic, and the paper would have to say which.

(d) The bounded null. The Commission’s own tables cap what the accumulation channel can do. Classifying 2011–2021 non-corporate-action initial 13Ds by how much of the reported stake was already accumulated by the amended five-business-day deadline: 80 percent of campaigns were unconstrained; 20 percent would have had some accumulation cut; 3 percent would have had at least a tenth of the stake cut; 1 percent at least a quarter (U.S. Securities and Exchange Commission, 2023). If only constrained campaigns can have their bid outcomes moved by the rule, the aggregate effect on the twelve-month bid hazard is capped by the constrained share, and the design reports

the whole ladder (20, 3 and 1 percentage points) rather than picking one. The 3-point rung is a judgement about what counts as a materially cut campaign, and it is stated as such. The bound covers the accumulation channel only; it is not the aggregate footprint of the rule, and the dollar figures that travel with it in the release are the Commission’s own, and the Commission disclaims them.

(e) The matched difference-in-differences on the bid hazard. The control-outcome leg asks whether a bid arrives within twelve months of the trigger, for 13D targets against *never-13D* firms matched three-to-one on size, illiquidity, industry and quarter, each control inheriting its treated firm’s trigger date as a pseudo-trigger so both are measured over the identical calendar window. Never-13D rather than 13G controls: 13G filers disclaim control intent, and selection between the two forms is itself plausibly responsive to the rule. Bid arrival is coded from tender-offer and merger filings, with a thirty-filing hand audit done blind. Base rates anchor the power arithmetic: 18.1 percent of campaign targets are acquired within twelve months against 7.2 percent of matched controls (Greenwood and Schor, 2009). The power arithmetic gives minimum detectable effects of 5.8 percentage points on today’s counts, 4.4 percentage points for S1 after the re-parse, and 9.9 percentage points for S2. The latter two use the optimistic projection $N \approx 4,950$, $w \approx 0.394$; they are projections, not estimates. The MDE exceeds the 3-point ceiling of the bounded null, so the difference-in-differences cannot separate the accumulation effect from zero even at that ceiling. Only an estimate’s excess above 3 percentage points can be attributed to incumbent defence or selection; the first 3 points remain compatible with accumulation.

(f) Bidder entry by liquidity. Within the 13D sample, the bidder-entry leg estimates the liquidity interaction

$$\text{BID12}_i = \alpha + \delta(\text{LIQ}_i \times \text{Post}_i) + \beta \text{LIQ}_i + \gamma \text{Post}_i + X_i' \theta + \text{FE} + \varepsilon_i.$$

Against the matched never-13D controls it estimates the matched triple difference

$$\text{BID12}_i = \alpha + \tau(\text{Treat}_i \times \text{Post}_i \times \text{LIQ}_i) + \text{all lower-order interactions} + X_i' \theta + \delta_{\text{match}} + \varepsilon_i.$$

The coefficient τ asks whether the effect of the February 2024 acceleration on bidder entry differs by liquidity. A rule-of-thumb power calculation gives projected MDEs of about 8.8 percentage points for S1 and 19.8 percentage points for S2. These are weakly powered projections, and the realised variance will replace the rule of thumb before an estimate is interpreted.

(g) The referee checklist. Control group or bounded null: both, as above. Confounds: ten named ones with a stated handling each, including the three the literature specifically warns about: the EDGAR cut-off moving on the same date as the window, calendar-day screens that stop being neutral under a business-day rule, and the incumbent-defence channel, which runs opposite to accumulation. Rights-plan coverage is untested: the proposed moderator uses best-effort EDGAR searches for 8-K Item 3.03 and 8-A12B; if coverage is insufficient, the defence channel is retained with its signed-bias fallback. Power: computed from counts on disk plus stated variance assumptions, with the arithmetic shown; the timing split detects 1.1–2.3 percentage points per standard

deviation of illiquidity after the re-parse, against a mean pre-rule run-up of 2.8 percent (Zeng, 2026), and anything smaller is reported as a bounded null. Placebos use all 568 holiday-adjusted dates before the 10 October 2023 adoption of the February 2024 acceleration, plus a pseudo-trigger placebo one quarter before the real crossing and a descriptive 13G placebo whose deadlines did not change until September 2024. Pre-trends: quarterly liquidity-slope estimates over the seven pre-adoption quarters, joint F -test, and a fixed blocking rule: a p -value below 0.10 removes causal language from the leg. Parser validation remains incomplete: the available pipeline output is pre-XML and the full-universe re-parse has not run; both limitations are reported.

6.3 Headline shell

The headline empirical result is a choice that has not been made, and this subsection is a shell: it fixes what the headline subsection must contain, lists the candidate headlines with what each would claim and what would kill each, and marks the choice pending rather than picking one.

[headline choice pending]

Candidate A: the partition test as headline. Claim: the run-up’s liquidity slope is non-zero, the jump’s is zero, and the difference is measured with a standard error: the paper’s identity in one regression pair, sign-free and needing no theory input. What it needs: the re-parsed sample and the stacked specification of (a). What kills it: a jump slope that moves with liquidity like the run-up’s, which would refute the partition itself and is reported first if it happens. Scope: it is an identity claim, not an estimate of the February 2024 acceleration, and a referee will ask what the 5 February 2024 filing-window change did.

Candidate B: the slope change after the February 2024 acceleration. Claim: the run-up’s liquidity slope changed after February 2024, in the direction the window record of Table 1 points (attenuation at every checked node, composition leg carrying it). What it needs: the same sample, the fixed-length run-up, and a decision on how much weight to put on a theory direction whose antecedent $A(\tau)$ is measured to fail at the implemented calibration, so the design keeps both branches live and pre-specifies what each sign means. What kills it: a bounded null (the minimum detectable effect is 1.1–2.3 percentage points on the post-re-parse projection), which is a reportable outcome rather than a failure. Scope: it is the candidate most directly about the anchor, and the one most exposed to the theory’s honest conditionality.

Candidate C: the bounded null as headline. Claim: the accumulation channel’s aggregate effect on the bid hazard is capped at about three percentage points by the Commission’s own tables, the matched design cannot separately detect even the ceiling, and that arithmetic, not a point estimate, is the paper’s answer to what the acceleration could do to control outcomes. What it needs: the matched sample and the ladder of rungs. What kills it: nothing in the data; the risk is presentation, since it reports a ceiling rather than an effect. Scope: it is the most defensible statement and the least newsy one, and it can be combined with either A or B as a supporting leg rather than standing alone.

The choice among A, B, and C is mine, and the December draft’s headline subsection will be written into this shell once it is made.

7 Conclusion

The disclosure rule is the market’s partition, and this paper takes the identity literally. A stake threshold and a filing window split the histories that reach a control decision into a flagged cell and a pooled one; the window is a primitive of the model rather than a rescaling of noise, and the stake at filing is an object the model produces. On that ground the paper proves what it can and records where it cannot. Existence is a conditional whose two main hypotheses about the equilibrium object are measured to fail at the implemented calibration, stated once, in plain words, because a reader deserves to know exactly what the proposition does and does not cover. The attenuation theorem separates the rule’s two margins honestly: the threshold margin attenuates with no dominance condition; the window margin attenuates if and only if a weight ratio times an unsigned composition ratio is at most one. The general-equilibrium survival of the fixed-policy sign is an implication with the region as a named hypothesis, plus numerical nodes. The numerical layer reads the same stack against the implemented calibration and reports an attenuation-direction window record alongside the measured failure of the chord restriction that would close the loop; the two facts sit in the same section because they are both true. The scope is stated as plainly as the results: the paper does not claim a global window-margin attenuation sign, κ -invariance of the filing-day jump, equilibrium uniqueness, a nonempty general-equilibrium region as a theorem, endogenous filing before the deadline, noisy or partially revealing flagged-round trading, continuous-time execution, or welfare or optimal rule design; nor that the predecessor model’s non-monotone result for minority gains survives in the two-round model, nor that the predecessor calibration at $\Omega \approx 0.037$ is economically meaningful.

The empirics rest on the February 2024 acceleration. The established descriptive fact is that filers moved: the median delay fell from seven business days to five, and the share inside the new deadline more than doubled. The specification for everything beyond that is written: a timing split that tests the partition directly, a bindingness dose, a stake-at-filing leg, a bounded null from the Commission’s own tables, a matched design on the bid hazard whose honest deliverable is already known to be a ceiling, and a bidder-entry-by-liquidity leg. The headline choice is recorded as pending.

Three open questions carry the research forward: whether any two-round pooled cell satisfies the chord restriction’s support condition, which is the largest single conditionality in the stack; whether the existence assumptions can be satisfied on a relevant menu, including the incentive-compatibility question for a fully separating plan; and whether the scoped continuity problem in the outer-map assumption can be repaired constructively. Alongside them stands one open calibration input that is nowhere load-bearing: the disclosed share of engagements ω_a , for which no dataset in the literature offers an anchor, and which the paper does not invent. Each question is about the domain of an assumption this paper states rather than hides, which is where a theory of a legal rule should leave its unfinished work.

References

Admati, Anat R. and Paul Pfleiderer (2009). “The “Wall Street Walk” and Shareholder Activism: Exit as a Form of Voice”. In: *Review of Financial Studies* 22.7, pp. 2645–2685.

- Albuquerque, Rui A., Vyacheslav Fos, and Enrique J. Schroth (2022). “Value Creation in Shareholder Activism”. In: *Journal of Financial Economics* 145.2, pp. 153–178.
- Back, Kerry, Pierre Collin-Dufresne, Vyacheslav Fos, Tao Li, and Alexander Ljungqvist (2018). “Activism, Strategic Trading, and Liquidity”. In: *Econometrica* 86.4, pp. 1431–1463.
- Bebchuk, Lucian A., Alon Brav, Robert J. Jackson Jr., and Wei Jiang (2013). “Pre-Disclosure Accumulations by Activist Investors: Analysis and Policy”. In: *Journal of Corporation Law* 39.1, pp. 29–52.
- Becht, Marco, Julian Franks, Jeremy Grant, and Hannes F. Wagner (2017). “Returns to Hedge Fund Activism: An International Study”. In: *Review of Financial Studies* 30.9, pp. 2933–2971.
- Bhide, Amar (1993). “The Hidden Costs of Stock Market Liquidity”. In: *Journal of Financial Economics* 34.1, pp. 31–51.
- Bond, Philip, Alex Edmans, and Itay Goldstein (2012). “The Real Effects of Financial Markets”. In: *Annual Review of Financial Economics* 4, pp. 339–360.
- Brav, Alon, Wei Jiang, Frank Partnoy, and Randall Thomas (2008). “Hedge Fund Activism, Corporate Governance, and Firm Performance”. In: *Journal of Finance* 63.4, pp. 1729–1775.
- Burkart, Mike, Denis Gromb, and Fausto Panunzi (1998). “Why Higher Takeover Premia Protect Minority Shareholders”. In: *Journal of Political Economy* 106.1, pp. 172–204.
- Burkart, Mike and Samuel Lee (2022). “Activism and Takeovers”. In: *Review of Financial Studies* 35.4, pp. 1868–1896.
- Celentano, Francesco and Oliver Levine (2025). “Shareholder Activism, Takeovers, and Managerial Discipline”. Swiss Finance Institute Research Paper No. 25-81; SSRN 5506659.
- Cetemen, Doruk, Gonzalo Cisternas, Aaron Kolb, and S. Viswanathan (2026). “Leader-Follower Dynamics in Shareholder Activism”. In: *Journal of Finance*. Forthcoming.
- Coffee, John C. (1991). “Liquidity Versus Control: The Institutional Investor as Corporate Monitor”. In: *Columbia Law Review* 91.6, pp. 1277–1368.
- Collin-Dufresne, Pierre and Vyacheslav Fos (2015). “Do Prices Reveal the Presence of Informed Trading?”. In: *Journal of Finance* 70.4, pp. 1555–1582.
- (2016). “Insider Trading, Stochastic Liquidity, and Equilibrium Prices”. In: *Econometrica* 84.4, pp. 1441–1475.
- Corum, Adrian Aycan (2025). “The Stick or the Carrot? The Role of Regulation and Liquidity in Activist Short-Termism”. Working paper, SSRN 4319599.
- Corum, Adrian Aycan and Doron Levit (2019). “Corporate Control Activism”. In: *Journal of Financial Economics* 133.1, pp. 1–17.
- Edmans, Alex, Vivian W. Fang, and Emanuel Zur (2013). “The Effect of Liquidity on Governance”. In: *Review of Financial Studies* 26.6, pp. 1443–1482.
- Edmans, Alex, Itay Goldstein, and Wei Jiang (2012). “The Real Effects of Financial Markets: The Impact of Prices on Takeovers”. In: *Journal of Finance* 67.3, pp. 933–971.
- (2015). “Feedback Effects, Asymmetric Trading, and the Limits to Arbitrage”. In: *American Economic Review* 105.12, pp. 3766–3797.
- Faure-Grimaud, Antoine and Denis Gromb (2004). “Public Trading and Private Incentives”. In: *Review of Financial Studies* 17.4, pp. 985–1014.
- Gantchev, Nickolay (2013). “The Costs of Shareholder Activism: Evidence from a Sequential Decision Model”. In: *Journal of Financial Economics* 107.3, pp. 610–631.

- Glosten, Lawrence R. and Paul R. Milgrom (1985). “Bid, Ask and Transaction Prices in a Specialist Market with Heterogeneously Informed Traders”. In: *Journal of Financial Economics* 14.1, pp. 71–100.
- Goldstein, Itay (2023). “Information in Financial Markets and Its Real Effects”. In: *Review of Finance* 27.1, pp. 1–32.
- Greenwood, Robin and Michael Schor (2009). “Investor Activism and Takeovers”. In: *Journal of Financial Economics* 92.3, pp. 362–375.
- Grossman, Sanford J. and Oliver D. Hart (1980). “Takeover Bids, the Free-Rider Problem, and the Theory of the Corporation”. In: *Bell Journal of Economics* 11.1, pp. 42–64.
- Hirschman, Albert O. (1970). *Exit, Voice, and Loyalty: Responses to Decline in Firms, Organizations, and States*. Cambridge, MA: Harvard University Press.
- Johnson, Travis L. and Nathan Swem (2021). “Reputation and Investor Activism: A Structural Approach”. In: *Journal of Financial Economics* 139.1, pp. 29–56.
- Kahn, Charles and Andrew Winton (1998). “Ownership Structure, Speculation, and Shareholder Intervention”. In: *Journal of Finance* 53.1, pp. 99–129.
- Kyle, Albert S. (1985). “Continuous Auctions and Insider Trading”. In: *Econometrica* 53.6, pp. 1315–1335.
- Kyle, Albert S. and Jean-Luc Vila (1991). “Noise Trading and Takeovers”. In: *The RAND Journal of Economics* 22.1, pp. 54–71.
- Maug, Ernst (1998). “Large Shareholders as Monitors: Is There a Trade-Off between Liquidity and Control?” In: *Journal of Finance* 53.1, pp. 65–98.
- Norli, Øyvind, Charlotte Østergaard, and Ibolya Schindele (2015). “Liquidity and Shareholder Activism”. In: *Review of Financial Studies* 28.2, pp. 486–520.
- Ordóñez-Calafí, Guillem and Dan Bernhardt (2022). “Blockholder Disclosure Thresholds and Hedge Fund Activism”. In: *Journal of Financial and Quantitative Analysis* 57.7, pp. 2834–2859.
- Polk, Mark, Lewis Buchheit, Jessica Riley, and Gavin Stone (2024). “Shrinking the 13D Disclosure Window Will Benefit Non-Activist Investors”. In: *Journal of Financial Regulation and Compliance* 32.4, pp. 516–538.
- Trivedi, Avaneendra (2026). “The Mandated Revelation Field and a Conserved-State Measure of Scheduled Disclosure in U.S. Equity Markets”. Working paper, SSRN 6866499.
- U.S. Securities and Exchange Commission (2022). *Beneficial Ownership Reporting on Schedules 13D and 13G*. Release No. 33-11030; File No. S7-06-22 (proposing release, February 10, 2022).
- (2023). *Beneficial Ownership Reporting on Schedules 13D and 13G*. Release No. 33-11253; File No. S7-06-22 (adopting release, October 10, 2023), 88 Fed. Reg. 76896.
- Zeng, Jean (Jieyin) (2026). “Do Managers Learn about Their Firm’s Ownership Changes before Public Disclosure?” In: *Review of Accounting Studies* 31, pp. 1301–1341.