

Disclosure, Liquidity, and Takeover Premia

Repositioning Framework v3.1: Angle, Simplified Model, and Empirical Design (revised after referee round)

Junyu Li

2026-08-19

Abstract

This memo specifies the repositioning of the paper formerly titled “Liquidity, Activism Disclosure, and Takeover Premia” (draft_v2) for submission to a top finance journal. The new framing: the Schedule 13D disclosure rule is a policy lever on the *revelatory* efficiency of prices in the market for corporate control — it partitions market inference about activism into a disclosed and an inferred branch, and the takeover premium is priced off both. The theory core is simplified to a three-branch model with an explicit disclosure primitive, a headline comparative static (disclosure attenuation) that is stated exactly as far as it is proved, an honest endpoint-plus-region treatment of the liquidity hump, and an empirical section built on the SEC’s February 5, 2024 acceleration of the 13D filing window. Target: JF/RFS; alternative JFE. **v3.1 (2026-08-19)**: revised after an internal referee round; see the change log in Section 7 and the referee report at `quality_reports/reports/2026-08-19_framework_v3_referee_report.md`. Honesty labels from v3 are preserved or strengthened, never weakened.

Table of contents

1	The Angle	2
1.1	The repositioning in one paragraph	2
1.2	Why this angle wins	2
1.3	Framing device	3
2	The Simplified Model	4
2.1	Notation (new in v3.1)	4
2.2	Players, timing, and information	4
2.3	Disclosure strictness as a model primitive — two margins	5
2.4	Equilibrium	6
3	Main Results and Proof Plan	6
3.1	Proved objects already in hand	6
3.2	T2 — Disclosure attenuation (the headline comparative static), stated as far as it is proved	6
3.3	T4 — The liquidity hump, honestly stated	8
3.4	T5 — The premium wedge (appendix microfoundation)	8
3.5	No-manipulation discipline	8

4 Empirical Design	9
4.1 What the referee round found in the disk data (must be absorbed, not argued with)	9
4.2 Facts section (facts the model must match)	9
4.3 H1 — Attenuation of the liquidity slope (a before/after test, with a dose)	10
4.4 H2 — Run-up to jump shift (the Collin-Dufresne–Fos template), in levels	11
4.5 H3 — The hump as a falsifiable cross-section (a prediction the theory offers; not promised)	12
4.6 Calibration moments (split into what the model can produce)	12
5 Edit Map: draft_v2 → draft_v3	13
6 Journal Targets and Risks	15
7 Change log v3 → v3.1 (2026-08-19)	16

1 The Angle

1.1 The repositioning in one paragraph

Move from “a static microstructure model of the exit-voice margin with numerically verified comparative statics” to “**the 13D disclosure rule as a policy lever on revelatory price efficiency in the market for corporate control.**” The disclosure rule — its threshold and its filing window — splits market inference about activism into a *disclosed* branch (engagement is common knowledge) and an *inferred* branch (engagement must be read from order flow), and a potential bidder prices the takeover off both. Pre-disclosure accumulation under a *fixed* disclosure horizon is well understood (Kyle and Vila 1991; Back, Collin-Dufresne, Fos, Li and Ljungqvist 2018, *Econometrica*; Cetemen, Cisternas, Kolb and Viswanathan 2026, *JF*), as is post-disclosure intervention choice *given* a toehold (Burkart and Lee 2022, *RFS*). What no paper does — theoretically or structurally — is make the disclosure *rule* the state variable that determines how liquidity maps into the takeover premium: Ordóñez-Calafi and Bernhardt (2022, *JFQA*) design the threshold without a bidder; Cetemen et al. stop at the crossing; Albuquerque, Fos and Schroth (2022), Johnson and Swem (2021) and Celentano and Levine (2025) estimate activism outcomes with no order-flow inference; Corum (2025) has liquidity and regulation but no premium. The empirical literature links liquidity to premia in reduced form (Massa and Xu 2013, *JFQA*; Huang, Maharjan and Nanda 2024, *JCF*) and the theory literature links liquidity to the *mode* of intervention (Maug 1998; Mello and Repullo 2004, “Shareholder Activism is Non-Monotonic in Market Liquidity”). We are the first to make the rule the partition and to price the premium off it.

(v3.1: the v3 sentences “the first formalization of how the market learns about activism-driven control events” and “Nobody ... links market liquidity or the disclosure rule to takeover premia” were overclaims — see the referee report, §2 — and are replaced by the paragraph above.)

1.2 Why this angle wins

1. **The gap is quotable, verbatim.** Back et al. (2018, p. 1454): “we assume that the horizon at which the activist’s stake is disclosed is fixed and common knowledge ... it might be interesting to endogenize the horizon” (the 5% threshold is the institutional context). Burkart and Lee (2022, p. 1891): “we do not endogenize the acquisition of the toehold in

anonymous, predisclosure markets ... We leave an analysis of how predisclosure and post-disclosure decisions interact to future work.” Note the sentence in between — the toehold problem “has been comprehensively studied by Back et al. (2018)” — which is why our claim is the *partition and the rule*, not market learning per se.

2. **Frontier taste.** The closest recent top-3 theory papers (Burkart–Lee RFS 2022; Levit–Malenko–Maug JF 2024 and JF 2026; Kakhbod–Loginova–Malenko–Malenko RFS 2023; Burkart–Lee–Voss, ECGI WP 956/2024, 2025 ECGI prize; Corum–Levit JFE 2019) are analytical and one-mechanism; most, not all, close with an empirical-implications section (Burkart–Lee ends on concluding remarks). Levit–Malenko–Maug (2024) price shares off the marginal shareholder with Walrasian clearing — a published precedent that a non-Kyle trading technology is acceptable, not a “proof” of it. Brouwer existence plus honestly labeled numerical uniqueness is acceptable *only if the headline comparative statics are proved* — which Section 3 treats as the binding constraint, not a formality.
3. **The policy hook is live — and the first-stage fact is already public.** The SEC’s Modernization of Beneficial Ownership Reporting (Rel. 33-11253; adopted October 10, 2023; effective **February 5, 2024**) shortened the initial 13D window from 10 calendar days to 5 business days. Trivedi (SSRN 6866499, June 2026) already runs a pre-registered difference-in-differences on this change (13G filers as control): +0.35 in the share filed within five business days, with nulls on mean lag, bid–ask spread and an illiquidity proxy. Polk, Buchheit, Riley and Stone (“Shrinking the 13D disclosure window will benefit non-activist investors”, *J. Financial Regulation and Compliance* 32(4):516–538, 2024) measure delay-period abnormal returns on pre-2024 data and project the effect of the shorter window. Corum (2025, SSRN 4319599, “The Stick or the Carrot?”) theorizes liquidity, regulation and activist short-termism (no premium). Bishop, Fos, Jiang and Partnoy (2026, SSRN 6061814) show activists avoid targets where a toehold would trigger HSR disclosure before the 13D. None has a takeover-premium object or the deadline as a lever on inference; the triple *deadline × liquidity × equilibrium premium* is open. Treat the window as ~12 months, not “open but closing”.
4. **Differentiation from the most dangerous competitor.** Celentano and Levine (2025, R&R at *RFS* per the authors’ CV, dated Oct. 2025) estimate activism’s equilibrium effect on M&A with **no order-flow inference and no price impact** — liquidity enters only as a reduced-form entry cost — and a stake fixed at the 5% statutory threshold. One sentence on page 1: *they take prices as given; we make liquidity and the disclosure rule the state variables, and the takeover premium an equilibrium object of order-flow inference and tender free-riding.*

1.3 Framing device

Following Bond, Edmans and Goldstein (2012, *ARFE*) and Goldstein (2023, *Review of Finance*), the paper studies **revelatory** price efficiency (RPE) — what the price reveals to a specific decision-maker, the bidder — rather than forecasting efficiency. A disclosure *rule* is a policy lever on RPE. Goldstein (2023, §5.2) lists RPE-oriented measurement as an open problem; the feedback–corporate-control interaction is Edmans, Goldstein and Jiang (2012, *JF*), whose discrete order-flow technology (EGJ 2015, *AER*) is the methodological precedent — with the caveat that EGJ prove uniqueness *analytically over parameter regions* (AER Prop. 1); we adopt their region-wise posture, not a “numerical uniqueness following EGJ” (draft_v2 l. 702, 2235 must be rewritten).

2 The Simplified Model

2.1 Notation (new in v3.1)

v firm value; $s = v + \varepsilon$ the blockholder's signal with $\sigma_s^2 = \sigma_v^2 + \sigma_\varepsilon^2$; (q, a) order and engagement; z noise order; $X = q + z$; D disclosure flag; κ noise intensity; $C(s)$ engagement cost, strictly positive and decreasing (A2); Δ_{eng} the value improvement from engagement, ρ its success probability, $\tilde{\Delta} = \rho \Delta_{\text{eng}}$; realized bid premium $m^R(a)$ with $m^R(0) = m_0$, $m^R(1) = m_1$, and $\tilde{m} = \mathbb{E}[m^R(1)] \leq m_1$; $\bar{m}(X, D) = m_0 + \pi(X, D)(\tilde{m} - m_0)$ the expected premium; $\pi(X, D) = \Pr(a = 1 \mid X, D)$; $\omega_E, \omega_H, \omega_Q, \omega_P$ ex-ante region masses; $h(\pi) = \pi p(\pi)$; $\bar{p}_1$ the disclosed-branch bid probability; $\bar{S}, K, \sigma_\xi$ the bidder's synergy scale, entry cost and synergy dispersion; $\Delta^{\min} = \mathbb{E}[m^R(a) \mathbf{1}\{\text{bid}\}]$ expected minority takeover gains, $\Delta^{\text{act}} = (\tilde{m} - m_0) \Pr(a = 1, \text{bid})$ its activism component; D7 tender game: r fringe-raid probability, γ portability, ϕ dilution (φ is the normal density in D7), ψ pivotality factor, λ appropriability. (v3 mislabeled the dilution and pivotality symbols and overloaded q ; fixed.)

2.2 Players, timing, and information

All players are risk neutral. There is one blockholder (B), one noise trader, one competitive market maker (M), and one potential bidder.

- $t = 0$. Nature draws firm value $v \sim \mathcal{N}(\mu_v, \sigma_v^2)$. The blockholder privately observes $s = v + \varepsilon$, $\varepsilon \sim \mathcal{N}(0, \sigma_\varepsilon^2)$, and holds an initial stake normalized to one share.
- $t = 1$ (**trade and disclose**). The blockholder chooses an action (q, a) from the three-branch menu

$$(q, a) \in \left\{ \underbrace{(-1, 0)}_{\text{Exit}}, \underbrace{(0, 1)}_{\text{Quiet}}, \underbrace{(+1, 1)}_{\text{Public}} \right\},$$

where q is the order and $a \in \{0, 1\}$ engagement. The four-branch menu of draft_v2 is pruned under (**A2'**): the engagement net benefit dominates the cost on the whole non-exit region, $A(Q) - A(H) \geq C(s)$ for all $s \geq k_1$ (equivalently $C(k_1) \leq A(Q) - A(H)$ given (A2)), so Quiet weakly dominates Hold and the menu is $\{E, Q, P\}$. *(v3 said “Hold folds into the trading decision; it collapses at baseline anyway”: the collapse is a baseline numerical fact — draft_v2

1. 1010; the text after Prop. cutoffs says the region “may collapse” — not a theorem. The prune sets $\bar{\pi} = \omega_Q / (\omega_H + \omega_Q) \equiv 1$, which makes the chord condition (C*) a genuinely primitive inequality, $h(1) < 2h(1/2)$; the baseline numerics must be re-solved on the three-branch model before any D8 number is re-quoted.)* Noise demand $z \in \{-1, 0, +1\}$ with

$$\Pr(z = 0) = 1 - \kappa, \quad \Pr(z = +1) = \Pr(z = -1) = \frac{\kappa}{2},$$

so $\kappa \in [0, 1]$ is the liquidity/noise intensity. The market maker observes total order flow $X = q + z \in \{-2, \dots, +2\}$ and the disclosure flag D (defined in §2.3), and sets

$$P(X, D) = \delta \mathbb{E}[Y \mid X, D],$$

with δ a normalization. (v3.1: draft_v2 uses $\delta = 0.95$ in §5 and *numerical/params.py*, and $\delta = 1$ in the transfer-netting lemma; pick one and, if $\delta = 1$, redo the (A5a) margin arithmetic.)

- $t = 1.5$ (**bidder entry**). The bidder observes (P, D) and a private synergy shock $\xi \sim \mathcal{N}(0, \sigma_\xi^2)$, and enters iff

$$\Pi_B = \bar{S} - P + \xi - \bar{m}(P, D) - K \geq 0, \quad \bar{m}(P, D) \equiv \mathbb{E}[m^R \mid P, D],$$

so the entry probability is

$$p(P, D) = \Phi\left(\frac{\bar{S} - P - \bar{m}(P, D) - K}{\sigma_\xi}\right), \quad \frac{\partial p}{\partial P} < 0 \quad (\text{bid monotonicity, proved at fixed } \bar{m}).$$

Because (P, D) is a sufficient statistic for the payoff-relevant pair $(\hat{V}, \bar{m})$ — the three disclosed cells share one price *and* one posterior, so their pooling is payoff-irrelevant, and the $D = 0$ cells carry generically distinct prices — conditioning on (P, D) coincides with conditioning on (X, D) . (*This replaces draft_v2’s Lemma dropA7, whose “injectivity at any nondegenerate calibration” is contradicted by the paper’s own Appendix B: $P^*(x, 1)$ is identical for $x \in \{0, 1, 2\}$.*)

- $t = 2$ (**payoffs**). If a bid occurs, $Y = P + m^R(a)$; otherwise $Y = v + a\tilde{\Delta}$. The premium wedge $m_1 > m_0$ is microfounded in the appendix tender game (D7): with fringe-raid probability r , portability γ , dilution ϕ and pivotality factor ψ ,

$$\lambda = 1 - r(1 - \gamma)\psi \in [0, 1] \quad (\text{proved, closed form}).$$

2.3 Disclosure strictness as a model primitive — two margins

In draft_v2 the flag is $D = \mathbf{1}\{q = +1\}$ and $\omega_P = 1 - \Phi(\alpha_D)$ is an *equilibrium mass* pinned by the cutoff k_D ; “ $\omega_P(\tau)$ with $\omega'_P(\tau) > 0$ ” (v3) was therefore a relabeling, not a primitive. v3.1 makes strictness explicit and separates the two institutional margins, because they enter the model differently and — see §3.2 — carry different signs. (Symbol hygiene: draft_v2 already uses τ for the discount horizon, $\delta = e^{-r\tau}$, l. 335; rename that one.)

- **Threshold margin** $\tau_\theta \in [0, 1]$ (5% vs 3%; who must file). A Quiet engagement is publicly disclosed with probability τ_θ : $D = \mathbf{1}\{q = +1\} \vee \mathbf{1}\{a = 1, \zeta \leq \tau_\theta\}$, $\zeta \sim U[0, 1]$ independent. Then the disclosed mass is $\omega_P + \tau_\theta \omega_Q$ (so $\partial/\partial\tau_\theta > 0$ genuinely), the $D = 0$ masses become $(\omega_E, \omega_H, (1 - \tau_\theta)\omega_Q)$, and the closed-form $D = 0$ posteriors of Prop. posteriors go through with $\omega_Q \mapsto (1 - \tau_\theta)\omega_Q$. Disclosed-branch κ -invariance survives: in both disclosed sub-blocks ($X = 1 + z$ and $X = z$), conditioning on X reveals only the noise draw.
- **Window/timing margin** $\tau_w \in [0, 1]$ (10 calendar days vs 5 business days; *when* the market learns). The Public buy is flagged before the market maker prices it with probability τ_w ; with probability $1 - \tau_w$ it is priced from order flow alone. $\tau_w = 1$ is draft_v2; $\tau_w = 0$ is draft_v2’s no-disclosure benchmark (§8.2 there). Back et al. (2018, pp. 1452–1453): “what matters is $\sigma^2 T$... reducing the trading horizon T is isomorphic to reducing noise trading volatility” — i.e., a shorter window also maps into **lower** κ , not only into higher τ_w . The February 2024 shock is therefore a *joint* movement ($d\kappa < 0, d\tau_w > 0$) and must be presented as one; the clean single-margin experiment is the cross-country threshold contrast (τ_θ).

Both margins move mass from the inferred to the disclosed branch. Only the threshold margin does so by removing *engaged* mass from the pool the market must read; the window margin removes the *buy* order, which is intrinsically informative ($X = +2$ is reachable only from $q = +1$) even without a flag. This is why their comparative statics differ (§3.2).

2.4 Equilibrium

A perfect Bayesian equilibrium consists of (i) a weakly ordered cutoff strategy $k_1 \leq k_D$ over the signal s (with $k_0 = k_1$ under (A2')); (ii) Bayes-consistent posteriors $\pi(X, D)$; (iii) the bidder's entry rule; (iv) the competitive price fixed point $P = \delta \mathbb{E}[Y \mid X, D]$. Existence follows from Brouwer's theorem on the cutoff polytope Θ under (A1), (A2), (A4) **and (A5a)** — the maintained single-valuedness of the inner price fixed point on each cell, without which the best-response map is a correspondence. (A5a) is an assumption, not a theorem: v3 called existence “unconditional”, which it is not; draft_v2's own Remark A5margins (l. 696) says the baseline *fails* the conservative sufficient bound ($0.805 + 0.947 > 1$), while §5 (l. 1010) says it *satisfies* the (A5) sufficient condition ($\delta/\sigma_\xi = 2.375 < 2.507$) — the two statements must be reconciled before draft_v3 says anything about (A5). Uniqueness is stated as a labeled numerical regularity (multistart convergence), never as a theorem — the EGJ (2015) posture of region-wise characterization rather than global uniqueness claims.

3 Main Results and Proof Plan

3.1 Proved objects already in hand

The posterior block is closed form: the disclosed branch is degenerate, $\pi(X, 1) = 1$, and κ -invariant (proved, modulo (A5a) for the inner fixed point); the inferred branch satisfies $\pi(-2, 0) = 0$ with $\pi(X, 0)$ rational in (ω, κ) (proved). Bid monotonicity, the dominance lemma pruning the action menu, and existence (as qualified above) are proved. Endpoint symmetry

$$\Delta^{\min}(0^+) = \Delta^{\min}(1^-)$$

is proved **at fixed cutoffs** (noise symmetry plus disjoint $D = 0$ supports at $\kappa \rightarrow 1$; the proof does not use Hold, so it survives the prune); the step that lifts it to full equilibrium (“the equilibrium indifference system is identical at the two extremes”, draft_v2 l. 815–817) is asserted and numerically checked (gap 6.9×10^{-7}), and must be labeled as such or proved.

3.2 T2 — Disclosure attenuation (the headline comparative static), stated as far as it is proved

At fixed cutoffs, the activism component decomposes across the two branches as

$$\Delta^{\text{act}}(\kappa; \tau) = (\tilde{m} - m_0) \left[\omega_P^{\text{disc}}(\tau) \bar{p}_1 + (1 - \omega_P^{\text{disc}}(\tau)) \mathbb{E}_{D=0}[h(\kappa; \tau)] \right],$$

where $\bar{p}_1$ is the disclosed-branch bid incidence (κ -invariant, proved) and the inferred-branch term carries all inference sensitivity. Hence

$$\left| \frac{\partial \Delta^{\text{act}}}{\partial \kappa} \right| = (\tilde{m} - m_0) (1 - \omega_P^{\text{disc}}(\tau)) \left| \frac{\partial}{\partial \kappa} \mathbb{E}_{D=0}[h(\kappa; \tau)] \right|,$$

and the τ -derivative of the right-hand side is a **two-term** product rule: a weight term $-\partial_\tau \omega_P^{\text{disc}} \cdot |\partial_\kappa \mathbb{E}_{D=0} h| < 0$ (this is what v3 called “one line”) **and** a composition term $(1 - \omega_P^{\text{disc}}) \partial_\tau |\partial_\kappa \mathbb{E}_{D=0} h|$, because the $D = 0$ posteriors are homogeneous of degree zero in $(\omega_E, \omega_H, \omega_Q)$ and therefore invariant to τ *only* under a proportional rescaling of all three non-disclosed masses — which no disclosure rule does. The status of T2 is therefore:

1. **Threshold margin τ_θ — attenuation holds, a fortiori, pending one lemma.** Moving Quiet engagement into the disclosed branch lowers $\bar{\pi}(\tau_\theta) = (1-\tau_\theta)\omega_Q/(\omega_H+(1-\tau_\theta)\omega_Q)$, and by the chord identity of Lemma d1-jensen the interior κ -motion of $\mathbb{E}[h]$ is the probability-weighted chord gap of h on $[0, \bar{\pi}]$, which vanishes as $\bar{\pi} \downarrow 0$; numerically at baseline masses, $|\partial_\kappa \mathbb{E}_{D=0} h|$ falls from 0.0153 ($\tau_\theta = 0$) to 0.00014 ($\tau_\theta = 0.95$). Both terms are negative. **To prove:** $\partial|\partial_\kappa \mathbb{E}_{D=0} h|/\partial\bar{\pi} \geq 0$ (two paragraphs from Lemma d1-jensen). Consistent check with the repo model: lowering k_D at fixed (k_1, k_0) lowers the total variation of Δ^{act} over $\kappa \in [0.15, 0.85]$ monotonically ($0.0176 \rightarrow 0.0040$ as ω_P : $0.037 \rightarrow 0.50$; independently reproduced) — but most of that fall occurs *without any flag* (no-disclosure regime: $0.0165 \rightarrow 0.0106$), because the buy is intrinsically revealing. Once proved, T2 on this margin is close to an accounting identity — sell it as the policy statement it is (a rule is a lever on revelatory efficiency), not as a discovered mechanism.
2. **Window margin τ_w — attenuation is *not* a partial-equilibrium theorem.** Comparing the flagged regime ($\tau_w = 1$) with the pooled regime ($\tau_w = 0$) at fixed cutoffs with the repo model, the κ -sensitivity of Δ^{act} (total variation over $\kappa \in [0.15, 0.85]$) is **higher** under disclosure for disclosed masses up to ≈ 0.29 (ratio disc/no-disc = 1.06 at the baseline $\omega_P = 0.037$; 1.19 at 0.13; 1.14 at 0.29) and lower only at $\omega_P = 0.50$ (ratio 0.38); pointwise, disclosure is steeper at $\kappa \geq 0.7$ for $\omega_P \leq 0.29$. The sign is calibration-dependent. This is the margin the February 2024 shock moves, so H1 (§4) cannot be called “the flagship test of T2” until the sign of the *disclosure-jump* liquidity slope (item 4) is derived under $(d\kappa < 0, d\tau_w > 0)$.
3. **Magnitude at the paper’s own baseline is negligible.** With $k_D = 2.26$, $\mu = 1$, $\sigma_s = 0.707$, the disclosed mass is $\omega_P \approx 0.037$; the two curves of draft_v2’s Figure “disclosure” (`disclosure_attenuation.csv`) have ranges 0.01107 vs 0.01117 over $\kappa \in [0.15, 0.85]$ — a difference below 1%, and by mean $|\text{slope}|$ the *disclosed* regime is slightly more sensitive (0.0251 vs 0.0236). The calibration must be re-anchored so that the disclosed share of engagement is empirically meaningful (13D-disclosed vs behind-the-scenes engagement), and the paper must report the attenuation *elasticity*, not only its sign.
4. **The model object that H1 tests must be defined.** T2 concerns the ex-ante expected activism premium; the 13D announcement return is a different object. Define the *disclosure jump* $J(X) = P(X, 1) - P_{ND}(X)$: the price with the flag minus the pooled price at the same order flow (the no-disclosure price is the pre-flag price). J is well defined; at baseline $\mathbb{E}[J \mid D = 1]$ increases in κ (0.33 at $\kappa = 0.2$, 0.39 at 0.5, 0.42 at 0.8) and $J(2) = 0$ (the buy is already revealed at $X = +2$). Its liquidity slope and how τ_w and κ jointly move it are the derivation H1 rests on. Same object gives the run-up share of H2 $([P_{ND}(X) - \mu]/[P(X, 1) - \mu])$.
5. **T3 — general-equilibrium region version: a to-do, not a result.** No GE attenuation derivation exists in the repo (v3 wrote it in the indicative). D8’s inversion-free bound controls $dk/d\kappa$ only; a cross-partial would need $d^2k/d\kappa d\tau$, which no contraction modulus reaches, and **prop:d8-counter** warns that GE can overturn a PE sign here. The statement to aim for compares two *first* derivatives: for $\tau' > \tau$, if the fixed-cutoff slope gap exceeds $\bar{B}(\kappa; \tau) + \bar{B}(\kappa; \tau')$ on $[\kappa_{lo}, \kappa_{hi}]$ (with $\bar{B}$ the D8 bound at each τ), then $|d\Delta^{\text{act}}/d\kappa|(\kappa; \tau') < |d\Delta^{\text{act}}/d\kappa|(\kappa; \tau)$ there — a triangle inequality that reuses D8 twice. Requires τ in **numerical/** first: a few days.

3.3 T4 — The liquidity hump, honestly stated

The hump is **not** the headline. It is presented as three objects, in descending order of strength (and it must be distinguished from Mello and Repullo 2004, whose non-monotonicity is in intervention incentives, ours in the premium through the inference partition):

1. **Endpoint symmetry (proved at fixed cutoffs; equilibrium limit asserted + numerically checked):** $\Delta^{\min}(0^+) = \Delta^{\min}(1^-)$.
2. **Certified region (computed, machine-checked):** strict single-peakedness of $\Delta^{\min}(\kappa)$ is certified on $\kappa \in [0.35, 0.825]$ (inversion-free bound; $[0.30, 0.85]$ under exact IFT) with contraction modulus $L \leq 0.836$, under the checkable conditions (R1) pointwise dominance and (R2) integral control (D8). Channel (A) is single-peaked at fixed cutoffs if the chord condition

$$\mathcal{C}(\bar{\pi}) = h(0) - 2h(\frac{\bar{\pi}}{2}) + h(\bar{\pi}) < 0, \quad h = \pi \cdot p,$$

holds (proved; $\mathcal{C} > 0$ gives a trough; $\mathcal{C} = 0$ is not covered). **Honesty note carried over from draft_v2 (Remark d5-vacuous), which v3 omitted:** on the paper’s calibration family (C*) cannot fail, so it selects nothing there — the hump/trough orientation is produced entirely by the GE cutoff-shift channel, and the chord second difference is retained only as a 16/20-cell diagnostic. Under the three-branch prune $\bar{\pi} \equiv 1$ and (C*) becomes the primitive inequality $h(1) < 2h(1/2)$; whether it can fail on the re-solved family is open.

3. **Falsifiable prediction (numerical elsewhere):** the interior peak $\kappa^\dagger \approx 0.58$ (full-equilibrium argmax in `baseline_series.csv`, 0.0206 grid; 0.60 on D8’s coarser 0.025 grid — a grid-resolution spread; v3’s 0.59 was a solver-check value between grid points) and the certified trough at $\sigma_\xi = 0.60$ are featured as an *economic boundary*: the hump lives where bidder entry is premium-sensitive (Remark d8-boundary). Non-monotonicity as an identification device is an *analogy* to Corum–Levit (2019, Prop. 4 — non-monotonicity in board private benefits separates selection from treatment), not the same device.

3.4 T5 — The premium wedge (appendix microfoundation)

The disagreement-node tender game delivers $\lambda = 1 - r(1 - \gamma)\psi$ in closed form (MC-verified), and the boundary of the wedge assumption is characterized: $m_1 > m_0$ holds iff not (certain raids $\wedge \gamma = 0 \wedge$ unblockable offers), given $\theta < 1$ and $\rho\Delta_{\text{eng}} > 0$ — all proved. The measured-premium reversal $M(\pi) = \bar{m}/P$ strictly decreasing in π (fixed cutoffs) holds for $\lambda < \lambda_{\text{crit}} \approx 0.07$ (D7 JSON), whereas the baseline calibration sits at $\lambda \approx 0.86$; D7 aims the result at Albuquerque–Fos–Schroth’s cross-sectional predictions. Whether it also accounts for Celentano–Levine’s -13.7% (-5.2pp) is a separate mapping question — their estimand is a selection-corrected treatment effect, $M(\pi)$ a cross-sectional ratio — so v3’s “rationalizes” is withdrawn: either move the calibration below λ_{crit} or state the reversal as a low-appropriability regime.

3.5 No-manipulation discipline

EGJ (2015, *AER*) Case 2 — a corrective action makes value non-monotone in the state, so a positively informed trader may sell to induce it; their sufficient condition against manipulation is $R_H - R_L > \frac{4}{3}x$ (verified in the published text). Here the profile is *infeasible* rather than unprofitable: Exit liquidates the entire unit stake (Lemma dominance (i)), and Public Voice buys, so no action shorts the firm while retaining exposure to the bidder’s decision. One sentence in `draft_v3`, not a lemma; it becomes substantive only if partial exit is added.

4 Empirical Design

Discipline. One dated shock, one signed prediction per test; institutional facts as anchoring, never as “tests” without exogenous variation; announcement CARs are *disclosure price impact*, never labeled value creation (preempting Ben-David, Bhattacharya, Huang and Jacobsen, JF 2026: “CAR is an unreliable measure of expected value creation”). **v3.1 addition:** every US 13D filer is treated on the same date, so H1 is a *before/after* comparison with rich controls; the word “difference-in-differences” is used only where a second difference exists (the bindingness dose in H1; the within-event window contrast in H2).

4.1 What the referee round found in the disk data (must be absorbed, not argued with)

- **No break in the outcome at the rule date.** Half-year median 13D CAR (WRDS market-model output, 2,285 events): 2022H1 4.4%, 2022H2 3.0%, 2023H1 1.7%, 2023H2 0.4%, **2024H1 1.0%**, 2024H2 6.0%, 2025H1 2.8%, 2025H2 0.8%. The v3 “1.6% → 2.6%” pre/post medians reproduce exactly but come from 2024H2 (outside the main post window and straddling the 13G change) and from a **21-day** CAR window (**nrets** = 21), not $[-1, +1]$; one post event has a 2,899% CAR. Start the pre window in 2022 and the sign flips.
- **Sample arithmetic.** 9,234 is the filing count; 4,638 have parsed event dates (0.68/0.66/0.64 for 2022/2023/2024, **0.00 for 2025** — structured-era parser); the memo’s own windows (post Feb 5–Aug 31, 2024; pre 2023-01-01–2024-02-04) give 990 matched events in the existing WRDS output (301 post, in 7 month-clusters; 1,226 post-window filings before matching).
- **F1 rests on 188 parsed filings** (98 pre / 90 post; parse rates 0.68/0.64); Manski bounds on the share filed within 5 bd are $[-7.9, +60.1]$ pp and overlap — a parse rate above 0.72 signs the effect under the worst case. **F1 replicates on the full universe with a clean pre-trend:** median delay 7.0 / 7.0 / 5.0 bd and share within 5 bd 31.9% / 35.7% / 70.6% for 2022 / 2023 / 2024.
- **Two parser bugs.** `pct_of_class` regex cannot match a three-digit percentage (100.0% → 0.0) and takes the first reporting person’s row (understating group stakes): 438 filings (7.1% of the 6,189 parsed) sit at exactly 2.59%, 307 at 0.0, and 22.9% below 5% — not usable for F2 until fixed. `np.busday_count` ignores federal holidays.

4.2 Facts section (facts the model must match)

- **F1 (executed):** the five-business-day deadline binds. On the full 2022–2025 universe of SC 13D originals the median disclosure delay is 7.0 business days in 2022 and 2023 and 5.0 in 2024; the share filed within five business days moves 31.9% → 35.7% → 70.6% — flat before the rule, a break at it. Report delays on a US-federal-holiday calendar, per-quarter parse rates (0.64–0.68) rather than silent drops, worst-case bounds, the bunching mass at exactly five days and the halving of the right tail (p90 23 → 11.1). Confounds to state: the EDGAR filing cut-off moved to 10 p.m. ET on the same date (a clock change can move a compliance share); 2024Q4 sits after the 13G retiming (report 2024Q2–Q3 as the clean post window). Extend to 2025 after the structured-era parser fix. Trivedi (2026) reports +0.35 on the same share with 13G controls — cite as corroboration.
- **F2 (gated on the parser fix):** stake-at-first-13D distribution; target moments from Collin-Dufresne–Fos (2015): median 6.2%, mean 7.68% (screened 2001–2010 activist sample

— replicate their screens; the raw 2022–2025 all-filer median is 9.55%); bunching just above 5%. Unit stake in the model $\rightarrow$ this is a motivating exhibit, not a calibration target. **The sharpest post-rule prediction is the intermediate outcome: does the disclosed stake fall after February 2024?** (CDF: $\sim 23\%$ of the final stake is accumulated between trigger and filing.) Fewer confounds than any CAR test; executable on disk data once the regex is fixed.

- **F3 (a signed prediction, not a fact):** the disclosure jump. A shorter window means less of the activist’s presence is impounded before filing, so the $D = 0 \rightarrow D = 1$ jump should be larger; the confound is the smaller disclosed stake, which pushes it down. Report $[-1, +1]$ CARs (not 21-day windows) by half-year and within stake deciles; state that the current half-year series shows no break; report quantiles, not means.
- **F4 (deferred, 1–3 weeks):** campaign $\rightarrow$ takeover transition via EDGAR submissions API (SC TO-T / DEFM14A / 8-K). Benchmarks: Greenwood–Schor (2009): 18.1% of activist targets acquired within 12 months vs 7.2% for matched firms and 12.6% for non-activist 13D filers (+11pp); Boyson, Gantchev and Shivdasani (2017, JFE): over one-third of targets receive a bid within two years (22% on the risk-arbitrage-excluded “activism merger” measure). (*v3’s “70% within 2 years” was wrong.*)

4.3 H1 — Attenuation of the liquidity slope (a before/after test, with a dose)

Hypothesis. Shortening the accumulation window moves information release from pre-filing inference to the filing event; the sensitivity of 13D announcement returns to pre-event liquidity should change after February 5, 2024. **The sign is not yet a theorem** (§3.2, items 2 and 4): T2 signs the ex-ante premium’s sensitivity on the threshold margin; the disclosure jump’s liquidity slope under the joint window shock ($d\kappa < 0, d\tau_w > 0$) must be derived first and the prediction pre-registered in the paper before the estimate is reported.

Specification.

$$\text{CAR}_i[-1, +1] = \beta_1 \text{Amihud}_i + \beta_2 (\text{Amihud}_i \times \text{Post}_i) + \gamma \text{Post}_i + \delta' X_i + \text{FE} + \varepsilon_i,$$

with the derived sign restriction on β_2 ; also on $\text{CAR}[-10, +1]$ and in quantile regressions (medians are the honest summary). Identification rests on three things, reported in this order: (i) a **calendar-time coefficient figure** ($\hat{\beta}_1$ by quarter, 2022Q1–2025Q4, with CIs) — the pre-trend test; (ii) a **bindingness dose**: for each filer, the pre-2024 leave-one-out median disclosure delay (3 pre-period filings); filers already inside five days are plausibly untreated, giving a genuine triple difference $\text{Amihud} \times \text{Post} \times \text{Dose}$ with its own first stage; (iii) **randomization inference over placebo rule dates** (every month-start; February of 2022, 2023, 2025) — the answer to 7 post month-clusters, not a month-clustered SE. Cluster by firm (403 repeat firms, max 9 events). State the minimum detectable effect: with $n \approx 990$, post share 0.30 and CAR SD ≈ 0.15 , $\text{SE}(\beta_2) \approx 1.0\text{pp}$ per SD of Amihud, so $t = 2$ needs $\approx 2.1\text{pp}/\text{SD}$ — larger than the entire Amihud level loading AFS estimate (+0.9pp/SD). H1 may end as a bounded null; say so in advance.

Sample and data. All SC 13D originals 2022-01-01 to 2025-12-31 (9,234 filings; 65% CRSP match projected, not yet executed on the full sample), US common stocks (shrcd 10/11); event day = first trading session after the EDGAR *acceptance* timestamp (the 10 p.m. ET cut-off change re-maps filing dates for exactly the treated period; `accepted_after_4pm` is already parsed). CARs from the on-disk CRSP daily snapshot (2021–2025, 14,092 PERMNOs) with Ken French factors; events and stakes from the repo’s EDGAR pipeline; filer type from a hand-coded random 150 with

reported precision/recall (a name regex is not a classifier); attrition table pre and post separately (filings → originals → CUSIP → CRSP → common stock → valid Amihud → final).

Illiquidity. Amihud (2002) as a **pre-event average over trading days** $[-250, -11]$ relative to filing (excluding the accumulation window: CDF show Amihud falls ~46% on filer trade days), $ILLIQ_i = D_i^{-1} \sum_t |r_{i,t}| / DVOL_{i,t}$ with $D_i \geq 100$; use $\ln(1 + ILLIQ)$ winsorized at 1/99; log-cap and log-turnover throughout; report the Amihud–log-cap correlation; preferred specification uses within-industry-and-size-decile ranks so β_1 is not a size slope.

Confound handling (full list, replacing v3’s four items). Same date (2024-02-05): 13D initial deadline (the shock); 13D amendment deadline “promptly” → 2 business days (flag events with a 13D/A inside the window); 13G→13D switch timing; **EDGAR cut-off 10 p.m. ET** (event-day realignment above); cash-settled derivatives and group-formation guidance (flag). Later: 13G retiming 2024-09-30 (main post window ends 2024-08-31; kills 13G as a control), structured data 2024-12-18 (measurement break). Inside the post window: T+1 settlement 2024-05-28 (report Feb–May and Jun–Aug halves separately; a dummy is not identified). Before: universal proxy (meetings after 2022-08-31; costs half the pre window if 2023-start is kept — show 2022-start as robustness); **anticipation** (proposed 2022-02-10, adopted 2023-10-10 — donut Oct 2023 → Feb 2024). Throughout: the 2022–23 hiking / 2024 easing cycle and activism-volume cycle (month FE where identified; a non-event matched-firm-quarter Amihud slope as the background-slope control). Reverse causality (activism deteriorates target liquidity: Meles et al. 2026) is why Amihud is measured strictly pre-event.

Selection honesty. Filing is a choice and the marginal filer changes with the rule; report filer-mix diagnostics pre/post; interpret β_2 as a reduced-form information-content change, with an AFS-style selection correction as robustness, not headline. 13G filers are a *composition diagnostic* (13D/13G share), not a placebo (different population, own treatment on 2024-09-30).

4.4 H2 — Run-up to jump shift (the Collin-Dufresne–Fos template), in levels

The 10-to-5-day compression should shift returns from the pre-filing run-up into the filing-day jump. **Do not use a ratio of signed CARs** (v3’s ρ has a denominator that is negative for ~40% of events and no finite mean). Anchor on the parsed *event* (crossing) date and stack three windows per event — run-up $CAR[-40, -1]$ (event-anchored), the crossing→filing segment whose *length is the treatment*, and the jump $CAR[-1, +1]$ (filing-anchored):

$$CAR_{iw} = a + b \text{Post}_i + c \mathbf{1}\{w = \text{Jump}\} + d (\text{Post}_i \times \mathbf{1}\{w = \text{Jump}\}) + \delta' X_i + \varepsilon_{iw},$$

clustered by event and firm; d is the differential post-rule change in the jump relative to the run-up — the one genuine within-event difference-in-differences in the design. Benchmark: the published Collin-Dufresne–Fos (JF 2015, 1994–2010) report a run-up of about 3% over $(-60, -1)$ and a two-day filing jump of about 2.5%; the 7% / 3% figures quoted in v3 are from the 2012 NBER WP (18452, 2001–2010) — cite whichever version is used, not both. CDF’s finding that Kyle’s λ , PIN and Amihud *fall* during accumulation is the load-bearing institutional fact for the model’s architecture. Sample = events with parsed crossing dates (none in 2025 until the parser fix).

4.5 H3 — The hump as a falsifiable cross-section (a prediction the theory offers; not promised)

For 13D-target deals receiving a bid within 12 months, hand-collect offer prices (SC TO-T / 8-K Item 1.01 / DEFM14A) and compute

$$\text{Premium}_i = \frac{P_i^{\text{offer}}}{P_i^{\text{unaffected}}} - 1,$$

with $P^{\text{unaffected}}$ at -42 trading days (Schwert 1996) and SDC’s $-1\text{d}/-1\text{w}/-4\text{w}$ as robustness, using $\min(-42\text{d}, \text{day before the 13D})$ because the 13D typically precedes the bid. The model counterpart is $\Delta^{\min}$ *conditional on a takeover* — the premium given a bid, not the announcement return; the unconditional analogue is $p \cdot \Delta^{\min}$, so report $\Pr(\text{bid} \mid q)$ alongside (conditioning on a bid conditions on the outcome κ moves). Test the non-monotone prediction with terciles or a quadratic in Amihud and the Lind–Mehlum (2010) U-test with a Fieller interval for the turning point; the v3 statement $\theta_3 > \theta_1 \wedge \theta_3 > \theta_5$ is an intersection–union hypothesis, not a Wald test, and five quintiles on a few hundred deals (~ 60 per cell, premium SD $\sim 35\text{pp}$) give $t \approx 1.25$ for an 8pp hump — state the MDE. Separate the channels: control for acquirer–target Amihud where the acquirer is public (Huang–Maharjan–Nanda’s currency channel is stock-deal-specific and null in cash deals — **require the hump in both cash and stock**), split public vs private acquirers (Massa–Xu’s channel is public-bidder-specific), keep the Amihud level (AFS $+0.9\text{pp}/\text{SD}$ loading on both filing types) so the hump is identified off curvature. Anchors: Massa–Xu (2013, JFQA): $+10\%$ premium per SD of target liquidity; Huang–Maharjan–Nanda (2024, JCF): -4.5pp per SD of acquirer–target liquidity difference, stock deals only; Celentano–Levine (2025): mean premium 36.6% , activism interaction -5.2pp . **Status:** weeks-to-months of hand collection with a selection problem — offered as the falsifiable prediction with the design spelled out; not promised in the submission draft.

4.6 Calibration moments (split into what the model can produce)

Panel A — targeted in calibration (model object in brackets).

Moment	Value	Source
13D announcement CAR level [disclosure jump $\mathbb{E}[J \mid D = 1]$]	6.34% (level only; the 75.2% treatment share has no model counterpart)	Albuquerque–Fos–Schroth 2022
Mean bid premium [$\Delta^{\min} \mid \text{bid; pins } (m_0, m_1) / D7 \text{ primitives}$]	36.6%	Celentano–Levine 2025
Activism $\rightarrow$ premium interaction [$\Delta^{\min} \mid D=1 \text{ vs } D=0$; say which of the two numbers is matched]	-13.7% (-5.2pp)	Celentano–Levine 2025
Takeover prob. lift from activism [bidder entry p ; a calibration input if p is a primitive]	$+11\text{pp}$ (12m: 18.1% vs 7.2% matched); $+7.7\%$ relative	Greenwood–Schor 2009; C&L 2025

Panel B — signs the model must reproduce.

Moment	Value	Source
Liquidity $\rightarrow$ targeting sign [$\partial \Pr(\text{engage})/\partial \kappa$]	positive (baseline probit: 0.33% $\rightarrow$ 0.73%; IV “somewhat larger”, not quantified)	Norli–Østergaard–Schindele 2015
Voice $\rightarrow$ exit mode shift in liquidity [$\partial \Pr(P)/\partial \kappa$ vs $\partial \Pr(E)/\partial \kappa$]	−6.9pp on 13D share per SD	Edmans–Fang–Zur 2013

Panel C — motivating facts the static, unit-stake model does not target.

Moment	Value	Source
Run-up vs filing-day jump	JF 2015 (1994–2010): $\sim 3\%$ over $(-60, -1)$ vs $\sim 2.5\%$ two-day jump; NBER WP 18452 (2001–2010): $\sim 7\%$ vs $\sim 3\%$	Collin-Dufresne–Fos 2015 / 2012
Stake at first 13D Pre-announcement accumulation	median 6.2%, mean 7.68% 54% of block, 8.5% trading profit	Collin-Dufresne–Fos 2015 Norli–Østergaard–Schindele 2015
Campaign costs by stage	negotiations \$2.9M / board seat \$1.8M / proxy contest \$5.9M	Gantchev 2013

5 Edit Map: draft_v2 $\rightarrow$ draft_v3

Area	Action
Title / abstract / intro	Reframe around the disclosure rule as a lever on revelatory efficiency; institutional hook first (1968\$ $\rightarrow$ \$2024 compression; the Feb-2024 shock); abstract states exactly what is proved where — no “theorem” for region-certified or to-do objects; the narrower checkable claim in the abstract, the “revelation technology” phrase (if kept) only in the introduction

Area	Action
Literature	Fix misattributions: the -13.7% premium finding is Celentano–Levine 2025 , not AFS (rename <code>prop:d7-afs</code>); delete the unsupported Johnson–Swem window-isomorphism sentence (attribute the $\sigma^2 T$ mapping to Back et al. 2018, pp. 1452–1453); reframe EGJ 2015 (analytic region-wise uniqueness, AER Prop. 1 — cite the posture, not “numerical uniqueness”); add Kyle–Vila 1991 (RAND), Mello–Repullo 2004 (FRL), Burkart–Lee 2022 (missing from <code>bibliography.bib</code>), Burkart–Lee–Voss (ECGI 956/2024), Levit–Malenko–Maug (JF 2024, JF 2026), Cetemen et al. (JF 2026), Bebhuk–Brav–Jackson–Jiang 2013, Massa–Xu 2013, Huang–Maharjan–Nanda 2024, Ben-David–Bhattacharya–Huang–Jacobsen (JF 2026, four authors; also missing from <code>bibliography.bib</code>), Trivedi 2026, Polk et al. 2024, Corum 2025, Bishop–Fos–Jiang–Partnoy 2026, Meles et al. 2026
Model	Three-branch menu under (A2’); τ_θ and τ_w as explicit primitives (§2.3); bidder conditions on (P, D) as a sufficient statistic (replaces lem:dropA7); one-sentence no-manipulation remark; ternary noise + flag defended as the <i>minimal</i> technology partitioning inference (LMM 2024 as precedent, not proof); D7 symbols relabeled (r, γ, ϕ, ψ) ; resolve δ (0.95 vs 1) and the (A5) contradiction (l. 696 vs l. 1010)
Equilibrium	Compact; T2 stated as the two-term product rule with the threshold-margin lemma proved and the window-margin sign left as a calibration statement; T3 (GE region version) added only if the first-derivative comparison is proved; existence under (A1),(A2),(A4),(A5a); meta-material collapsed to one remark
Comparative statics	R1 repackaged as endpoint lemma (fixed cutoffs) + certified region + falsifiable prediction; Remark d5-vacuous kept verbatim; sensitivity figures trimmed; the disclosure figure re-drawn at a calibration with meaningful ω_P and with the elasticity reported

Area	Action
Testable implications	Five $\rightarrow$ four numbered predictions, each mapped to a dataset and to a model object (the disclosure jump for CARs; $\Delta^{\min}$ bid for premia); cross-country thresholds as the out-of-sample agenda and the clean τ_θ experiment
Welfare	Halved: first-best subsection cut; disclosure-wedge proposition downgraded to a remark (its hypothesis (ii) is admittedly untested); positive framing only — optimal-threshold welfare language belongs to Ordóñez-Calafi–Bernhardt
Extensions	Benchmark regimes become the $\tau_w \in \{0, 1\}$ endpoints; noisy rumors kept as the wolf-pack bridge, labeled numerical
Empirical section (new)	Facts + H1 + H2 per Section 4 (before/after with dose; stacked-window levels); H3 as the offered prediction; confound table; attrition table; pre-registration box; calibration panels A/B/C
Appendices	D7/D8 stay as microfoundation/region-theorem records; bargaining appendix compressed to ~2 pages; duplicated curvature and endpoint proofs folded; honesty labels preserved verbatim — never weakened

6 Journal Targets and Risks

Target	Fit	Key risk
JF / RFS (primary)	Theory-forward + institutional anchor + one clean shock; matches the Levit–Malenko–Maug / Burkart–Lee taste	T2 must become a real theorem on a stated margin with a non-trivial magnitude; the hump must be honestly labeled; differentiate from Celentano–Levine, Cetemen et al. and Corum on page 1
JFE (alternative)	If the empirical section grows toward 30–35% (premium subsample + transitions)	JFE’s activism space is structural-estimation-dominated; the before/after must carry the dose design and the placebo-date inference
JFQA / RAPS (fallback)	Certain fit	—

Risk register (revised). (i) The first-stage fact is already public (Trivedi 2026; Polk et al. 2024) and Fos–Jiang–Partnoy are working the disclosure-friction margin — the differentiator must be the premium object and the model, not the shock alone; execute F1 (universe) + the stake-at-filing test + the H1 calendar-time figure first. (ii) The disk data show no CAR break at the rule date — plan for H1 as a possibly bounded null and say so in the paper. (iii) T2’s sign is margin-dependent and its baseline magnitude is negligible — the theory work must precede any claim that the shock “tests T2”. (iv) The R1 counterexample invites “the hump is fragile” — preempt by featuring the economic boundary and Remark d5-vacuous. (v) CAR-critique referees — CARs are only ever disclosure price impact; premia come from hand-collected offer prices. (vi) “Too many moving parts” — the three-branch cut and appendix exiles are the answer, but note that the honest τ_θ device re-splits the disclosed branch. (vii) Reproducibility: the Fact-2 execution code and the parser fixes must be committed; the gitignored CRSP snapshot must be backed up off-repo. (viii) Reverse causality (activism $\rightarrow$ liquidity) — Amihud strictly pre-event.

Execution sequence (post-approval, revised). (1) Owner approves the angle and chooses the paper’s primary margin (τ_θ vs τ_w); (2) theory: add τ to `numerical/`, re-solve the three-branch model, derive the disclosure-jump slope, prove the threshold-margin lemma, resolve (A5)/ δ , re-anchor the calibration; (3) empirics: parser fixes (percent regex, holidays, structured era) $\rightarrow$ F1 universe $\rightarrow$ stake-at-filing $\rightarrow$ event-day realignment $\rightarrow$ H1 calendar-time figure $\rightarrow$ dose design and placebo inference $\rightarrow$ H2 stacked levels; (4) premium subsample only if resourced; (5) compile gates (0 errors, 0 undefined references) plus `make clean && make all`. Time estimate for (3): weeks, not “3–4 days”.

7 Change log v3 $\rightarrow$ v3.1 (2026-08-19)

Referee round: five parallel referee agents (theory, literature/originality, novelty scan, empirical design with execution on disk data, facts verification) plus two adversarial verifiers; full report at `quality_reports/reports/2026-08-19_framework_v3_referee_report.md`; agent reports in `research/review_v3/`. Edits: (1) §1 originality claims restated; competitors named (Trivedi, Polk et al., Corum, Bishop et al., Cetemen et al., Kyle–Vila, Mello–Repullo). (2) §2 notation block; (A2′) for the prune; bidder information set as a sufficient statistic; τ split into τ_θ and τ_w with the Back-et-al. κ -mapping; existence assumption list; δ and (A5) contradictions flagged. (3) §3 T2 restated as a two-term product rule with numeric evidence (threshold margin attenuates; window margin does not at fixed cutoffs; baseline magnitude $<1\%$); disclosure-jump object introduced; T3 labeled to-do; T4 with Remark d5-vacuous, $\kappa^\dagger = 0.58$, Mello–Repullo; T5 λ_{crit} caveat and C&L “rationalizes” withdrawn; no-manipulation reduced. (4) §4 rewritten around what the disk data show; F1 universe numbers and bounds; F2 parser gate and stake-at-filing test; F3 as a signed prediction; F4 Boyson corrected; H1 as before/after with dose, calendar-time figure, placebo inference, MDE, full confound table; H2 in stacked levels; H3 demoted; Amihud definition; calibration panels; numbers corrected (+7.7% relative; Norli baseline probit; CDF window). (5) Edit map, targets, risks and execution sequence updated. Honesty labels preserved or strengthened throughout.