

The Two-Round Blockholder Disclosure Model

A standalone model note

v4 theory lane, ticket 30 (T2j)

Card stamp 2026-08-28, follow-up curation

Abstract

This note states the v4 two-round model in one place: primitives, timing, the partition the disclosure rule induces, the equilibrium notion, the standing hypotheses, and the eight-result ledger with its honesty labels. It then records what the v4 model takes from the frozen manuscript `draft_v2.tex`, what it simplifies, and what it drops; and it lists the executed evidence, failures included. Every statement of content here is transcribed from `research/model_v4/MODEL_CARD.md` (version stamp 2026-08-28, follow-up curation) and `research/model_v4/LABEL_LEDGER.md`. Four card events since the 2026-08-23 stamp are carried below: P1 restored to PROVED on 2026-08-25 on its ticket-35 repair; $A(\tau)$'s ticket-33 evidence note of 2026-08-25; the 2026-08-27 A6 panel ruling; and the 2026-08-28 follow-up curation — the last two moving no label. The card is the single source of truth: where a proof file and the card differ, the card wins. The conditionality attached to each result is not decoration. It is the result.

1 Position and object

The disclosure rule *is* the market's partition. The model asks how liquidity κ moves bidder entry and the expected takeover premium when a stake threshold τ and a filing window T split control-node histories into a **flagged** cell (the filing has landed before the control decision) and a **pooled** cell (it has not).

The control-outcome object is the expected engagement-related premium $\Delta^{\text{act}}(\kappa, \tau, T)$. The price-path objects are the run-up path R_d , the cumulative run-up R , and the filing-day jump J . Lower τ means a tighter threshold margin; lower T means a tighter window margin.

Throughout, $\kappa \in [0, 1]$ is *noise-trading intensity* — never depth, volume or turnover. Upright T is the filing window; $\mathcal{T}$ is the outer best-response map. Bare λ is not available (it is the appropriability coefficient of the tender-game appendix), and neither is a bare C , W , $\mathcal{S}$, u or g : every one of those letters is overloaded in this project and always carries its subscript.

2 Primitives

Fundamentals, signal, bidder

- $v \sim N(\mu_v, \sigma_v^2)$ is the target's standalone value; the blockholder observes $s = v + \varepsilon$ with $\varepsilon \sim N(0, \sigma_\varepsilon^2)$ independent of v . The Gaussian projection coefficient is $\beta = \sigma_v^2 / (\sigma_v^2 + \sigma_\varepsilon^2) \in (0, 1)$, so that $\mathbb{E}[v \mid s] = \mu_v + \beta(s - \mu_v)$. (The letter is β , the frozen manuscript's name; bare λ is reserved.)
- The bidder draws a private synergy shock $\xi \sim N(0, \sigma_\xi^2)$ independent of (v, s) ; $\bar{S}$ is mean bidder synergy and $K > 0$ the entry cost.
- m_0, m_1 are the takeover premia without and with engagement, with $m_1 > m_0$ and $m_0 \geq 0$; $\Delta_m = m_1 - m_0 > 0$ is the premium wedge, and $\Delta_V \geq 0$ is the non-takeover value engagement creates. The restriction $m_0 \geq 0$ is load-bearing: it is what makes the inner pricing fixed point exist, be unique and be continuous (Section 6, A5).

Trading, stakes, the legal clock

- Noise is a ternary mark of size $\bar{z} > 0$: $\mathbb{P}(z_d = 0) = 1 - \kappa$ and $\mathbb{P}(z_d = \pm \bar{z}) = \kappa/2$.
- b_0 and $\bar{b}$ are the initial and maximum stake, $0 \leq b_0 \leq \bar{b}$, with $b_0 < \tau$ maintained: a pre-existing crossing is outside the core.
- $\mathcal{J}$ is a finite ordered menu of complete contingent plans, least to most aggressive, $|\mathcal{J}| = J < \infty$; $a_j \in \{0, 1\}$ is the engagement attached to plan j , with $a_j = 1$ for Voice plans and 0 for Exit and Hold.
- $B_j(s, d) \in [0, \bar{b}]$ is the cumulative pooled stake at business day d , with $B_j(s, -1) = b_0$; for Voice plans $\partial_d B_j \geq 0$ and $\partial_s B_j \geq 0$, Hold is constant and Exit weakly decreasing. For *every* plan and every d , $s \mapsto B_j(s, d)$ is Borel — automatic for Voice and Hold, a genuine requirement for Exit, and needed because pooled pricing integrates over every type including Exit types.
- $b_j^*(s) = B_j(s, H)$ is the terminal target stake; H is the finite control-decision horizon in business days.
- $c_j(s; \tau) = \inf\{d : B_j(s, d) \geq \tau\}$ is the threshold-crossing date ($+\infty$ if it never happens); $f_j = c_j + T$ is the legal filing date; the flag lands iff $f_j \leq H$, equivalently iff $B_j(s, H - T) \geq \tau$.
- $D_j(s; \tau, T) = \mathbf{1}\{a_j = 1, c_j < \infty, f_j \leq H\}$; $D = 1 \Rightarrow a = 1$. $B_j^F = B_j(s, f_j)$ is the stake at filing and $Q_j^F = b_j^*(s) - B_j^F \geq 0$ the flagged-round order for a Voice plan.
- Γ is a finite ordered coarsening from stake increment to pooled order mark, $q_{jd}(s) = \Gamma(B_j(s, d) - B_j(s, d - 1))$, and observed pooled order flow is $X_d = q_{jd} + z_d$. (Γ , not ψ : ψ is tender-game pivotality and χ is the frozen manuscript's cost parameter.)

Information, prices, control outcome

$\mathcal{H}_d^P = (X_0, \dots, X_d; \text{flag landed by } d)$ is the pooled public history, finite. $F = (B^F, a = 1)$ is the truthful filing message and $\mathbf{S}_F = (B^F, Q^F, a = 1)$ is F augmented by the flagged-round order. The control-node information set is

$$\mathcal{I}_H = \mathcal{H}_H^P \text{ on the pooled cell } \{D = 0\}, \quad \mathcal{I}_H = \mathbf{S}_F \text{ on the flagged cell } \{D = 1\},$$

the bidder's own ξ being private. The engagement posterior is $\pi(\mathcal{I}) = \mathbb{P}(a = 1 \mid \mathcal{I})$, equal to 1 on the flagged cell; entry probability is $p(\mathcal{I}) = 1 - \Phi((P + K + m_0 + \pi\Delta_m - \bar{S})/\sigma_\xi) \in (0, 1)$; and with entry indicator $\mathbf{B}$ the terminal shareholder payoff is

$$Y = (1 - \mathbf{B})(v + a\Delta_V) + \mathbf{B}(P + m_0 + a\Delta_m).$$

Prices are competitive, $P(\mathcal{I}) = \mathbb{E}[Y \mid \mathcal{I}]$, with the convention $P_{-1}^P := \mathbb{E}[Y]$ for the pre-trading pooled price (needed whenever $c = 0$, which $T = H$ forces on every flagged history). The genuine fixed point sits at control nodes: at an earlier pooled date $d < H$ the price is a tower expectation of already-solved control-node values, with no self-reference. $P_{\text{ND}}(\mathcal{H}_{f-}^P) = P_{f-}^P$ is the *not-yet-disclosed* price at the same realised order flow — not a never-disclosed counterfactual, under which reading the run-up/jump identity acquires a residual term. Then

$$R_d = P_d^P - P_{c-}^P, \quad R = P_{f-}^P - P_{c-}^P, \quad J = P^F - P_{\text{ND}}, \quad P^F - P_{c-}^P = R + J \text{ exactly.}$$

J is *not* claimed κ -invariant.

The blockholder's objective

$$U_j(s) = \mathbb{E}[b_j^*(s)Y - \mathcal{C}_j^{\text{trade}} - a_j C_j(s) \mid s, j],$$

the expected terminal value of the position the plan builds, net of what it costs to build and to engage: $\mathcal{C}_j^{\text{trade}}$ is the execution outlay (increments valued at the pooled prices P_d^P up to the

plan's last pooled date, plus $Q_j^F(s)P^F$ when $D_j = 1$) and $C_j(s) \geq 0$ the engagement cost. Only two properties of U_j are ever used: *plan-locality* — U_j depends on j only through the executed stake path, the prices paid on it, the terminal stake, the engagement flag and the cost — and *integrability*, $\mathbb{E}[\max_j |U_j|] < \infty$.

Premium and sensitivity objects

$$\begin{aligned} h(\mathcal{I}) &= \pi(\mathcal{I})p(\mathcal{I}), \quad h \geq 0, \quad h(0) = 0; & \Delta^{\text{act}} &= \Delta_m \mathbb{E}[h(\mathcal{I}_H)] \geq 0; \\ M_F &= \Delta_m \mathbb{E}[h \mid D = 1], \quad M_P = \Delta_m \mathbb{E}[h \mid D = 0]; & \Omega &= \mathbb{P}(D = 1) = \mathbb{P}(a = 1) \omega_a; \\ \mathcal{S} &= |\partial_\kappa \Delta^{\text{act}}|, \quad \mathcal{S}_P = |\partial_\kappa M_P|; & C_h(\bar{\pi}) &= h(0) - 2h(\bar{\pi}/2) + h(\bar{\pi}). \end{aligned}$$

$\omega_a = \mathbb{P}(D = 1 \mid a = 1)$ is the disclosed share of engagements, the calibration target; Ω is the unconditional flagged weight and is the frozen manuscript's ω_P . $\bar{\pi}$ is the **upper support point** of the pooled engagement posterior in the $A(\tau)$ representation — *not* the pooled engagement share. The share is the mean $\mathbb{E}[\Pi_\kappa]$, which under $A(\tau)$ is κ -invariant (a mean-preserving spread) and is strictly below $\bar{\pi}$ in any non-degenerate case, equalling $\bar{\pi}/2$ only under level symmetry $A_0 = A_1$. Reading $\bar{\pi}$ as the mean forces a point mass at $\bar{\pi}$ with $A'_\kappa = 0$ and zero interior motion — degenerate, and excluded. A'_κ is the common derivative of the $A(\tau)$ weights ($A'_0 = A'_1 = A'_\kappa$, $A'_{1/2} = -2A'_\kappa$). Weight-effect ratios are W_τ, W_T (e.g. $W_T = (1 - \Omega(\tau, 5))/(1 - \Omega(\tau, 10))$) and composition-effect ratios are C_τ, C_T (e.g. $C_T = \mathcal{S}_P(\tau, 5)/\mathcal{S}_P(\tau, 10)$), unsigned.

For the general-equilibrium dominance-and-contraction node: $k = (k_1 \leq \dots \leq k_{J-1})$ is the cutoff vector on the compact ordered polytope Θ with parameter vector ϑ ; $\mathcal{T}(k; \vartheta)$ the outer best-response map; $L_{\mathcal{R}} = \sup_{\mathcal{R}} \|D_k \mathcal{T}\|$ the contraction bound; $r_\tau = -\tau$, $r_T = -T$ the strictness coordinates (higher r = tighter); $g_r^{PE} = -\text{sgn}(d\Delta^{\text{act}}/d\kappa) \partial_{\kappa r} \Delta^{\text{act}}$ the direct fixed-policy attenuation margin; $\bar{k}_x = |\partial_x \mathcal{T}|/(1 - L_{\mathcal{R}})$ and $\bar{k}_{\kappa r}$ the inversion-free derivative bounds;

$$\mathcal{B}_r^{GE} = |\Delta_{\kappa k}| \bar{k}_r + (|\Delta_{kr}| + |\Delta_{kk}| \bar{k}_r) \bar{k}_\kappa + |\Delta_k| \bar{k}_{\kappa r}, \quad \eta_r = g_r^{PE} - \mathcal{B}_r^{GE}$$

the GE remainder bound and the slack on a dominance-and-contraction region $\mathcal{R}_r$.

3 Timing: the two rounds

1. Nature draws v and the blockholder's signal s ; the blockholder picks one complete contingent plan j from the finite ordered menu $\mathcal{J}$.
2. **Round 1 — pooled trading.** The plan's stake path executes over business days $d = 0, \dots, H$. Market makers see pooled order flow and set $P_d^P = \mathbb{E}[Y \mid \mathcal{H}_d^P]$. There is **no within-window re-optimisation**, and hence no feedback from realised order flow or prices into the path: $B_j(s, d)$, $q_{jd}(s)$ and Q_j^F are functions of (j, s, d) and (j, s, τ, T) alone. This is a numbered hypothesis, not a modelling aside: two steps of L2 fail without it.
3. **Disclosure node.** The flag lands iff $D = 1$, i.e. iff the plan engages, crosses τ at some date $c < \infty$, and $c + T \leq H$. The filing reveals $F = (B^F, a = 1)$. **The flag terminates the pooled round.** Pooled trading stops when the filing lands; the flagged round follows it and the bidder acts after that. Without this reading $Q^F = b_j^*(s) - B_j^F(s)$ is not the blockholder's whole residual position, and the flagged-round step of P1 fails.
4. **Round 2 — flagged trading, then the bidder.** If $D = 1$ the blockholder submits Q^F , the market prices $P^F = P(F, Q^F)$, and then the bidder decides. If $D = 0$ there is no flagged round and the bidder acts on the pooled history.

Sequence: pooled round $\rightarrow$ flag or no flag $\rightarrow$ flagged round if applicable $\rightarrow$ bidder decision.

4 The partition

The disclosure indicator $D = \mathbf{1}\{a = 1, c(\tau) + T \leq H\}$ maps every control-node history into exactly one of two cells, $\mathcal{C}_F$ (flagged) and $\mathcal{C}_P$ (pooled), exclusive and exhaustive by construction. The cells are keyed by the disclosure rule and by nothing else:

- On $\mathcal{C}_F$ the control-node information set is the flagged tuple $\mathbf{S}_F = (B^F, Q^F, a=1)$: the market knows engagement happened, $\pi = 1$, and the pooled residual is (under A7') pure noise conditional on $\mathbf{S}_F$.
- On $\mathcal{C}_P$ the control-node information set is the pooled public history $\mathcal{H}_H^P$: the market has an inference problem, and the posterior $\pi(\mathcal{H}_H^P)$ is where κ enters.
- The clock equivalence: for every Voice plan, $f_j \leq H \iff B_j(s, H - T) \geq \tau$. A tighter threshold (lower τ) or a tighter window (lower T) moves histories from $\mathcal{C}_P$ into $\mathcal{C}_F$ and nowhere else.
- Each flagged history yields B^F, R_d, R, J , with $P^F - P_{c^-}^P = R + J$ exactly.
- $\Omega = \mathbb{P}(D = 1)$ is the cell weight; A8 asks for $0 < \Omega < 1$, which is required only for positive cell mass, never for the structural partition itself.

5 Equilibrium notion

Cutoff perfect Bayesian equilibrium: (i) a weakly ordered cutoff vector $k = (k_1 \leq \dots \leq k_{J-1})$ mapping s into a plan; (ii) sequentially optimal pooled and flagged components; (iii) Bayes-consistent beliefs on path; (iv) competitive pooled and flagged prices at their fixed points; (v) the bidder-entry rule; (vi) off-path beliefs as limits of full-support perturbations. Weak inequalities permit collapsed action regions, Hold included. Existence is Brouwer on the compact ordered polytope Θ for the outer map $\mathcal{T}(k; \vartheta)$, i.e. $k = \mathcal{T}(k; \vartheta)$. **Uniqueness is not claimed.**

6 Standing hypotheses

A1 Independent primitives. v, ε, ξ and all z_d mutually independent; all variances strictly positive.

A2' Finite model, amended boundedness. Finiteness is unchanged: the plan menu $\mathcal{J}$, the image of Γ , the noise support $\{-\bar{z}, 0, +\bar{z}\}$ and the horizon H are all finite. The old flat boundedness clause was *false* and is replaced by: prices and payoffs are locally bounded in (s, ϑ) on the maintained parameter set, and $\mathbb{E}[\max_{j \in \mathcal{J}} |U_j|] < \infty$ for every $k \in \Theta$. Flat global boundedness is inconsistent with the rest of the model — v is Gaussian and the flagged region is unbounded in s under A7', so Y , prices and U_j are unbounded. Integrability is all any proof consumes.

A3 Ordered plans, single crossing. At every belief/price system, adjacent-plan payoff differences cross zero at most once in s , and the preferred plan is weakly increasing in s .

Evidence note added 2026-08-27 (A6 panel, in passing — a separate finding, deliberately not folded into the A6 note). At the implemented calibration **A3 itself fails, at two independently-found loci, upstream of A6.** (i) At $(\kappa=0.5, \tau_{50}, T=5)$ with k_2 on an **open set** above cell edge 6 (verified at offsets 10^{-9} through 2×10^{-2}), $U_V - U_H$ has **three strict sign changes** ($s = 1.5754434/1.5833333/1.5902426$; middle excursions $2.4\text{--}2.8 \times 10^{-4}$ against a 10^{-9} payoff tolerance), the pointwise argmax runs H,V,H,V single-valued on each interval, so **no weakly increasing selection exists** — $\mathcal{S}(k) = \emptyset$ and Step 13's $\mathcal{T}$ is **undefined** there, not merely discontinuous. (ii) At $(\kappa=0.15, 0.05, 5)$ — a ticket-34 UNRESOLVED node — the argmax reverses **VOICE** $\rightarrow$ **HOLD** across cell edge $s = 1.659062163$ at **both** located fixed points: the preferred plan decreases in s . The route is the s -direction step of U_{VOICE} ($n(s)$ is integer-valued — Step 15(i) / WHERE IT FAILS 4's card-legal

counterexample, instantiated by the solver’s own N_GRID note) interacting with the off-path price snap. **No conflict with ticket 34’s “the A3 and A6 proxies pass at every achieving seed”**: those proxies are local screens — the A3 proxy tests residual slope signs at the two candidate cutoffs and the A6 proxy tests Θ -corner non-pinning at the closest seed (t2_p1_fournode_recheck.py) — and neither measures argmax monotonicity over s nor continuity of $\mathcal{T}$ in k , so both are silent on these findings. **Candidate mechanical account of ticket 34’s four UNRESOLVED nodes**, on file and UNCHECKED beyond the one node probed: at the $\kappa = 0.15$ node one fixed point sits exactly on the edge where $U_H - U_V$ jumps through zero without crossing it, and the panel’s residuals (payoff 3.06×10^{-4} – 1.77×10^{-3} at cutoff residuals of 10^{-11} -grade) **bracket ticket 34’s recorded range exactly**; the k -direction jump mechanism does **not** explain those nodes (no proximity correlation — the substantiate panellist’s own recorded negative). *Swept 2026-08-28 over the other three nodes (quality_reports/fixes/t2_t34_account_sweep.py/.json, pre-registered three-way rule): the account HOLDS at all three. At ($\kappa=0.15, 0.075, 1$) and ($\kappa=0.85, 0.075, 1$) a located fixed point sits on an $n(s)$ cell edge — 1.460178993 (offset $\sim 10^{-13}$, where 10 of the 30 recheck seeds land) and 1.517932397 (offset $\sim 10^{-12}$, reached by **no** seed and found only by the direct edge test) — with $U_H - U_V$ jumping through zero without crossing. Neither pin is its node’s achieving basin: their payoff residuals, 1.398×10^{-3} and 1.314×10^{-3} , sit above the recorded bests 1.059×10^{-3} and 3.061×10^{-4} , each equalling the larger one-sided jump to at most 2.7×10^{-4} relative — a recorded, non-gating quantity. At ($\kappa=0.85, 0.05, 5$) **no pin was found at any candidate edge in** [1.29, 2.11]; the achieving basin’s worst deviation instead sits in the cell immediately above edge 1.583333333 ($0.0250 \sigma_s$ from it), where the same jump through zero occurs, at a deviation/jump ratio of 0.366 — inside the pre-registered factor of 3. Every pin is $n(s)$ -family; the τ -crossing pullbacks yielded none. Probe 5(b)’s distances replicate ($0.0258/0.0437/0.0295 \sigma_s$ vs $0.026/0.044/0.030$). **No node yields a second independent fixed point, so node 15’s residual bracket does not recur** — criterion (ii) rests on reproduction of every recheck basin alone. Diagnostic evidence at one calibration; existence at these nodes stays neither claimed nor denied. No label moves — A3 is a hypothesis; P1 stays PROVED as a conditional. Records: the same panel files as the A6 note.*

A4 Legal-clock discipline. c is the first date the path reaches τ ; the filing lands exactly at $c + T$; filings truthfully reveal stake and purpose; only Voice plans cross in the core.

A5 Inner pricing regularity, mostly demoted to a theorem. Each public-history pricing map has a unique fixed point, continuous in beliefs, cutoffs and parameters. *Under $m_0 \geq 0$ — a card restriction — existence, uniqueness and continuity of the inner fixed point are theorems, not assumptions*: with $\bar{y}(\mathcal{I}) = \mathbb{E}_\mu[v] + \pi\Delta_V$ and $\bar{m}(\mathcal{I}) = m_0 + \pi\Delta_m$, the pricing map is $P \mapsto \bar{y} + \bar{m}p(P)/(1 - p(P))$, strictly decreasing in P wherever $\bar{m} \geq 0$, so it crosses the identity exactly once. A5 is retained only as its continuity clause: the pricing family is continuous in the cutoff vector and the parameters and measurable in the flagged tuple. Because the flagged information sets are continuum-indexed, “unique fixed point” must be read as a measurably selected *family*, not a finite list.

A6 Compact outer self-map. All best-response cutoffs lie in a common compact ordered polytope Θ ; $\mathcal{T}$ is continuous and maps Θ into itself.

Evidence note added 2026-08-27 (A6 panel, Austin-authorized; ruling at item 4 of the open list under What this note does not claim). Two opposed-brief agents (substantiate / defuse) examined the re-derivation’s withheld change 6 (N11) and **converged**; the orchestrator’s adjudication is on file. **The continuity clause fails for the declared construction, and the locus is not the one N11 named.** All k -dependence of U_j runs through the pooled price vector (the flagged layer is k -free under A7-J), and Step 9(b) gives Bayes where $\Lambda_k(h) > 0$ but a k -free plan-uniform posterior on the frontier, so the price system can be discontinuous exactly on $\bigcup_h \partial\{k : \Lambda_k(h) > 0\}$ — a set inside (the finitely many **cell-edge**

hyperplanes $\{k_i = a\} \cup$ (the **collapse faces whose dying plan is the sole generator** of some reachable pooled history). The jump **reaches $\mathcal{T}$ with non-vanishing weight**: U_j integrates those prices against the deviator’s own noise law (weight $\geq \min(\kappa/2, 1 - \kappa)^{d+1}$, independent of the dying plan’s population mass), so the vanishing-mass defusal is **refuted** — **by both panellists, independently**; the largest-weakly-increasing-selection tie-break is pointwise in k and passes the jump through; and no k -independent perturbation family reconciles the limits (at fixed n the system is continuous in k ; the discontinuity is created only as $t_n \rightarrow 0$ — an order-of-limits problem the family choice cannot fix). On collapse faces proper: for $J \geq 3$ menus where a middle plan owns a reachable **exclusive** pooled history entering some U_j , the interior limit $\mu_v + \beta(c - \mu_v)$ varies over the face while any k -free family supplies one constant, so continuity fails at **every face point but at most one** (continuum-face lemma — single-pass panel derivation, **not gate-checked**). The implemented menu is **not** in that class: Exit and Hold pool perfectly in order flow, and its Hold-collapse face is **measured clean** (pooled prices within 4.4×10^{-16} and $\mathcal{T}$ bit-identical as k_1 sweeps to full collapse). At the implemented calibration the failure is live at the **interior $n(s)$ cell edges** instead: measured $\mathcal{T}_2$ jumps of $6.33 \times 10^{-3} / 1.09 \times 10^{-2} / 2.83 \times 10^{-2}$ across $\leq 2 \times 10^{-9}$ steps in k_2 at $(\kappa=0.5, \tau_{50}, T=5)$ — **measured independently by both panellists with separate scripts, agreeing to 3 s.f.**, the belief snap matching the Step 9(b) prediction to $\sim 10^{-8}$, surviving-type controls $\sim 3 \times 10^{-9}$, robust at 1000 $\times$ the breakpoint-merge tolerance; at $(\kappa=0.15, 0.05, 5)$ jumps reach 0.16 and a diagonal crossing of $\mathcal{T}_2$ is **destroyed**. A chamber-interior $\Theta^+ = [1.23, 1.245] \times [1.5253, 1.5506]$ (exhibited) is compact, self-mapping and jump-free at the baseline — Brouwer runs verbatim on it and it contains k^* — but it is **not the Θ Steps 13–14 construct** (they build from the bracket $[s_{lo}, s_{hi}]$, which contains the edges), cannot be exhibited without approximately locating the fixed point first, and **no such chamber exists at the $\kappa = 0.15$ node**, where a fixed point sits exactly on the edge $k_2 = 1.659062163$. **No label moves and none is licensed** — A6 is a hypothesis; P1 stays PROVED as a conditional, in the $A(\tau)$ pattern: what is on record is that its antecedent, read with the Θ the proof constructs, is not satisfied by the implemented calibration. Repairs on file, both outside the equilibrium notion’s declared Brouwer-with-one-fixed-family route: the t -constrained game + Kakutani + $t \downarrow 0$ (**proofs/P1_proof.md** Step 18), and a k -indexed concentration family (constructible; its 0/0 corner unresolved). The implementation’s **OFF_PATH_EPS = 10^{-14}** is the fixed- t constrained game — the standard repair already shipped, with the switch relocated by $\sim 10^{-9}$ rather than removed. Coverage: probes at one node per claim class plus the 27-node census, **not swept over** (κ, τ, T) ; nonexistence is neither claimed nor shown (23/27 sweep nodes converge; a discontinuous self-map may still have fixed points). Records: **threads/2026-08-27_A6_panel_substantiate.md**, **threads/2026-08-27_A6_panel_defuse.md**; probes **quality_reports/fixes/a6_panel_probes_2026-08-27/** (analysis-grade, not curated t2 checks).

Curation note added 2026-08-28. The three decisive measurements are now executed t2 checks: **quality_reports/fixes/t2_a6_edge_jump_check.py/.json** (both panellists’ routes replayed at their own filed brackets — $\mathcal{T}_2$ jumps $6.33 \times 10^{-3} / 1.09 \times 10^{-2} / 2.83 \times 10^{-2}$, agreeing across routes to a relative 1.3×10^{-4} , controls $2.8\text{--}3.6 \times 10^{-9}$, $\pm 10^{-6}$ robustness intact), **t2_a6_node15_check.py/.json** (jump 0.1647, destroyed crossing $+1.0 \times 10^{-7} \rightarrow -6.70 \times 10^{-2}$, edge fixed point to 1.06×10^{-12}) and **t2_a6_collapse_face_check.py/.json** (pooled prices within 4.441×10^{-16}). **Every figure these checks touch reproduces; two wordings above are corrected, the numbers are not.** The belief snap matches the Step 9(b) prediction to $\sim 10^{-8}$ at all three edges at the truncation/cancellation crossover bracket 10^{-8} ; at the probes’ own 10^{-9} bracket the first edge still holds (4.0×10^{-8} , Analyst A’s “7–8 dp”), but the second and third are 1.2×10^{-7} and 1.7×10^{-7} — floating-point cancellation over a 10^{-9} -wide sliver, not a gap in the prediction. And “ $\mathcal{T}$ bit-identical” holds for U but not for $\mathcal{T}_2$, which moves 6.66×10^{-16} (3

ulps) at the one k_1 where the price signature itself deviates most (4.441×10^{-16}); invariance holds at the map's own root-finder resolution. The analytic weight bound $\min(\kappa/2, 1 - \kappa)^{d+1}$ is **not** curated — no probe computes it; its measured counterpart (the jump entering the adjacent-plan payoff difference undiminished) is. **No label moves and none is licensed.**

A7 / A7' / A7-J Filing sufficiency. On flagged histories $(B^F, Q^F, a = 1)$ identifies the informed component of the selected plan; conditional on it the pooled order-flow residual is pure noise, independent of (v, s, ξ) . L2 uses A7' on path; the weak wording is not sufficient — it permits two (j, s) pairs with different pooled paths, which is exactly L2's first failure case. The two injective forms are distinct. **A7' (on-path composed target)** requires the composed terminal target $s \mapsto b_{j(s)}^*(s)$ to be strictly increasing on the flagged signal region, quantified over every cutoff vector $k \in \Theta$; strictness is required only for flag-capable composed targets, with no backtracking across admissible Voice-plan switches. **A7-J (joint tuple injectivity)** requires the full map $(j, s) \mapsto (B_j^F, Q_j^F, a_j)$ to be injective on the full flagged-pair set, including pairs not selected on path. A7-J is stronger than on-path A7' and is the form consumed by the pre-review P1 proof. Strictness of B^F is neither necessary (it fails at crossing-date jumps on the pro-rata menu) nor sufficient (multi-Voice backtracking). Under A7' the flagged tuple is continuum-valued as a tuple: injectivity forces (B^F, Q^F) to be continuum-valued, while the coordinates may trade the burden. Injectivity plus measurability already gives the measurable inverse (standard Borel spaces); no separate assumption is needed. *Satisfiability is resolved for A7'*: A7' plus a fixed cutoff policy plus $\Omega > 0$ deliver the on-path injective form with an explicit inverse, and the pro-rata single-Voice menu with terminal target strictly increasing on all of $\mathbb{R}$ also satisfies A7-J. A7-J additionally needs strict increase off the Voice region; a target flat below the Voice cutoff breaks A7-J while leaving A7' intact. Failure boundary: a binding stake cap, quantized stakes, a composed target repeating values across Voice-plan switches, $\Omega = 0$, and policy-dependence when the condition is stated at only one equilibrium's cutoffs. A7'-satisfying menus are fully separating on the flagged set, so the burden moves to incentive compatibility rather than away.

A8 Interior crossing. $0 < \Omega(\kappa, \tau, T) < 1$.

A(τ) Threshold chord restriction. The pooled posterior law has the symmetric ternary representation

$$\mathbb{E}[h] = A_0(\kappa)h(0) + A_{1/2}(\kappa)h(\bar{\pi}/2) + A_1(\kappa)h(\bar{\pi}), \quad A'_0 = A'_1 = A'_\kappa, \quad A'_{1/2} = -2A'_\kappa,$$

with maintained orientation $C_h(\bar{\pi}) \leq 0$ and $|C_h|$ weakly increasing in $\bar{\pi}$. Two clauses:

- (τ -i) *The kernel depends on the information set only through the engagement posterior:* $h(\mathcal{I}) = h(\pi(\mathcal{I}))$, so $h(0)$, $h(\bar{\pi}/2)$, $h(\bar{\pi})$ are well defined and κ -free. This is a restriction, not a reading: $h = \pi p$ and p depends on the price as well as on π , so in the model $h = \pi p(\hat{v}, \pi)$ is a function of *two* scalars. The clause says the standalone-value channel and the engagement channel do not co-move inside the pooled cell in a way that moves h at a fixed posterior.
- (τ -ii) *The support and $\bar{\pi}$ are κ -free; only the weights move.* Without the second half, L3's conclusion is **false**: the derivative gains a term first order in $\bar{\pi}$ and the vanishing fails.

Where the bite is: given a κ -invariant three-point support, the derivative restrictions are equivalent to κ -invariance of the pooled block's total mass and of its unnormalised engagement moment, both of which the model delivers at fixed policies. A(τ)'s entire remaining content is the support condition. A one-round ternary-noise market with informed mark $2\bar{z}$ and pre-order engagement share $\frac{1}{2}$ satisfies it; the frozen manuscript's own no-disclosure structure (informed mark $\bar{z}$) does **not** — its pooled law has four atoms, two of which move with κ . **Whether the two-round pooled cell of the timing above satisfies the support condition is OPEN.** Every A(τ)-conditional result — and therefore L4 leg 3 and T1 Part B — inherits that conditionality.

Evidence note added 2026-08-25 (ticket 33). **At the implemented calibration the**

support condition FAILS — and it fails on the support, not on the derivative pattern. The pooled cell’s engagement-posterior law was enumerated exactly (all $4^{H+1} = 4,194,304$ order-flow paths, the same law `pooled_premium` integrates) at **200 nodes**: $\kappa \in \{0.05, \dots, 0.95\} \times$ the five frozen τ percentiles $\times T \in \{1, 2, 5, 10\}$, frozen policies, $H = 10$. Two gates pass first, so the object measured is $A(\tau)$ ’s own: an independent re-enumeration reproduces `pooled_pass` to **0.0 exactly**, and the enumerated mean $\mathbb{E}[\Pi]$ equals the pooled share $\bar{\pi}_{\text{pr}} = \mathbb{P}(a = 1 \mid D = 0)$ to 1.7×10^{-16} . Neither Example A’s $|A'_\kappa| = 0.25$ nor level symmetry is imposed anywhere, and $\bar{\pi}$ is read as the upper support point throughout, per the binding ruling. **20 nodes are degenerate** ($\bar{\pi}_{\text{pr}} = 0$ at $T \in \{1, 2\}$ with τ at the 10th percentile: no engaging atom survives into the pooled cell, the law is the point mass at 0, $M_P = 0$ and $C_h(0) = 0$, so $A(\tau)$ holds vacuously and the node decides nothing). At **all 180 non-degenerate nodes** $A(\tau)$ **fails**; at none does it hold.

- **(τ -ii), support half — FAILS, by some eleven orders of magnitude.** The support carries **23–767 distinct posterior values**, never three (0 of 180 nodes), and there is **no mass at $\bar{\pi}/2$ at any node** ($A_{1/2} \equiv 0$). Between **0.57% and 91.8% of the pooled mass sits off $\{0, \bar{\pi}/2, \bar{\pi}\}$** — 13.9% at the median node ($T = 5$, median τ , $\kappa = 0.55$: 107 atoms, $A_0 = 0.768$, $A_1 = 0.093$). The atoms are not dust: coarsening the cluster tolerance to 10^{-3} still leaves **6–332** of them, and the floor-free law (the $\varepsilon \downarrow 0$ limit of clause (vi) of the equilibrium notion, the law reported here) counts at most 51 atoms fewer than the floored law the package prices. The interior atoms move with κ : the two-sided Hausdorff distance between adjacent- κ support sets reaches **0.4608** — unchanged when restricted to atoms carrying mass $\geq 10^{-6}$ — against $A(\tau)$ ’s predicted $< 10^{-12}$, at **0 of 18** series. This refutes L3 Step 18’s (S1) and (S2) together at this calibration.
- **(τ -ii), $\bar{\pi}$ half — HOLDS.** $\bar{\pi} = 1$ to 1.5×10^{-13} at every non-degenerate node, and κ -free to the same order (18 of 18 series). This is a separate finding and it is not a partial rescue: $\bar{\pi} = 1$ is the **one-round** outcome L3 Step 18 derives from the card’s §4.2 mark structure, and that step’s conjecture that “the two-round timing ... leav[es] the pooled cell with a top atom strictly below 1” is **false at this calibration** — unflagged Voice types still generate fully revealing order flows. $\bar{\pi} \in \{0, 1\}$ across the whole grid and never interior, so L3’s small- $\bar{\pi}$ corollary has no instance here either.
- **Derivative pattern — FAILS, and independently of the support.** $A'_0 = A'_1$ holds at **0 of 180 nodes**: $|A'_0 - A'_1| \in [0.041, 2.306]$ against a predicted $< 10^{-10}$, with $A'_0 \in [-2.146, 2.374]$ against $A'_1 \in [-0.014, 0.429]$ — an order of magnitude apart in level, and both change sign over the grid, which independently corroborates that A'_κ carries no sign (audit finding 2). $A'_{1/2} = -2A'_\kappa$ also fails at all 180, but with $A_{1/2} \equiv 0$ that residual is exactly $2|A'_0|$ and is recorded as **inherited** — a restatement of the support failure, not a second piece of evidence.
- **Chord identity — FAILS.** $|\mathcal{S}_P - \Delta_m |A'_\kappa| |C_h(\bar{\pi})|$, with A'_κ **recovered** from the enumerated weights and $\bar{\pi}$ the **actual** upper support point, is **0.0013–0.0717 (up to 7.17 premium pp)** against $< 10^{-10}$, at 0 of 180 nodes and on the most favourable of three kernel conventions. Recovered $|A'_\kappa| \in [0.042, 2.374]$; the value the identity would *require* is $[0.00023, 0.392]$, **disjoint** from block 3’s implied $[0.997, 1.158]$ — which is a different object (mean absolute slope over the κ grid, and the level-symmetric $\bar{\pi} = 2\bar{\pi}_{\text{pr}}$), and the distance between the two measures what the level-symmetry assumption was doing.
- **(τ -i), reported as a diagnostic and not part of the verdict.** Within a Π -cluster (Π constant to 10^{-12}) the enumerated entry probability still spreads by up to **0.085**, and h by up to **0.018** mass-weighted. The kernel does not reach the information set only through the posterior at this calibration either.

What this changes, and what it does not. NUMERICAL-class applicability evidence

at one calibration; **no label moves**, and none is licensed — $A(\tau)$ is an assumption, not a labelled claim. L3, L4 leg 3 and T1 Part B stay **PROVED as conditionals** with their proofs untouched; what is now on record is that their antecedent is **not satisfied by the implemented pooled cell at this calibration**, so at this calibration those legs say nothing about the implemented cell. The question stated above stays open as a question about $A(\tau)$'s **domain** — a different menu, a different H , or a different calibration could still satisfy (S1)–(S2) — and the two prior “failures” remain misformulated tests; this is the first test that measures $A(\tau)$'s own object. Coverage caveats carried forward: the 18 non-degenerate series are only **6 distinct pooled cells** ($T = 1$ and $T = 2$ induce identical D -partitions at every τ ; $T = 5$ joins them at the three highest τ percentiles and repeats itself at the two lowest; all five $T = 10$ quantiles coincide), and all six fail; the 50 $T = 10$ nodes sit at $\Omega = 0.000681$, below MIN_CELL_MASS (HANDOFF_sign.md §8.1). Script and record: `quality_reports/fixes/t2_atau_support_check.py` → `t2_atau_support_check.json` (200 nodes, 920 pooled enumerations, 1002 s; top-level verdict field **FAILS at calibration**).

A(br) **Chord-sensitivity bridge**. Consumed by L4 leg 3 and T1 Part B, and by nothing else. For two compared thresholds $\tau' < \tau$ at fixed policies and a common κ :

- (br-i) *Representation at both policies*. $A(\tau)$'s symmetric ternary representation holds for the pooled class under τ and under τ' , with endpoints $\bar{\pi}(\tau)$, $\bar{\pi}(\tau')$ and coefficients $A'_\kappa(\tau)$, $A'_\kappa(\tau')$.
- (br-ii) *κ -localisation*. At fixed policies all κ -dependence of M_P sits in the $A(\tau)$ weights, so $\partial_\kappa M_P = \Delta_m A'_\kappa C_h(\bar{\pi})$ exactly, with no composition-through- κ remainder. Written against the honest reading $h = \pi p(\hat{v}, \pi)$, this repairs an ambiguity rather than adding a fourth independent restriction; the trailing “hence” is derivable, not assumed.
- (br-iii) *Coefficient stability across the threshold margin*. $|A'_\kappa(\tau')| \leq |A'_\kappa(\tau)|$; weakest sufficient form is equality — reclassification changes *which* histories are pooled, not the κ -responsiveness of the pooled weights. *This is the clause with the least justification behind it; it is the one to attack first.*
- (br-iv) *Endpoint linkage*. $\bar{\pi}$ is $A(\tau)$'s chord endpoint — the upper support point — and is a weakly increasing function of the pooled prior engagement share $\bar{\pi}_{\text{pr}} = \mathbb{P}(a = 1 \mid D = 0)$, the same function at τ and τ' . The identity branch $\bar{\pi} = \bar{\pi}_{\text{pr}}$ is excluded as degenerate.
- (br-v) *Comparability of the chord functional across thresholds*. $C_h(\cdot)$ — and the kernel h it is built from — are the same functions of the posterior at both compared thresholds. Without it, leg 3 compares $|C_h|$ across two different functionals and the comparison is meaningless. Required independently by three agents, and confirmed not implied by (br-i)–(br-iv).

Sharpening on file, recorded not assumed: $\bar{\pi} = \bar{\pi}_{\text{pr}}/\rho$ with $\rho := \frac{1}{2}A_{1/2} + A_1$ provably κ -free, so (br-iv) is equivalent to $\rho(\tau')/\rho(\tau) \geq \bar{\pi}_{\text{pr}}(\tau')/\bar{\pi}_{\text{pr}}(\tau)$. Under the level-symmetric reading $\rho = \frac{1}{2}$ and $\bar{\pi} = 2\bar{\pi}_{\text{pr}}$, forcing $\bar{\pi}_{\text{pr}} \leq 1/2$ — an inherited restriction on $A(\tau)$'s domain that is not resolved.

AGE GE differentiability and contraction. On a candidate region $\mathcal{R}$ the outer map is twice continuously differentiable, $L_{\mathcal{R}} < 1$, and the sign of the equilibrium liquidity derivative is constant on $\mathcal{R}$.

7 The result ledger

All eight results now carry two-pass evidence — an adversarial proof-read PASS and an independent statements-only re-derivation PASS, by different agents — and are labelled PROVED with their conditionality written into the label. (C1 moved on 2026-08-22, after its own proof-read, re-derivation, and the independent re-run of every check script — ALL REPRODUCE, `quality_reports/fixes/t2_rerun_verify_note.md`; **P1 was demoted on 2026-08-23** and

restored on 2026-08-25 on a fresh pair of passes over the amended statement.) The protocol (the *Labels* subsection below) requires both passes, by different agents, before a label moves. That gate is satisfied for D1, L1, L2, L3, L4, T1, C1 and — as of 2026-08-25 — P1 again. **P1’s 2026-08-21 chain never satisfied it:** the proof consumed A7-J while the row and re-derivation carried A7’, so the two passes covered two different statements, which is what the 2026-08-23 demotion turned on. The pair on file now is `threads/2026-08-25_P1_proofread_retry.md` (0 FAIL) and `rederive/P1_rederivation_2026-08-25.md` (PASS-WITH-CHANGES, changes folded into the row). Labels are transcribed from `LABEL_LEDGER.md`; every statement below is the *amended* statement — the hypothesis sets are named in full and descriptively, and **no statement was weakened silently**: each difference from the pre-regeneration row is traceable to the named finding beside it.

ID	Statement, with its full hypothesis set	Label
D1	Under A1, A2’, A4, A5, the primitive/plan table restrictions, the cutoff selection map, and the pricing conventions ($P_{-1}^P = \mathbb{E}[Y]$; $P_{ND} = P_{f-}^P$), plus two hypotheses this regeneration wrote onto the card — Borel regularity of $s \mapsto B_j(s, d)$ for every plan including Exit (needed only for part (c), since pooled pricing integrates over all types) and a content for $\mathcal{I}_H$ — the indicator $D = \mathbf{1}\{a = 1, c(\tau) + T \leq H\}$ is measurable and maps every control-node history into exactly one cell; for every Voice plan $f_j \leq H \iff B_j(s, H - T) \geq \tau$; and each flagged history yields B^F, R_d, R, J with $P^F - P_{c-}^P = R + J$.	PROVED
L1	Under D1, the premium definitions, A5 (which pins <i>the</i> version of $\mathbb{E}[Y \mathcal{I}]$), A2’ with Δ_m finite, and A1 (one probability space): whenever $0 < \Omega < 1$, $\Delta^{\text{act}} = \Omega M_F + (1 - \Omega)M_P$; at $\Omega = 1$ it degenerates to $\Delta^{\text{act}} = M_F$ and at $\Omega = 0$ to $\Delta^{\text{act}} = M_P$, the null-cell average being <i>undefined rather than imputed</i> — proved as a non-identification statement, not asserted.	PROVED
L2	At fixed cutoff and execution policies, under A1; A2’ with the primitive/plan table restrictions; A4; A5; A7’ in its on-path injective form, consumed almost surely on the flagged set; D1; the no-feedback timing; and $\Omega > 0$ — together with an explicit bidder-entry rule carried as bookkeeping: $(B^F, Q^F, a=1)$ makes the pre-filing pooled history conditionally independent of (v, s, ξ) on the flagged set, so the flagged posterior, price, entry probability and M_F are invariant to κ . (“Almost surely” is the only coherent reading once B^F is continuum-valued, since then no individual flagged tuple has positive probability.)	PROVED
L3	Under $A(\tau)$ — including $(\tau\text{-i})$ kernel-through-posterior and $(\tau\text{-ii})$ κ -free support <i>and</i> κ -free $\bar{\pi}$ — plus $h(0) = 0$; κ -free pooled mass and engagement moment at fixed policies; D1 by statement; minimal regularity (h continuous on $[0, \bar{\pi}]$, twice differentiable on the open interval — Darboux does the rest, no continuity of h''); for the small- $\bar{\pi}$ corollary only, a second-order Peano expansion of h at $0+$ and one and the same kernel along the shrinking family; and, for the seam where L4 consumes L3, $ A'_\kappa $ bounded <i>uniformly in $\bar{\pi}$</i> . Then $\partial_\kappa \mathbb{E}_\kappa[h] = A'_\kappa C_h(\bar{\pi})$ exactly; $C_h(\bar{\pi}) = \frac{1}{4}\bar{\pi}^2 h''(\zeta)$ for some $\zeta \in (0, \bar{\pi})$ — an identity, not an approximation; and $C_h = \frac{1}{4}h''(0)\bar{\pi}^2 + o(\bar{\pi}^2)$, so the interior motion vanishes at rate $\bar{\pi}^2$. An “if”, never an “iff” ($A'_\kappa = 0$ also kills the motion). Whether the two-round pooled cell satisfies $A(\tau)$ ’s support condition is OPEN .	PROVED under $A(\tau)$

ID	Statement, with its full hypothesis set	Label
L4	At fixed policies, for $b_0 < \tau' < \tau$ at a common window T and common κ, with $\Omega(\tau', T) < 1$: (leg 1, unconditional) $\mathcal{C}_F(\tau, T) \subseteq \mathcal{C}_F(\tau', T)$ with every newly flagged history generated by a Voice plan, hence $\Omega(\tau', T) \geq \Omega(\tau, T)$; (leg 2, unconditional) the pooled engagement share falls, $\bar{\pi}_{\text{pr}}(\tau') \leq \bar{\pi}_{\text{pr}}(\tau)$, with an exact identity for the gap; (leg 3, under A(br)) $\mathcal{S}_P(\tau', T) \leq \mathcal{S}_P(\tau, T)$, with equality whenever $C_h(\bar{\pi}(\tau)) = 0$. Legs 1–2 need only D1’s clock equivalence, the no-feedback timing, fixed policies, $b_0 < \tau' < \tau$ imposed at both thresholds, A1, A4, $D = 1 \Rightarrow a = 1$, and $\Omega(\tau') < 1$ — the “nestedness” clauses are <i>conclusions, not hypotheses</i>. Leg 3 additionally needs L3 by statement, A(τ)’s maintained magnitude monotonicity of C_h (the sign half $C_h \leq 0$ is never used at this leg) and A(br) (br-i)–(br-v).	PROVED (legs 1–2 outright; leg 3 under A(br))
P1	Under A1, A2', A3, A4, A6, A7-J (joint tuple injectivity — the joint (j, s) form of A7 above, on the whole flagged-pair set $\{(j, s) : D_j = 1\}$ <i>including pairs no cutoff vector selects</i>; strictly stronger than the on-path A7', and the form the proof consumes where it pins <i>off-path</i> flagged beliefs. Amended from A7' 2026-08-25: the pre-review row carried the on-path form while proofs/P1_proof.md h.7 consumed the joint form, so the two 2026-08-21 passes covered two different statements), D1 by statement <i>with its own hypotheses travelling</i>, the no-feedback timing read with the flag-terminates-the-pooled-round clause, the definitional round-2 action-set hypothesis (the flagged-round action set is the plan-generated set $\{Q_{j'}^F(s)\}$ over menu elements agreeing with j on everything already played — <i>not</i> a closure condition; the closure form is jointly unsatisfiable with finiteness by cardinality), continuation-cost equivalence on that same set (the proof’s h.16, added 2026-08-25: menu elements sharing j’s pooled path up to $f_j(s)$ with $a_{j'} = a_j$ carry the same engagement cost, $C_{j'}(s) = C_j(s)$). Trivially true on any single-Voice menu, where that set is a singleton. What it buys, under the plan-completion reading of the $C_j(s)$ timing convention below — under the sunk reading the continuation is constant on the deviation set with no clause at all and h.16 is not consumed, so the hypothesis is listed because the row does not commit to a reading, and it is what makes the conclusion hold under both: on that set the flagged price does not move and the order cancels, so the engagement cost is the only thing that can differ between staying and deviating — and at a flagged pair the cutoff vector does not select there is no date-0 optimality to fall back on, so without this clause the deviator takes the class member with the smallest cost and item (ii) of the equilibrium notion fails at that node. Live only on menus with two or more Voice plans sharing a pooled path),	PROVED

ID	Statement, with its full hypothesis set	Label
	$m_0 \geq 0$, the card's §4.3 blockholder-objective definition U_j (whose $-a_j C_j(s)$ display proofs/P1_proof.md h.14 now carries verbatim, which is what the row's "displayed there in full" asserts; timing convention, stated here because the card's §4.3 does not date $C_j(s)$: the engagement cost may be booked either on completing the plan or as sunk once the filing has landed — the two give the same round-2 comparison on the round-2 deviation set, which is what the continuation-cost clause above buys, so the result does not depend on the choice), and the card's §4.1–§4.3 table restrictions the argument consumes — in particular §4.3's Y row with the price convention $P(\mathcal{I}) = \mathbb{E}[Y \mid \mathcal{I}]$ and the entry row for $p(\mathcal{I})$; §4.2's Borel-regularity clause for <i>every</i> plan including Exit (needed directly, not via D1, whose conclusion is measurability of D and the cell map); §4.2's $D = 1 \Rightarrow a = 1$, the $c/f/B^F/Q^F/b^*$ definitions and $\partial_s B_j \geq 0$ for Voice; and §4.1's distributional forms with $\Delta_m > 0$:	

ID	Statement, with its full hypothesis set	Label
	at every $\kappa \in [0, 1]$, a cutoff PBE over complete contingent plans exists — $k^* \in \Theta$ with $k^* = \mathcal{T}(k^*; \vartheta)$, prices at their inner fixed points, Bayes-consistent on-path beliefs, off-path beliefs as limits of one full-support perturbation family over plans — fixed once and used to define the price system at every $k \in \Theta$, not only at k^*, since the deviation payoffs that define $\mathcal{T}$ read off-path pooled histories — at every pooled history reachable <i>with positive probability</i> under some plan profile (at the boundary values $\kappa \in \{0, 1\}$ the card's §4.1 noise support degenerates to $\{0\}$ and to $\{\pm \bar{z}\}$ respectively; a pooled history needing a mark outside it is null under <i>every</i> profile, so it is off nature's path rather than off the players', carries no clause-(vi) requirement, and is read by no step. This is the extension route, not the restriction one: no cut to $\kappa \in [0, 1)$ is taken, and the pre-repair claim of a belief at <i>every</i> pooled history — false at $\kappa = 1$ — is withdrawn); flagged-tuple beliefs supplied by A7-J at every tuple in the image of the flagged-pair map $(j, s) \mapsto (B_j^F, Q_j^F, a_j)$ — on path and off, since the image includes tuples generated by pairs the cutoff vector does not select — as the point mass at the unique generating pair, which is a version of the conditional law at every image tuple (the signal is continuous, so a version is what a conditional law is; any a.e.-equal version serves clauses (iii)/(vi) of the equilibrium notion equally) and is the version this equilibrium selects, with no tuple outside that image arising because the round-2 action-set hypothesis leaves no off-menu order to produce one; the card's §4.3 entry rule; and a sequentially optimal flagged component at every flagged pair (j, s), whether or not the cutoff vector selects it — the flagged price is invariant across the round-2 deviation set (A7-J pins the belief at the same s and $\pi = 1$), so the order cancels out of the continuation and the continuation-cost clause makes what remains constant. A5 is not assumed: its existence and uniqueness content is derived from $m_0 \geq 0$, its continuity content from the same scalar reduction, and its measurable-selection content from A7-J plus the card's §4.2 Borel clause (see A5). A6 is read as asserting that $\mathcal{T}$ — under a named tie-break-and-corner selection, without which a correspondence cannot be called continuous — is a well-defined single-valued continuous self-map of Θ, with Θ nonempty per the card's §4.5. At any such equilibrium at which A8 holds, both cells carry strictly positive probability and are on path; for A8's restatement as a single signal threshold add H-ord (Voice stake monotonicity across plans — the writer's h.13, renamed here to avoid collision with the objective row) and the upper-set engagement-flag hypothesis. A8 is a condition <i>on the fixed point exhibited</i>: existence of an equilibrium <i>at which</i> A8 holds is not claimed. Uniqueness is not claimed.	

ID	Statement, with its full hypothesis set	Label
	Evidence. Statement <code>threads/thread1_turn1_answer.md</code>; proof <code>proofs/P1_proof.md</code> (repairs applied through P1-R35, ticket 35 rounds 1–2, close-out and confirm-pass sweep); proof-read PASS 2026-08-25 <code>threads/2026-08-25_P1_proofread_retry.md</code> (0 FAIL; 3 REPAIRs + 4 OBSERVATIONs, all applied; the reader verified the Step 12 lemma part by part on the merits and records that his own round-1 FAIL witness is refuted — round-1 FAIL and the sanctioned repair round at <code>threads/2026-08-25_P1_proofread_round1.md</code>); re-derivation PASS-WITH-CHANGES 2026-08-25 <code>rederive/P1_rederivation_2026-08-25.md</code> (fresh agent, card row alone; changes 1–5 folded into this statement cell — the card’s §4.1–§4.3 citation block, D1’s hypotheses travelling with the three-part A5 sentence, the one-family/every-k/positive-probability off-path clause, A6’s tie-break-and-corner reading, the C_j timing convention; change 6 withheld for Austin — a proposed OPEN item on whether A6’s continuity of $\mathcal{T}$ is satisfiable at the collapsed cutoff vectors the equilibrium notion admits; ruled 2026-08-27: answered rather than filed OPEN — item 4 of the open list under <i>What this note does not claim</i> and the A6/A3 evidence notes above carry the panel record, no label moved). The 2026-08-21 chain is retained below and did not satisfy the gate for the recorded statement: proof-read PASS 2026-08-21 <code>threads/2026-08-21_batch1_proofread_audit.md</code> §4 (0 FAIL; P1-R1...R8; inner fixed point executed on 20k random draws — 0 multiplicity, 0 sign failures) and re-derivation PASS 2026-08-21 (PROVED-WITH-CHANGES) <code>rederive/P1_rederivation.md</code> (changes C1–C8) covered two different statements — the proof’s h.7 consumed the joint injective form of A7 while the row and re-derivation carried the on-path form — which is what the 2026-08-23 demotion turned on, together with Step 12’s missing continuation-cost clause and the false positivity claim at $\kappa = 1$; all three are repaired and independently reproduced by the 2026-08-25 re-derivation. Numerical status, stated honestly and separately from the label (ticket 34, quality_reports/fixes/t2_p1_fournode_recheck.json): the four sweep-unresolved nodes $(\kappa \in \{0.15, 0.85\} \times (\tau, T) \in \{(0.05, 5), (0.075, 1)\})$ remain STILL UNRESOLVED after 30 seeds each — best payoff-scale residual 3.1×10^{-4}–1.5×10^{-3} against a 10^{-9} criterion, best cutoff-scale residual 10^{-14}–10^{-11}; the A3 and A6 proxies pass at every achieving seed. UNCHECKED: existence at those four nodes is neither claimed nor denied by this evidence, and the label rests on the proof plus the two 2026-08-25 passes, not on the grid.	

ID	Statement, with its full hypothesis set	Label
T1	At fixed plan and cutoff policies, with $0 < \Omega < 1$ and $S_P > 0$: (A) $\mathcal{S} = (1 - \Omega)\mathcal{S}_P$ exactly, and the same factorisation holds for the total-variation aggregate of Δ^{act} over any κ -grid with no differentiability required; (B) threshold tightening attenuates — $\mathcal{S}(\tau')/\mathcal{S}(\tau) = W_\tau C_\tau \leq 1$ because <i>both</i> ratios lie in $[0, 1]$, no dominance condition needed; (C) window tightening attenuates iff $W_T C_T \leq 1$, where $W_T \leq 1$ is <i>proved</i> (from D1’s clock equivalence and the monotone Voice stake path) and C_T is <i>unsigned</i> — “equivalently $\partial_{r_T} \mathcal{S}_P / \mathcal{S}_P \leq \Omega_{r_T} / (1 - \Omega)$ ” holds on average along the tightening path (integrated over $[-T, -T']$), exactly in the infinitesimal limit, and is false read pointwise . Hypotheses: fixed policies; A8 at each compared policy; $S_P > 0$; L1; L2 (its own hypotheses travelling); D1; PE-Ω ($\partial_\kappa \Omega = 0$ at fixed policies — derivable, not assumed, and it fails in GE, which is C1’s term); κ-differentiability of M_P (no card hypothesis supplies this — carried in-proof); $A(\tau)$ at both compared policies with the $\bar{\pi}$ ruling; L3; $A(\text{br})$ (br-i)–(br-v) at the threshold pair; L4; the no-feedback timing; a smooth window interpolation for the local form; threshold-side smoothness (confirmed non-load-bearing). No unconditional window sign is claimed.	PROVED at fixed policies
C1	On a named region where $L_{\mathcal{R}} < 1$ and $g_r^{PE} > \mathcal{B}_r^{GE}$, the fixed-policy attenuation sign survives in equilibrium. The dominance-and-contraction implication is two-pass complete — an adversarial proof-read PASS (0 FAIL, 13 repairs, 7 observations) and an independent statements-only re-derivation PROVED-WITH-CHANGES (a norm convention and two-sided openness of $\mathcal{R}_r$ added; at $J - 1 \geq 2$ the bare $ \cdot $ in $\mathcal{B}_r^{GE}$ is not well defined and a mismatched reading makes part (B) false) — and the executed 80-node run returns 18 pointwise dominance-and-contraction nodes , slack η_r from 0.0595 (min) through 0.3467 (median) to 1.7227 (max), $L_{\mathcal{R}} \leq 0.5008$ at every node. These nodes verify the pointwise inequalities and supporting diagnostics only; they do not verify the full C1 antecedent or a named nonempty region.	PROVED

Remaining aspiration, not a claim: promotion from pointwise dominance-and-contraction nodes to a named nonempty region, which would require the D8 ε -ball and integral-control pattern.

Labels

PROVED — a complete proof, independently re-derived and proof-read. **NUMERICAL** — verified on a grid by an executed, committed check script with committed output. **ESTIMATED** — an empirical estimate with a standard error and a stated design. **CONJECTURE** — everything else, including anything whose proof is deferred. A **dominance-and-contraction node** is not a fifth label: it records pointwise $L_{\mathcal{R}} < 1$ and $\eta_r > 0$ with supporting diagnostics, not verification of the full C1 antecedent. A named-region promotion is not claimed. Labels are never weakened by editing; only an executed check or an independent re-derivation may move one.

7.1 Three known facts on file

These are facts the record carries, not claims the model makes. They are the reason the labels above are written with their conditions attached.

1. O-1 is a disclosure-regime analogy, not a window-margin test. This is a known fact on file, not a claim the model note makes: O-1 compares the public buy flagged versus pooled at fixed policies in the static repo model. Its ratios

$$1.06397 / 1.18373 / 1.13631 / 0.37798 \quad \text{at} \quad \Omega = 0.037252 / 0.128950 / 0.285804 / 0.50$$

are regime-comparison composition outcomes, not $W_T C_T$ and not a window pair. The analogy is useful because it shows that a composition factor can exceed one, motivating T1's genuine window-margin iff. The genuine window-margin record is `t2_t1_check` block 4: $W_T C_T < 1$ at every checked node at this calibration. The O-1 boundary $\Omega^* \approx 0.3428$ remains a disclosure-regime boundary, not a window boundary. Provenance: this is the referee's O-1 experiment run on the *repo* model (the frozen manuscript's static structure, fixed cutoffs, partial equilibrium), re-executed in `quality_reports/fixes/t1_o1_rerun_check.py` with every committed number reproducing to the last printed digit; it is not the two-round build.

2. Whether the implemented two-round pooled cell satisfies $A(\tau)$: at the implemented calibration it FAILS. Two executed diagnostics are informative but neither decides the support condition, and both are retained without smoothing.

- `t2_l2_check.json`, check `l2_placebo_M_P_sign_A_tau`: **FAIL** as a sign placebo. The enumerated pooled M_P is *hump-shaped* in κ — 10 of 18 increments positive, one sign change, peak near $\kappa = 0.55$. But $A(\tau)$'s maintained $C_h(\bar{\pi}) \leq 0$ does not sign A'_κ ; the card's own Example A has $A'_\kappa = -1/4$. The demanded sign is therefore not implied by $A(\tau)$, and the placebo is misformulated rather than a refutation of the support representation.
- `t2_t1_check.json`, check `t1_block3_chord_magnitude`: **FAIL** for the hard-coded Example-A coefficient. The residual $|\mathcal{S}_P - \Delta_m|A'_\kappa||C_h(\bar{\pi})|$ is 0.006279 (0.6279 premium percentage points) against a predicted 10^{-10} — the implied coefficient is $|A'_\kappa| \in [0.997, 1.158]$, while the check hard-coded 0.25. The ratio $|C_h|/\bar{\pi}^2$ remains constant to 0.48% between the two smallest $\bar{\pi}$ nodes, well inside the 5% criterion. This rejects the hard-coded calibration, not $A(\tau)$'s general support representation.

The decisive support-enumeration check is ticket 33, and it has landed: at the implemented calibration the support condition **FAILS**, at every non-degenerate node, while the $\bar{\pi}$ half of $(\tau\text{-ii})$ holds. The measurement is Section 6's $A(\tau)$ evidence note, and it moved no label. The two diagnostics above remain test-design artifacts, not a wiring error or a theorem failure — and neither of them is the test that decided it. What stays open is $A(\tau)$'s **domain**: a different menu, a different H or a different calibration could still satisfy the support condition (item 1 of the open list under *What this note does not claim*).

3. L4's sign predictions hold numerically with zero violations. `t2_l4_check.json` reports 10 checks, 0 failures. Across all 16 tightening steps and both windows: Ω rises at every step (0 violations, minimum increment +0.02263); $\bar{\pi}_{\text{pr}}$ falls at every step (0 violations, largest increment -0.018766); and $\mathcal{S}_P$ falls at every step (0 violations, largest increment -9.47×10^{-5}) — and, as the JSON notes, a single violation would be a failed hypothesis rather than sampling error, because nothing in that computation is sampled. Ω and $\bar{\pi}_{\text{pr}}$ are flat in κ to exactly zero across 71 grid points, which is what the no-feedback timing predicts.

What the pair of facts means. The chord mechanism is **hypothesis-bound**: it needs $A(\tau)$, whose applicability at the implemented pooled cell is now answered — it **fails** at this calibration, no label moved, so L3, L4 leg 3 and T1 Part B stay **PROVED** as conditionals that say nothing about the implemented cell here; whether $A(\tau)$ has a domain that fits stays open. The threshold phenomenon is **not** hypothesis-bound: the reclassification legs of L4 are unconditional, and they hold numerically without exception.

8 What this note does not claim

A global window-margin attenuation sign; κ -invariance of J ; equilibrium uniqueness; a named nonempty GE region *as a theorem* (the 18 pointwise dominance-and-contraction nodes are an executed record, not verification of the full C1 antecedent); endogenous filing before the deadline; noisy or partially revealing flagged-round trading; continuous-time execution; welfare or optimal rule design; that the frozen manuscript's hump result survives; that the prior calibration ($\Omega \approx 0.037$) is economically meaningful; any empirical value for ω_a .

Three items are explicitly open (*a fourth, item 4, added 2026-08-27*):

1. **Whether the two-round pooled cell satisfies $A(\tau)$ — OPEN.** L3 proves the representation's entire remaining bite is the support condition, exhibits a one-round market that satisfies it and the frozen manuscript's own no-disclosure structure that does not, and names the weakest sufficient conditions. L3, L4 leg 3 and T1 Part B are all conditional on $A(\tau)$, so this is the largest single conditionality the ledger carries — and the executed diagnostics of fact 2 above do not decide it at baseline; the support-enumeration check is ticket 33.
2. **Whether an equilibrium in which the blockholder chooses the fully separating plan exists on a given calibration — OPEN, and P1-adjacent.** A7' menus are fully separating on the flagged set, so the burden did not disappear when A7 satisfiability was resolved; it *moved* to incentive compatibility, which P1 does not settle.
3. **O-1 is a disclosure-regime analogy, not a window-margin test** (fact 1 above). The genuine window-margin record is `t2_t1_check` block 4, with $W_T C_T < 1$ at every checked node at this calibration.
4. **Whether A6's continuity of $\mathcal{T}$ holds for the declared construction — ANSWERED IN SUBSTANCE 2026-08-27 (panel evidence; locus corrected; the open remainder is scoped below).** The re-derivation's withheld change 6 proposed filing this as OPEN at collapsed cutoff vectors. Ruled on Austin's 2026-08-27 authorization after a two-agent opposed-brief panel (substantiate vs defuse) with orchestrator adjudication — the briefs **converged on every load-bearing point** and **cross-replicated the decisive measurement** (the same three $\mathcal{T}_2$ jumps, independent scripts, 3 s.f. agreement). *Answered:* the discontinuity mechanism is real and reaches $\mathcal{T}$ with non-vanishing weight (the vanishing-mass defusal is refuted, twice independently); the locus is the **cell-edge hyperplanes $\cup$ sole-generator collapse faces**, not the collapsed vectors as such; the implemented menu's Hold-collapse face is measured clean, while A6 read with the Θ Steps 13–14 construct **fails at the implemented calibration** (measured $\mathcal{T}_2$ jumps 6.3×10^{-3} – 2.83×10^{-2} at the baseline node, up to 0.16 at $\kappa = 0.15$, against a 10^{-10} tolerance); and on $J \geq 3$ menus where a middle plan owns a reachable exclusive pooled history, **no k -independent perturbation family** restores continuity at more than one collapse-face point (continuum-face lemma; panel derivation, not gate-checked), so on that menu class P1 asserts nothing — the $A(\tau)$ pattern. *Still OPEN:* (a) whether a constructive Θ — or the Step 18 t -constrained Kakutani route already on file, which removes h.6's continuity half — replaces h.6 in the statement; the repair is identified, not executed; (b) the complementary menu class (every middle plan's histories shared with a survivor — the implemented menu is one such for its collapse face), where the collapse-face clause may be satisfiable; (c) nonexistence, which is neither claimed nor shown anywhere — 23/27 sweep nodes converge and two fixed points survive even at the worst probed node. *Separately recorded, deliberately not folded in* (different hypothesis, different route): A3's own failure at this calibration and the candidate ticket-34 account — the A3 evidence note in Section 6. **No label moves on any of this;** P1's label rests on the proof and the two 2026-08-25 passes and is untouched. Evidence: the A6 note in Section 6; `threads/2026-08-27_A6_panel_substantiate.md` / `threads/2026-08-27_A6_panel_defuse.md`;

quality_reports/fixes/a6_panel_probes_2026-08-27/. Follow-ups: the decisive-probe curation landed 2026-08-28 (Section 6’s A6 note, curation note; the `t2_a6_*` checks); gate-check the continuum-face lemma only if it is ever promoted.

9 What v4 takes from draft_v2, what it simplifies, what it drops

The frozen manuscript `draft_v2.tex` is a record the supervisor has seen; it is not edited and it is not cross-referenced here. The table below describes its objects rather than citing them.

draft_v2 object	Verdict	What carries over, what changes
Four-action structure: Exit, Hold, Quiet Voice, Public Voice, with weakly ordered cutoffs (k_1, k_0, k_D) ; and the dominance lemma making the four-action menu a result rather than an assumption (engaged exit and silent buy are never chosen).	REUSE	Becomes the <i>special case</i> of the v4 ordered plan menu: $\mathcal{J}$ finite and ordered with $J \geq 3$, cutoff vector $k = (k_1, \dots, k_{J-1})$, which maps to (k_1, k_0, k_D) exactly when the menu is the four named actions. What changes: menu elements are now <i>complete contingent plans</i> over the calendar $d = 0, \dots, H$, not single one-shot (q, a) pairs, and A3’s single crossing is stated for adjacent plans. The dominance argument survives as the reason the menu is ordered and finite.
Brouwer existence architecture: boundedness of best responses, the self-map and order-preservation lemma, and existence of a monotone equilibrium by Brouwer on a nonempty compact convex Θ .	REUSE	Reused whole as P1’s architecture: Θ compact, ordered, convex; $\mathcal{T}$ continuous $\Theta \rightarrow \Theta$; $k^* = \mathcal{T}(k^*; \vartheta)$; collapse faces permitted (the manuscript’s own baseline collapses Hold). What changes: the fixed point ranges over plans; off-path beliefs are limits of full-support perturbations <i>over plans</i> ; and the inner pricing fixed point is no longer an assumption (see the A5 row).
The (C*) chord condition $\mathcal{C}(\bar{\pi}) = h(0) - 2h(\bar{\pi}/2) + h(\bar{\pi})$, and the lemma signing the posterior/pricing channel under it.	REUSE	Is v4’s $C_h(\bar{\pi})$, same expression, same maintained orientation $C_h \leq 0$ with $ C_h $ weakly increasing in $\bar{\pi}$. What changes: in <code>draft_v2</code> it selects hump versus trough in κ ; in v4 it is the magnitude that drives pooled sensitivity through L3’s exact identity $\partial_\kappa \mathbb{E}_\kappa[h] = A'_\kappa C_h(\bar{\pi})$, and the $A(\tau)$ representation that licenses that identity is stated as a restriction whose support condition is open.
The tender-game microfoundation of the premium wedge (the disagreement-node continuation deriving the appropriability coefficient $\lambda = 1 - q(1 - \gamma)\psi$ and the resulting wedge).	REUSE	Kept as the <i>input</i> that supplies (m_0, m_1) . v4 takes m_0, m_1 as primitives with $m_1 > m_0$ and $m_0 \geq 0$ and does not re-derive them; the tender game is where those numbers come from and why $\Delta_m > 0$ is not free. Its symbols are reserved and unavailable: bare λ is the appropriability coefficient, ψ is pivotality.

draft_v2 object	Verdict	What carries over, what changes
The GE dominance-check pattern: contraction modulus L along the path, an inversion-free Neumann bound $\bar{B}$, pointwise dominance off a ball plus a ball-integral condition, a certified interval, and a counterexample showing the global claim false.	REUSE	Is C1's precedent. $\mathcal{B}_r^{GE}$ is the <i>cross-derivative</i> analogue of $\bar{B}$; $\eta_r = g_r^{PE} - \mathcal{B}_r^{GE}$ is the slack; the executed record names pointwise dominance-and-contraction nodes, not a verified region, and a named-region promotion remains open. What changes: the object bounded is a cross-derivative in (κ, r) rather than a single-argument profile, and the strictness coordinate $r_T = -T$ is not differentiable because the window is an integer.
Stake-triggered disclosure as the market's partition, with the flag a function of the order ($D = \mathbf{1}\{q = +1\}$: a buy cannot be concealed).	REUSE	The idea that the rule <i>is</i> the partition is the whole v4 position. What changes is the trigger: $D = \mathbf{1}\{a = 1, c(\tau) + T \leq H\}$, a stake threshold plus a filing window on a business-day calendar, with D1's clock equivalence $f \leq H \iff B(s, H - T) \geq \tau$. The window T becomes a genuine primitive instead of a flag that is either on or off.
Assumption (A2)'s flat boundedness of prices and payoffs.	SIMPLIFY	Replaced by A2': finiteness clauses kept verbatim; boundedness weakened to local boundedness plus integrability $\mathbb{E}[\max_j U_j] < \infty$. The flat bound was <i>false</i> in this model (v Gaussian, flagged region unbounded in s), and no proof ever used it.
Assumption (A5)'s inner pricing regularity (existence, uniqueness and continuity of the price fixed point, assumed, with a $\delta/\sigma_\xi < 1/\phi(0)$ style regularity condition).	SIMPLIFY	Demoted to a theorem under $m_0 \geq 0$: the pricing map is strictly decreasing in P wherever $\bar{m} \geq 0$ and crosses the identity exactly once. A5 survives only as its continuity clause, read as a measurably selected family because the flagged information sets are continuum-indexed. Three independent confirmations, including executed counterexamples that produce zero roots and three roots once $m_0 < 0$.
The discount factor δ inside the pricing fixed point $P = \delta[(1-p)\hat{V} + p(P + \bar{m})]$, and the sensitivity analysis around it.	DROP	v4 prices are undiscounted conditional expectations $P(\mathcal{I}) = \mathbb{E}[Y \mid \mathcal{I}]$. The regularity work δ was doing is done instead by $m_0 \geq 0$.
The one-round pooled posterior: $\pi(X, 0)$ in closed form from a single order-flow draw $X \in \{-2, -1, 0, 1\}$ with $p_0 = 1 - \kappa$, $p_1 = \kappa/2$ and the action weights $(\omega_E, \omega_H, \omega_Q)$; Hold and Quiet Voice pooling because both trade $q = 0$.	DROP	Replaced by a pooled history $\mathcal{H}_H^P$ over $H + 1$ business days with plans, Γ -coarsened increments and a flag that can land mid-calendar. The closed-form posteriors do not survive the calendar. Consequential: L3 shows the manuscript's own no-disclosure structure (informed mark $\bar{z}$) yields a <i>four-atom</i> pooled law, two of whose atoms move with κ , so it lies outside $A(\tau)$ — the one-round posterior is not merely superseded, it is a named example of the representation failing.

draft_v2 object	Verdict	What carries over, what changes
The hump headline: $\Delta^{\min}(\kappa)$ single-peaked in liquidity, as the conditional-hump proposition and the paper’s first main result.	DROP	Not claimed anywhere in v4; the card states outright that it does not claim the hump survives. The v4 headline is the partition and the attenuation factorisation $\mathcal{S} = (1 - \Omega)\mathcal{S}_P$, not a shape in κ . Note the hump has not been <i>refuted</i> either: the enumerated two-round M_P is itself hump-shaped in κ at baseline (peak near 0.55) — it is simply no longer the result being sold, and its role now is to break $A(\tau)$ ’s orientation clause.
The unconditional attenuation claim: $ \partial\Delta^{\text{act}}/\partial\kappa $ decreasing in ω_P , asserted as a monotone consequence of shifting mass to the disclosed cell, with no composition term and no margin distinction.	DROP	Split and conditioned. The weight effect is kept and proved ($W_T \leq 1$, W_τ likewise), but the composition effect is a separate ratio: T1(B) gets threshold attenuation <i>unconditionally</i> because <i>both</i> W_τ and C_τ lie in $[0, 1]$; T1(C) makes the window margin an iff , $W_T C_T \leq 1$, with C_T unsigned. O-1 is a disclosure-regime analogy whose composition ratio can exceed one; it is not a window measurement. The genuine window comparison is <code>t2_t1_check</code> block 4. An unconditional window sign is not available.
Equilibrium uniqueness as a numerical regularity (a 30-seed multistart converging to one fixed point per κ , labelled a regularity rather than a theorem).	DROP	v4 does not claim uniqueness at all, as a theorem or as a regularity. The multistart survives only as an existence diagnostic in the check scripts, and it does not always clear its own binding criterion (Section 10, P1).
The welfare apparatus: transfer netting and the present-value form of total surplus, the liquidity planner’s wedge κ^* versus $\kappa^\dagger$, the disclosure-threshold planner, the first-best benchmark.	DROP	Out of scope by the card’s own statement: no welfare, no optimal rule design. Note this is also why a bare W is unavailable as a symbol — it was total surplus there, and it is the weight-effect ratio W_τ, W_T here.

10 Evidence

The check-script inventory

Eight scripts, each with committed JSON output, in `quality_reports/fixes/`. Verdicts are as the JSON reports them, failures included. “Wiring” and “substantive” are the design’s own classification: a wiring check runs both sides of an identity through the same enumeration and is therefore *not* evidence for the result, only for the code.

Script	Verdict	What it establishes, and what it does not
t2_d1_check	PASS 9 checks, 0 FAIL, 1 vacuous	Clock equivalence by three independent routes (date-by-date scan, one evaluation at $H - T$, closed-form crossing date) agreeing to exactly 0; partition exclusivity and exhaustion at 0 residual; the $T = H$ corner; robustness at $H = 12$. 50 nodes. <i>Vacuous</i> : the Q^F window-monotonicity check, because $\max Q^F = 0$ at the implemented calibration — the flagged round carries no residual order there, so the prediction has nothing to bite on.
t2_l1_check	PASS 4 checks, 0 FAIL	Decomposition residual $\Delta^{\text{act}} - (\Omega M_F + (1 - \Omega)M_P)$ at 0, both degenerate endpoints, and the non-evaluability of the null-cell average. <i>All four are classified WIRING by the design</i> : both sides run through the same enumeration, so the residual is machine noise by construction and is not evidence for L1 .
t2_l2_check	FAIL (1 of 5)	L2's own content passes exactly: the flagged range of M_F over 19 κ values is 0 and its κ -derivatives are 0 — the flagged path never touches the κ -dependent array. Two placebos confirm the pooled side does move. The failure is the fifth check , an <i>ancillary</i> $A(\tau)$ -orientation placebo: see Section 7.1, fact 2. Because $A(\tau)$ does not sign A'_κ (Example A has $A'_\kappa = -1/4$), the placebo is misformulated and does not decide $A(\tau)$'s support condition; it is not a failure of L2.
t2_l3_check	PASS 10 checks, 0 FAIL	The derivative identity, the mean-value form $C_h = \frac{1}{4}\bar{\pi}^2 h''(\zeta)$, the quadratic rate as $\bar{\pi} \downarrow 10^{-4}$, the affine-kernel zero-chord case, and three failure witnesses (a tent kernel with no root, the four-atom Example B gap, an unbounded- h'' witness). <i>Scope</i> : it runs on the three- and four-atom analytic laws and the standalone chord module — it does <i>not</i> enumerate the two-round model, so it verifies L3's mathematics, not $A(\tau)$'s applicability.
t2_l4_check	PASS 10 checks, 0 FAIL	All three sign legs with zero violations over 16 tightening steps (Section 7.1, fact 3); the exact reclassification identity at 10^{-16} ; flatness in κ at exactly 0 over 71 nodes; the τ -grid spanning a ninefold range in ω_a ; L3's quadratic corollary within 2.1%. Prediction 5 (the size of the A'_κ channel) is <i>reported, not gated</i> , by the request's own instruction: median absolute residual 0.0352, max 0.1213.

Script	Verdict	What it establishes, and what it does not
t2_p1_check	FAIL (1 of 10)	Nine substantive checks pass: inner-root single crossing and transversality, the flagged family single-valued with the right slope, sequential optimality of the flagged component, both cells on path, the threshold reformulation, monotonicity of Ω in (τ, T), and <i>existence at the core nodes</i> under the full 30-seed multistart. The failure is the grid sweep: 23 of 27 policy nodes met the binding payoff-scale criterion (10^{-9}) and 4 did not, at best payoff scales 1.5×10^{-3}, 1.1×10^{-3}, 4.0×10^{-4}, 3.1×10^{-4} — all four at extreme liquidity $\kappa \in \{0.15, 0.85\}$, at $\tau \in \{0.05, 0.075\}$ and $T \in \{5, 1\}$ (the independent re-run corrected an earlier gloss here: these are <i>not</i> zero-engagement-share or $T=H$ corner nodes — every actual corner row converged, and the recorded pooled prior engagement share is 0.102). Seed coverage is <i>ruled out</i> as the cause: the check re-ran all four nodes at the full 30 seeds and found identical best residuals, with most seeds landing on the same cutoff vector. All 27 nodes do converge on the cutoff-scale criterion. The cause is undiagnosed — candidates are a failure of A3's single crossing at those κ extremes (the adjacent-plan payoff gap is a sawtooth on this menu, and single crossing is a calibration fact, not structural) or a genuine boundary non-existence outside the repaired P1 hypothesis set. The check is retained as it ran, and P1 is PROVED: it was demoted on 2026-08-23 on the form-mismatch, sunk-cost and $\kappa = 1$ review findings, and restored on 2026-08-25 on its ticket-35 repair, the amended statement carried through a fresh two-pass gate; the 2026-08-21 chain is retained and did not satisfy the gate for the recorded statement. The four nodes are named, not smoothed.
t2_t1_check	FAIL (1 of 9)	The factorisation $\mathcal{S} = (1 - \Omega)\mathcal{S}_P$ holds pointwise to 1.2×10^{-16} and in total variation to 2.2×10^{-19}; Ω is flat in κ at 0 residual; the threshold margin, the window margin, the O-1 benchmark and the composition factors all pass; $H = 12$ robustness passes. The failure is block 3, the chord-magnitude route with hard-coded Example-A coefficient 0.25: see Section 7.1, fact 2. It is not a refutation of $A(\tau)$'s general support representation. One check is <i>vacuous</i> (the local form, which needs a smooth window interpolation the integer calendar does not provide).

Script	Verdict	What it establishes, and what it does not
t2_c1_region_check	PASS 12 checks, 0 FAIL, pointwise nodes <i>nonempty</i>	80 nodes ($8 \kappa \times 5 \tau \times 2 T$), 2.9 hours of wall time. Two gated verdicts, both PASS with 0 violations: the bound contains the four-corner re-solve remainder at the 8 validation nodes (max ratio 0.511), and it contains the implicit-function remainder at all 80 (max ratio 1.0). 18 pointwise dominance-and-contraction nodes , all at $T = 5$ and the upper three τ percentiles; $\eta_r \in [0.0595, 1.7227]$, median 0.3467; $L_{\mathcal{R}} \in [0.264, 0.501]$ with no node at $L \geq 1$. These nodes verify the two inequalities and supporting diagnostics only, not the full C1 antecedent or a named region. <i>Failure attribution</i> , recorded not hidden: 56 nodes have $g_r^{PE} = 0$ to 10^{-10} because the implemented legal clock quantises the crossing date, so $\Omega(\tau)$ is locally constant and the cross-derivative vanishes identically there; 4 are positive but dominated; 2 are negative; 40 nodes (all $T = 10$) are degenerate on flagged mass below 0.01. The sharper, non-inversion-free bound would certify 22 rather than 18 — reported as a measure of the Neumann step’s price, and certifying nothing. No tolerance was weakened to make the pointwise node set nonempty.

Provenance caveat. Seven of the eight scripts record the card stamp *2026-08-20* in their provenance block; only `t2_c1_region_check` was run against the 2026-08-21 stamp. The card changes between those stamps are hypothesis-explicitness moves, not model changes, but the mismatch is on the record and is not papered over here.

The numerical_v4 package: smoke facts

The implementation is `numerical_v4/` (modules `params`, `menu`, `policy`, `flagged`, `pooled`, `premium`, `solver`, `smoke`), with committed output in `numerical_v4/smoke_output.txt`. At the smoke configuration $H = 10$, $T = 5$, $M = 2$ order marks, $J = 3$ plans, $b_0 = 0.03$, $\bar{b} = 0.1$, seed $\tau = 0.05$ (the statutory 13D 5%):

- **Enumeration gate:** 4,194,304 histories, 826,686 feasible, $N_\theta = 12$; feasible $\times N_\theta = 9.92 \times 10^6$ against a gate limit of 10^8 . PASS. One evaluation takes 0.372 s; the whole smoke runs in 64.8 s.
- **τ freezing:** the seed equilibrium is $k = (1.31246481, 1.41888743)$ with $|k - \mathcal{T}(k)| = 1.34 \times 10^{-11}$; the median Voice $b^*(s)$ is 0.09076406, and τ is frozen there at every node.
- **Baseline equilibrium** ($\kappa = 0.5$): $k = (1.2405757283, 1.5310222869)$; outer residual 2.60×10^{-11} on the diagnostic cutoff scale and *exactly zero* on the binding payoff scale.
- **Masses:** $\Omega = 13.84\%$, $\mathbb{P}(a = 1) = 22.63\%$, $\omega_a = 61.15\%$, $\bar{\pi}_{\text{pr}} = 10.21\%$. *Note* $\omega_a = 61\%$ *here is a property of the calibration, not an estimate*: the card claims no empirical value for ω_a , and the empirical anchor for it is absent from the literature.
- **Premium objects:** $M_F = 0.5528$ pp, $M_P = 0.2228$ pp, $\Delta^{\text{act}} = 0.2685$ pp; L1’s decomposition residual is 0 (wiring).
- **Price path:** mean run-up $R = 389.18$ bp, mean filing jump $J = 2590.52$ bp, D1 identity residual 0 against a 10^{-12} tolerance.
- **A7 certificate:** passes — minimum $|db^*/ds|$ on the flagged set is 2.20×10^{-4} (gate 10^{-8}), no flat sub-intervals, no (B^F, Q^F) collisions, no cross-plan collisions, 10 flagged atoms of 19.
- **Numerical hygiene:** pooled mass error 2.2×10^{-16} ; maximum inner price residual 2.2×10^{-16} ; maximum b^* inversion error 5.9×10^{-14} ; GL-20 versus GL-40 quadrature on M_F

differing by 2.6×10^{-18} ; **zero multiple-root nodes** anywhere, which is the structural consequence of $m_0 \geq 0$.

- **L2 verdict, substantive:** the range of M_F across the frozen-policy κ sweep is *exactly* 0.
- **The shape that breaks $A(\tau)$:** M_P is hump-shaped in κ with its peak near $\kappa = 0.55$, so $\mathcal{S}_P = |\partial_\kappa M_P|$ is V-shaped about that peak, ranging from 4.99×10^{-3} to 2.66×10^{-1} premium pp per unit κ ; at the peak $\mathcal{S} = (1 - \Omega)\mathcal{S}_P = 6.37 \times 10^{-2}$.
- **Chord module:** $C_h \leq 0$ at every $\bar{\pi}$ tested, and $|C_h|/\bar{\pi}^2$ settles at ≈ 0.2219 as $\bar{\pi}$ falls from 10^{-1} to 10^{-4} — L3’s quadratic rate, on the standalone route.
- **One structural caveat:** $\max Q^F = 0$ at this calibration. The flagged round exists in the timing and in the code, but carries no residual order here, which is why the Q^F monotonicity check is vacuous.

Sources, all in the v4-theory worktree: `research/model_v4/MODEL_CARD.md` (stamp 2026-08-28, follow-up curation) for Sections 1–6 and 7; `research/model_v4/LABEL_LEDGER.md` for the labels; `research/model_v4/proofs/` and `research/model_v4/rederive/` for the two-pass evidence chains; `research/model_v4/HANDOFF_sign.md` and `quality_reports/fixes/t1_o1_rerun_check.json` for the O-1 numbers; `quality_reports/fixes/t2_*_check.json` for the check inventory; `numerical_v4/smoke_output.txt` for the smoke facts.